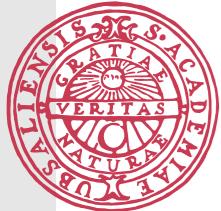


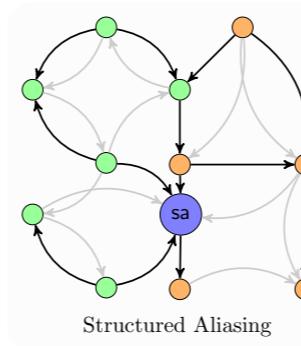
# Relaxed Linear References for Lock-free Data Structures

Elias Castegren, Tobias Wrigstad

ECOOP'17, Barcelona



**UPMARC**





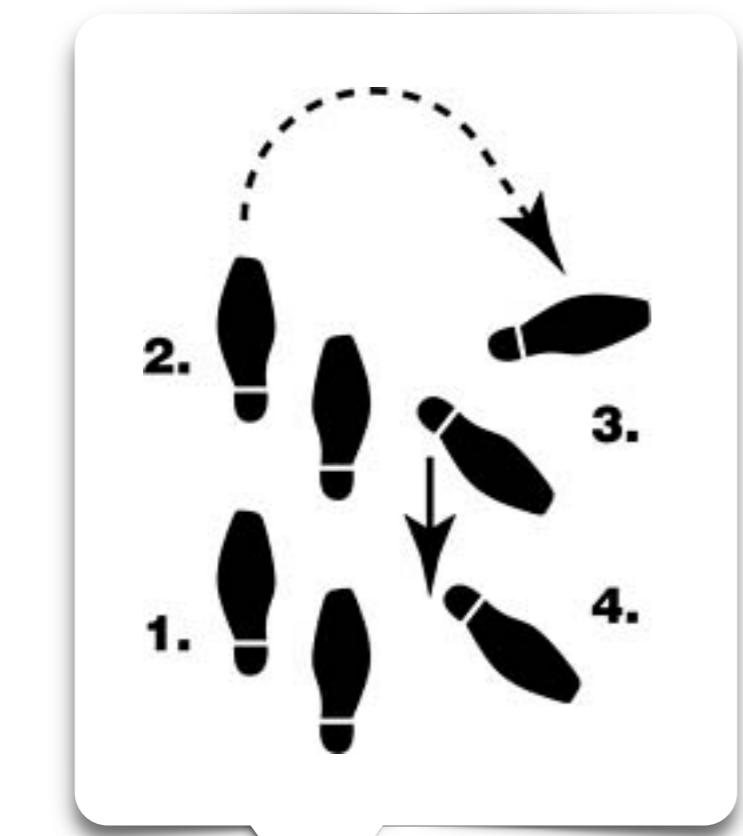
<http://www.kristianstadik.se/www/wordpress/wp-content/uploads/2016/06/midsommar.jpg>

<http://www.vikstroms.nu/cms/wp-content/uploads/2014/06/Midsommar-2014.png>

<http://pix.tagapi.aller.se/sf.php?src=http://stilexperten.mabra.com/wp-content/uploads/sites/36/2016/06/383ab62d304910ab9e9bb9a0f197d918.jpg&w=640&h=400&a=c>



<http://2.bp.blogspot.com/-ed71CELc5pY/T-h5pPLBR-I/AAAAAAAACK/9hILp4JdDJg/s1600/annicasxperia+juni+026.jpg>



”Små grodorna är lustiga att se”



[https://upload.wikimedia.org/wikipedia/en/e/e7/The\\_Swedish\\_Chef.jpg](https://upload.wikimedia.org/wikipedia/en/e/e7/The_Swedish_Chef.jpg)

[https://68.media.tumblr.com/5cab6229774bfb3fa6a35bf3479b6f8c/tumblr\\_inline\\_nht61yoHMA1r0y63m.jpg](https://68.media.tumblr.com/5cab6229774bfb3fa6a35bf3479b6f8c/tumblr_inline_nht61yoHMA1r0y63m.jpg)

[https://s3.amazonaws.com/gs-geo-images/3738763b-2dc4-47fb-81d2-345908ddbc07\\_l.jpg](https://s3.amazonaws.com/gs-geo-images/3738763b-2dc4-47fb-81d2-345908ddbc07_l.jpg)



# Lock-free Programming

---

- Fine-grained concurrency achieved by following some protocol
  1. Read a shared value (speculation)
  2. Perform some local computation
  3. If the speculation is still valid, make the local value global (publication)
- Requires atomic operations, e.g. CAS (compare-and-swap)

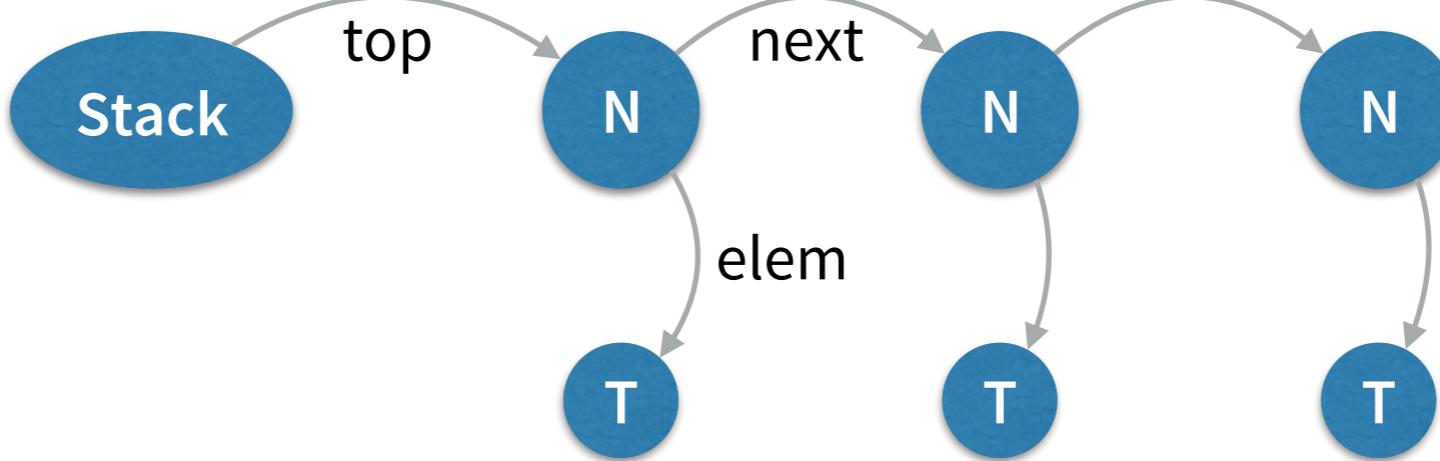
```
def CAS(x.f, y, z) : bool {  
    if (x.f == y) { }  
    x.f = z;           } } Performed atomically  
    return true;  
} else  
    return false;  
}
```

# Example: A Lock-free Stack [Treiber]

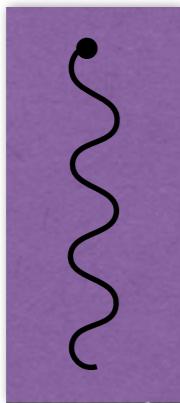
---

```
struct Stack {  
    var top : Node;  
}
```

```
struct Node {  
    var elem : T;  
    val next : Node;  
}
```



# Example: A Lock-free Stack [Treiber]



```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAS(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

Stack

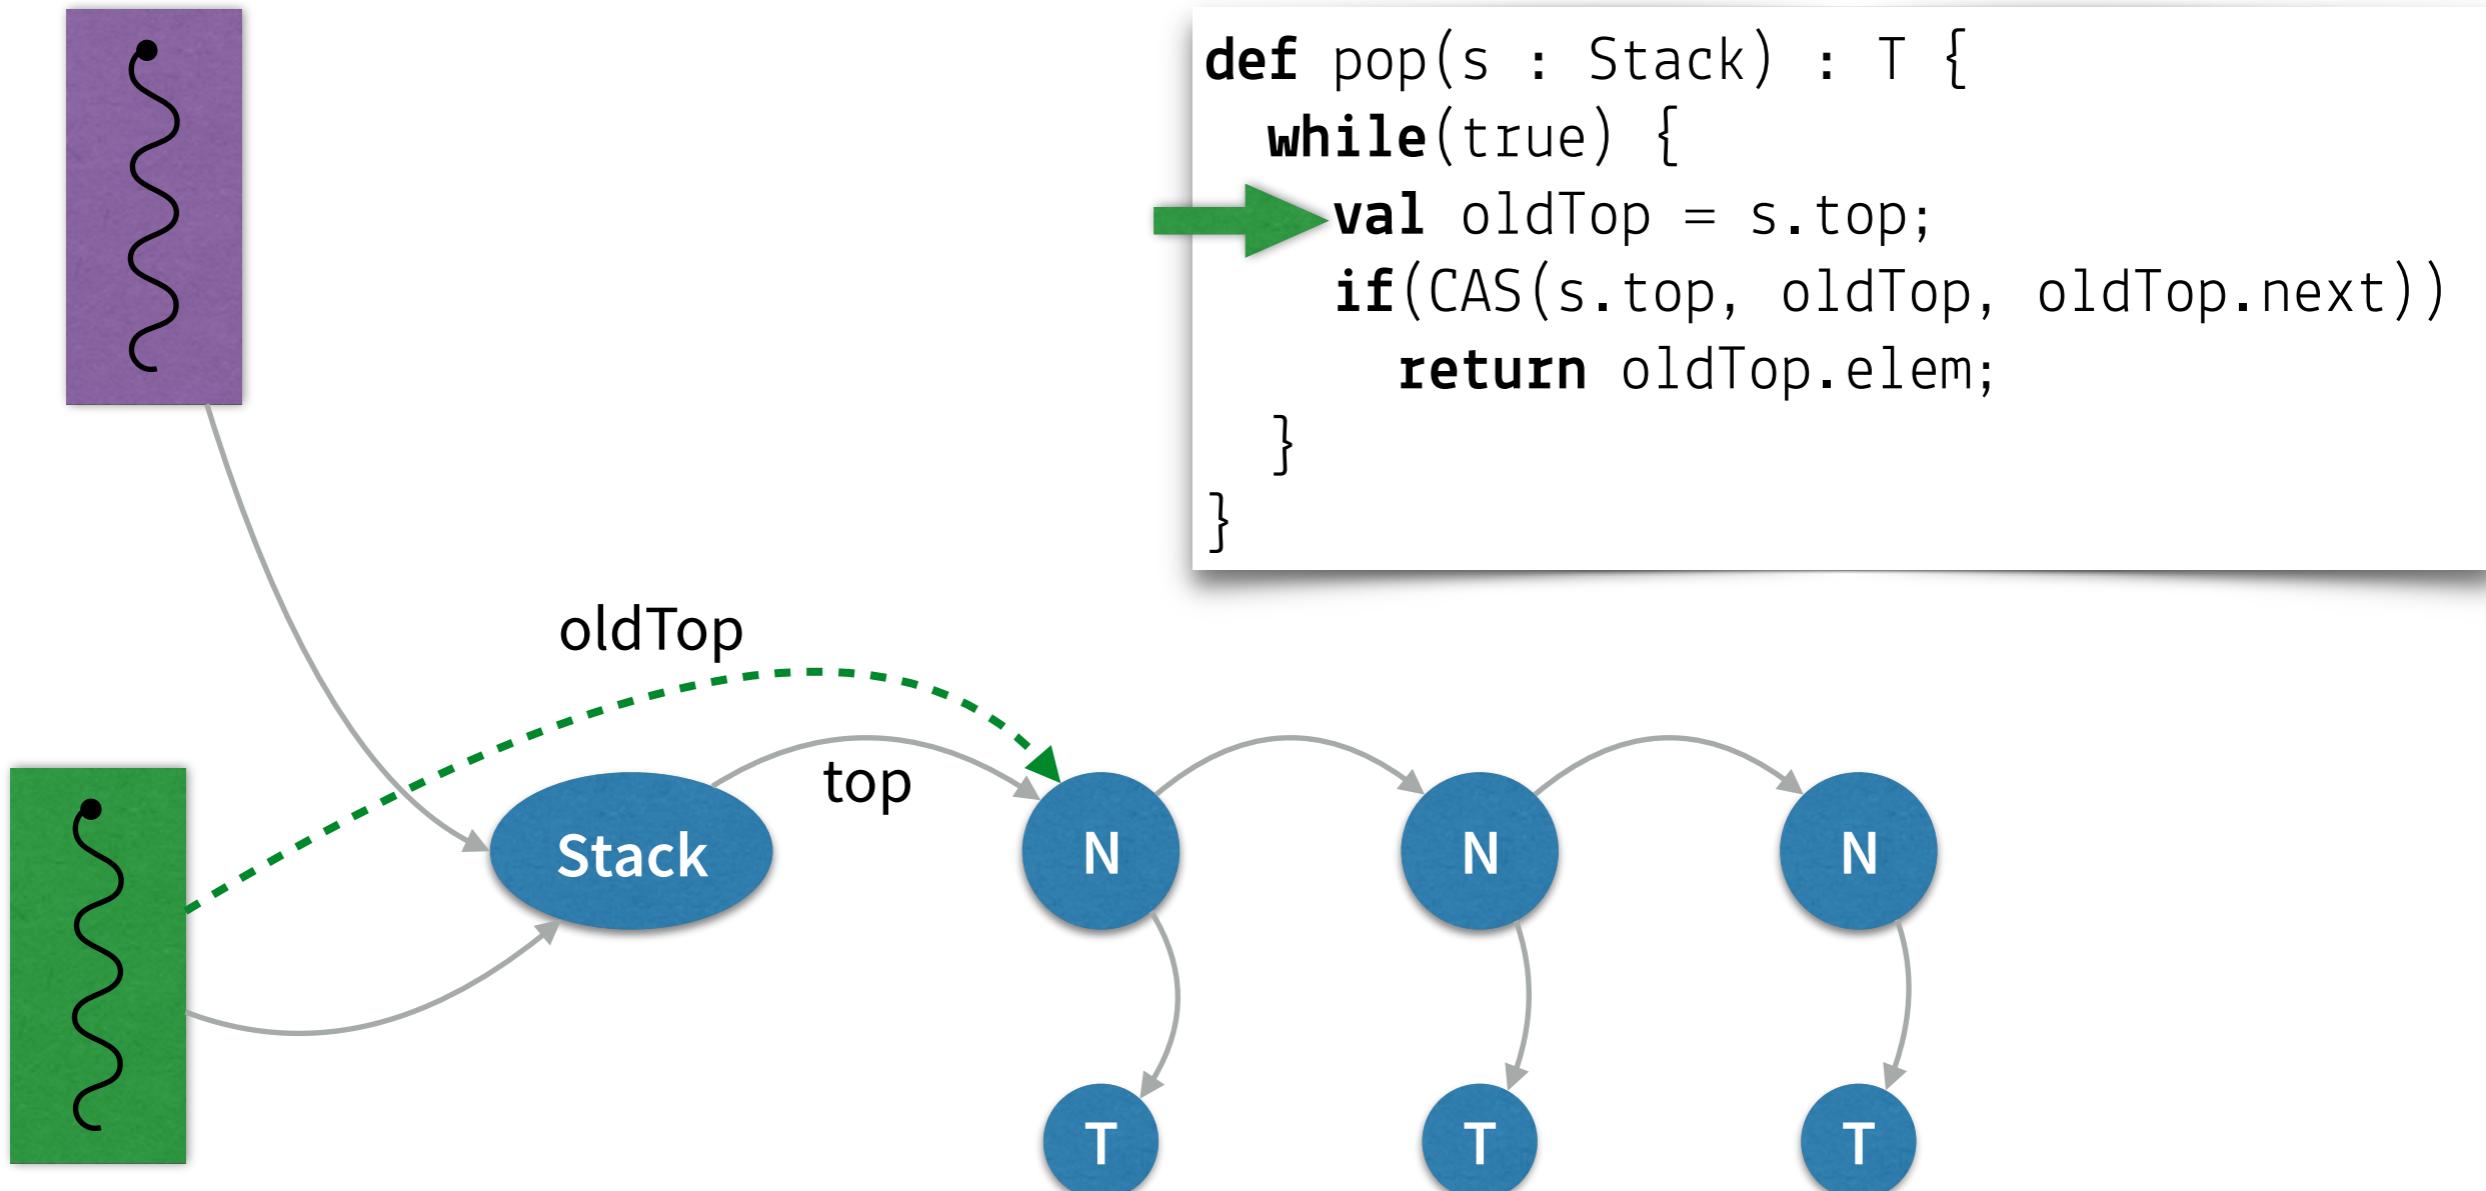
top

N  
T

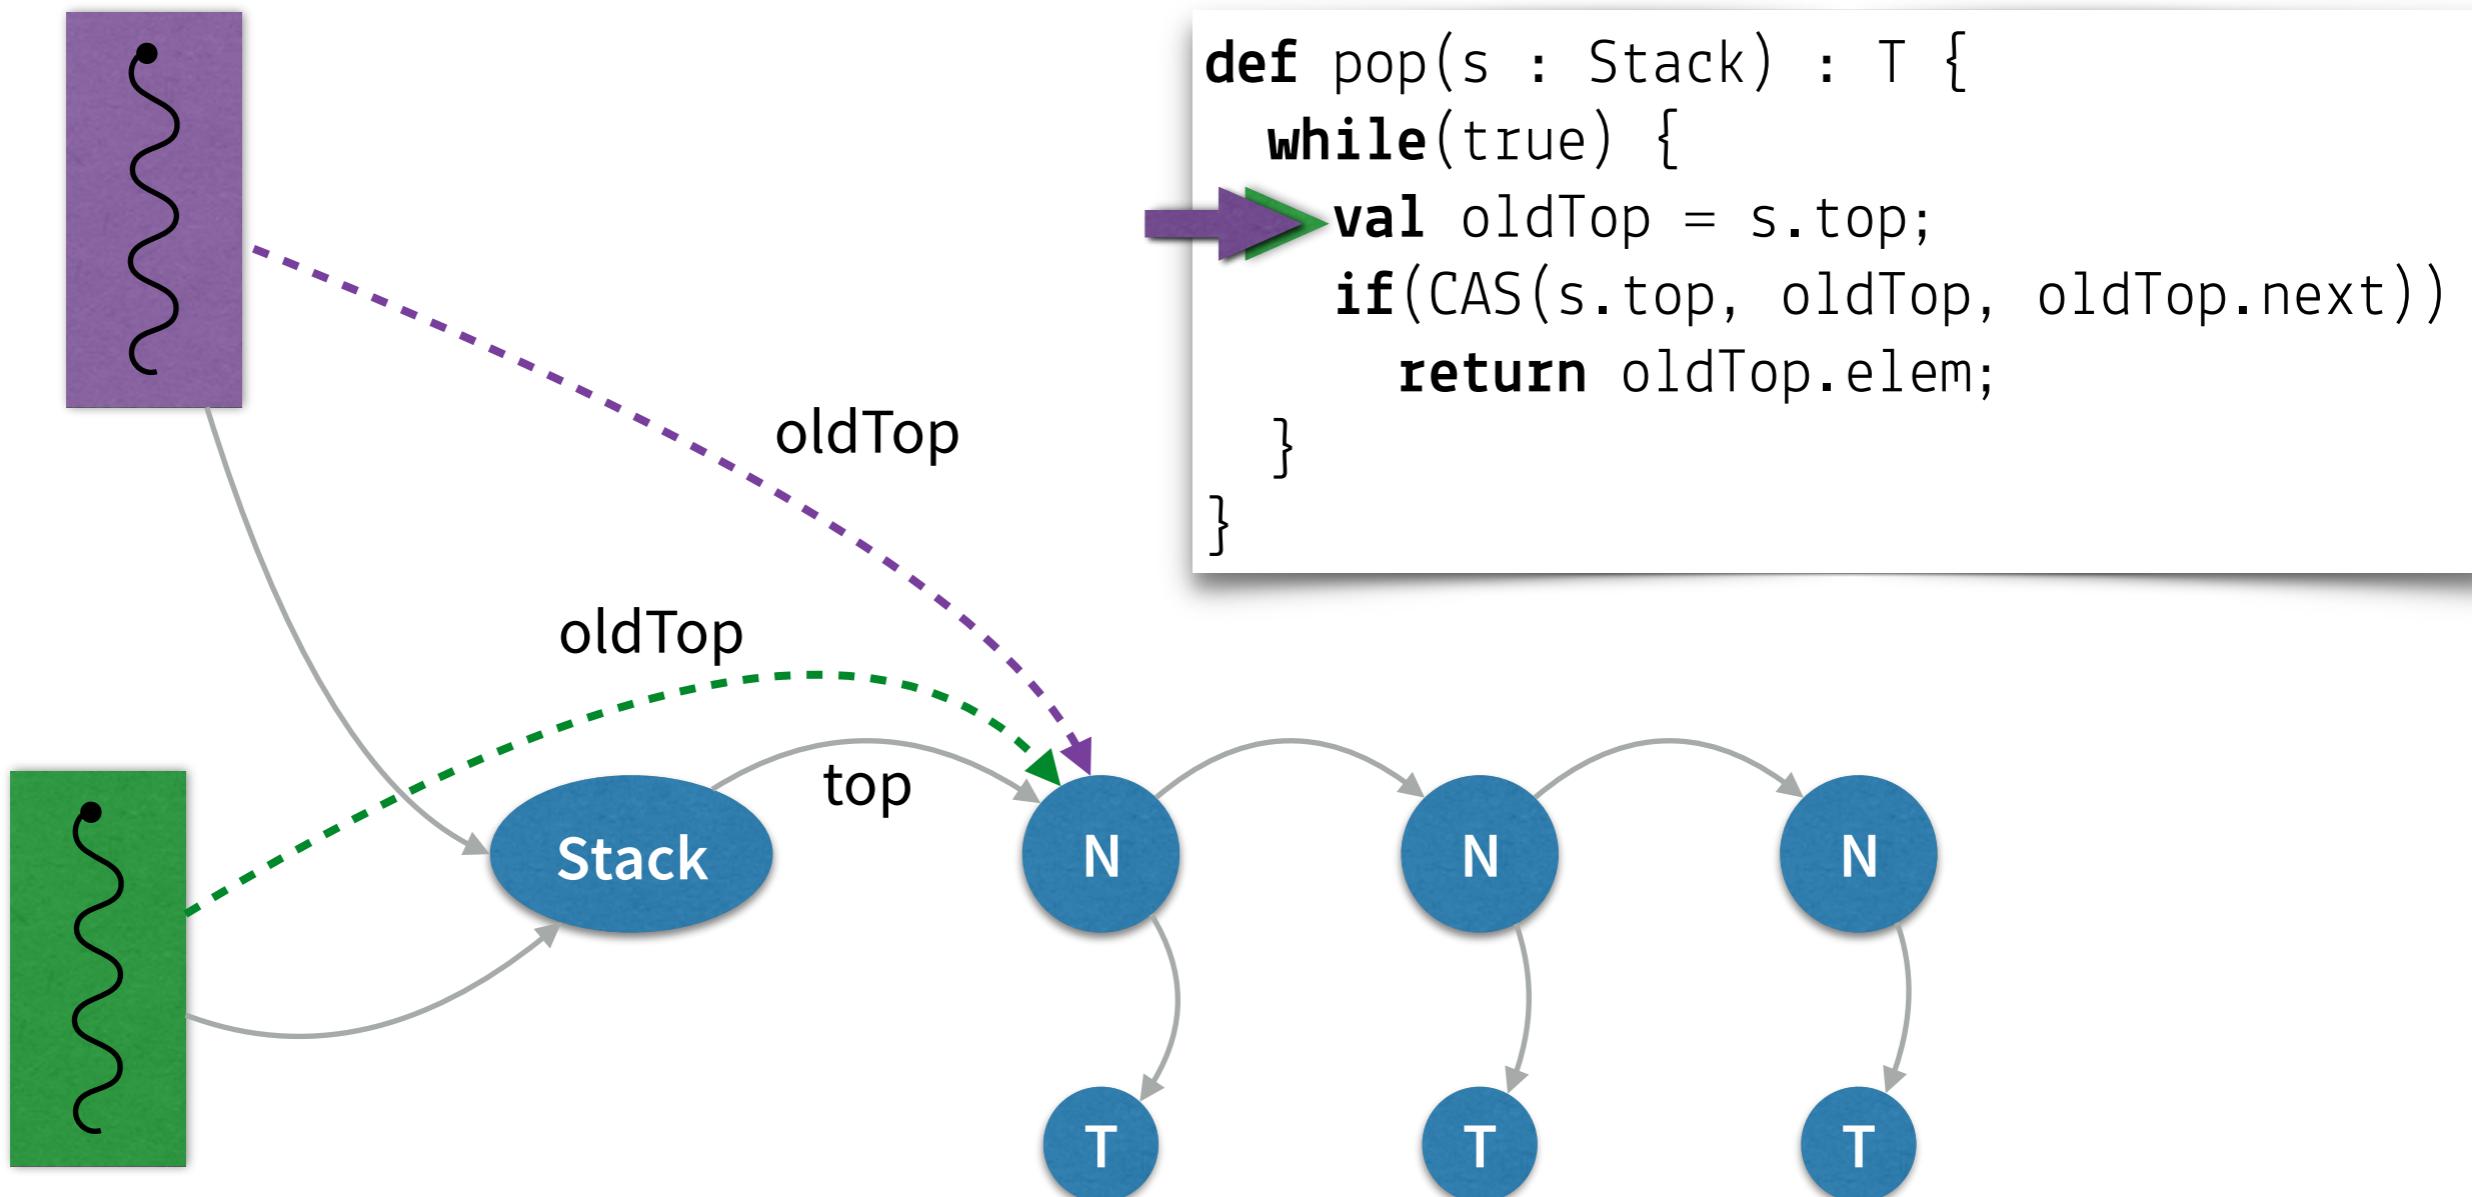
N  
T

N  
T

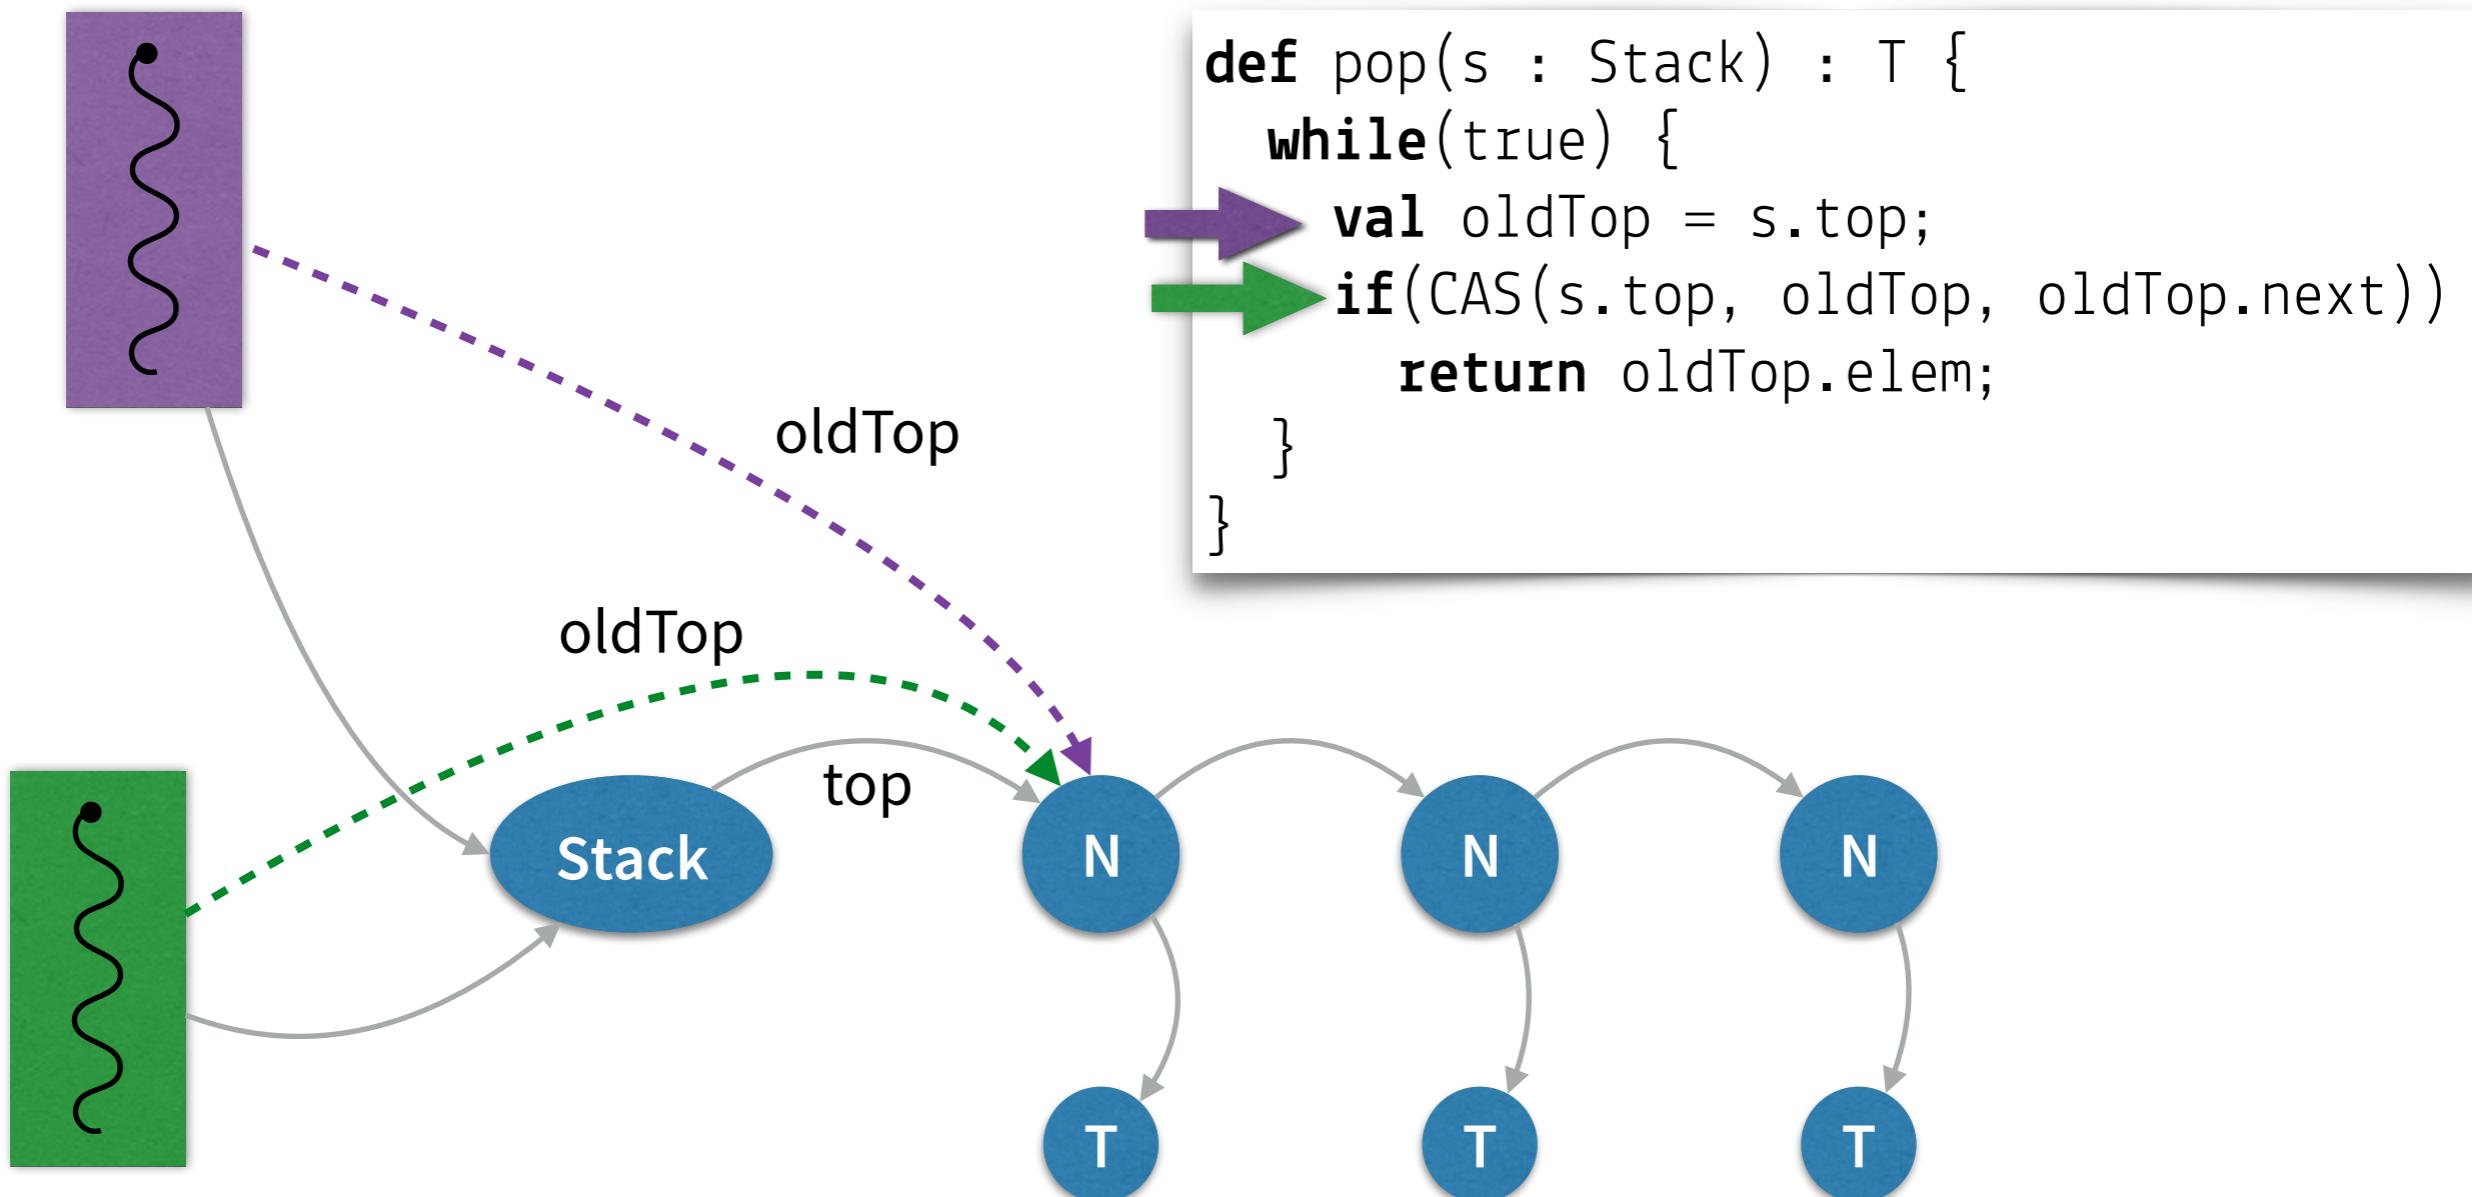
# Example: A Lock-free Stack [Treiber]



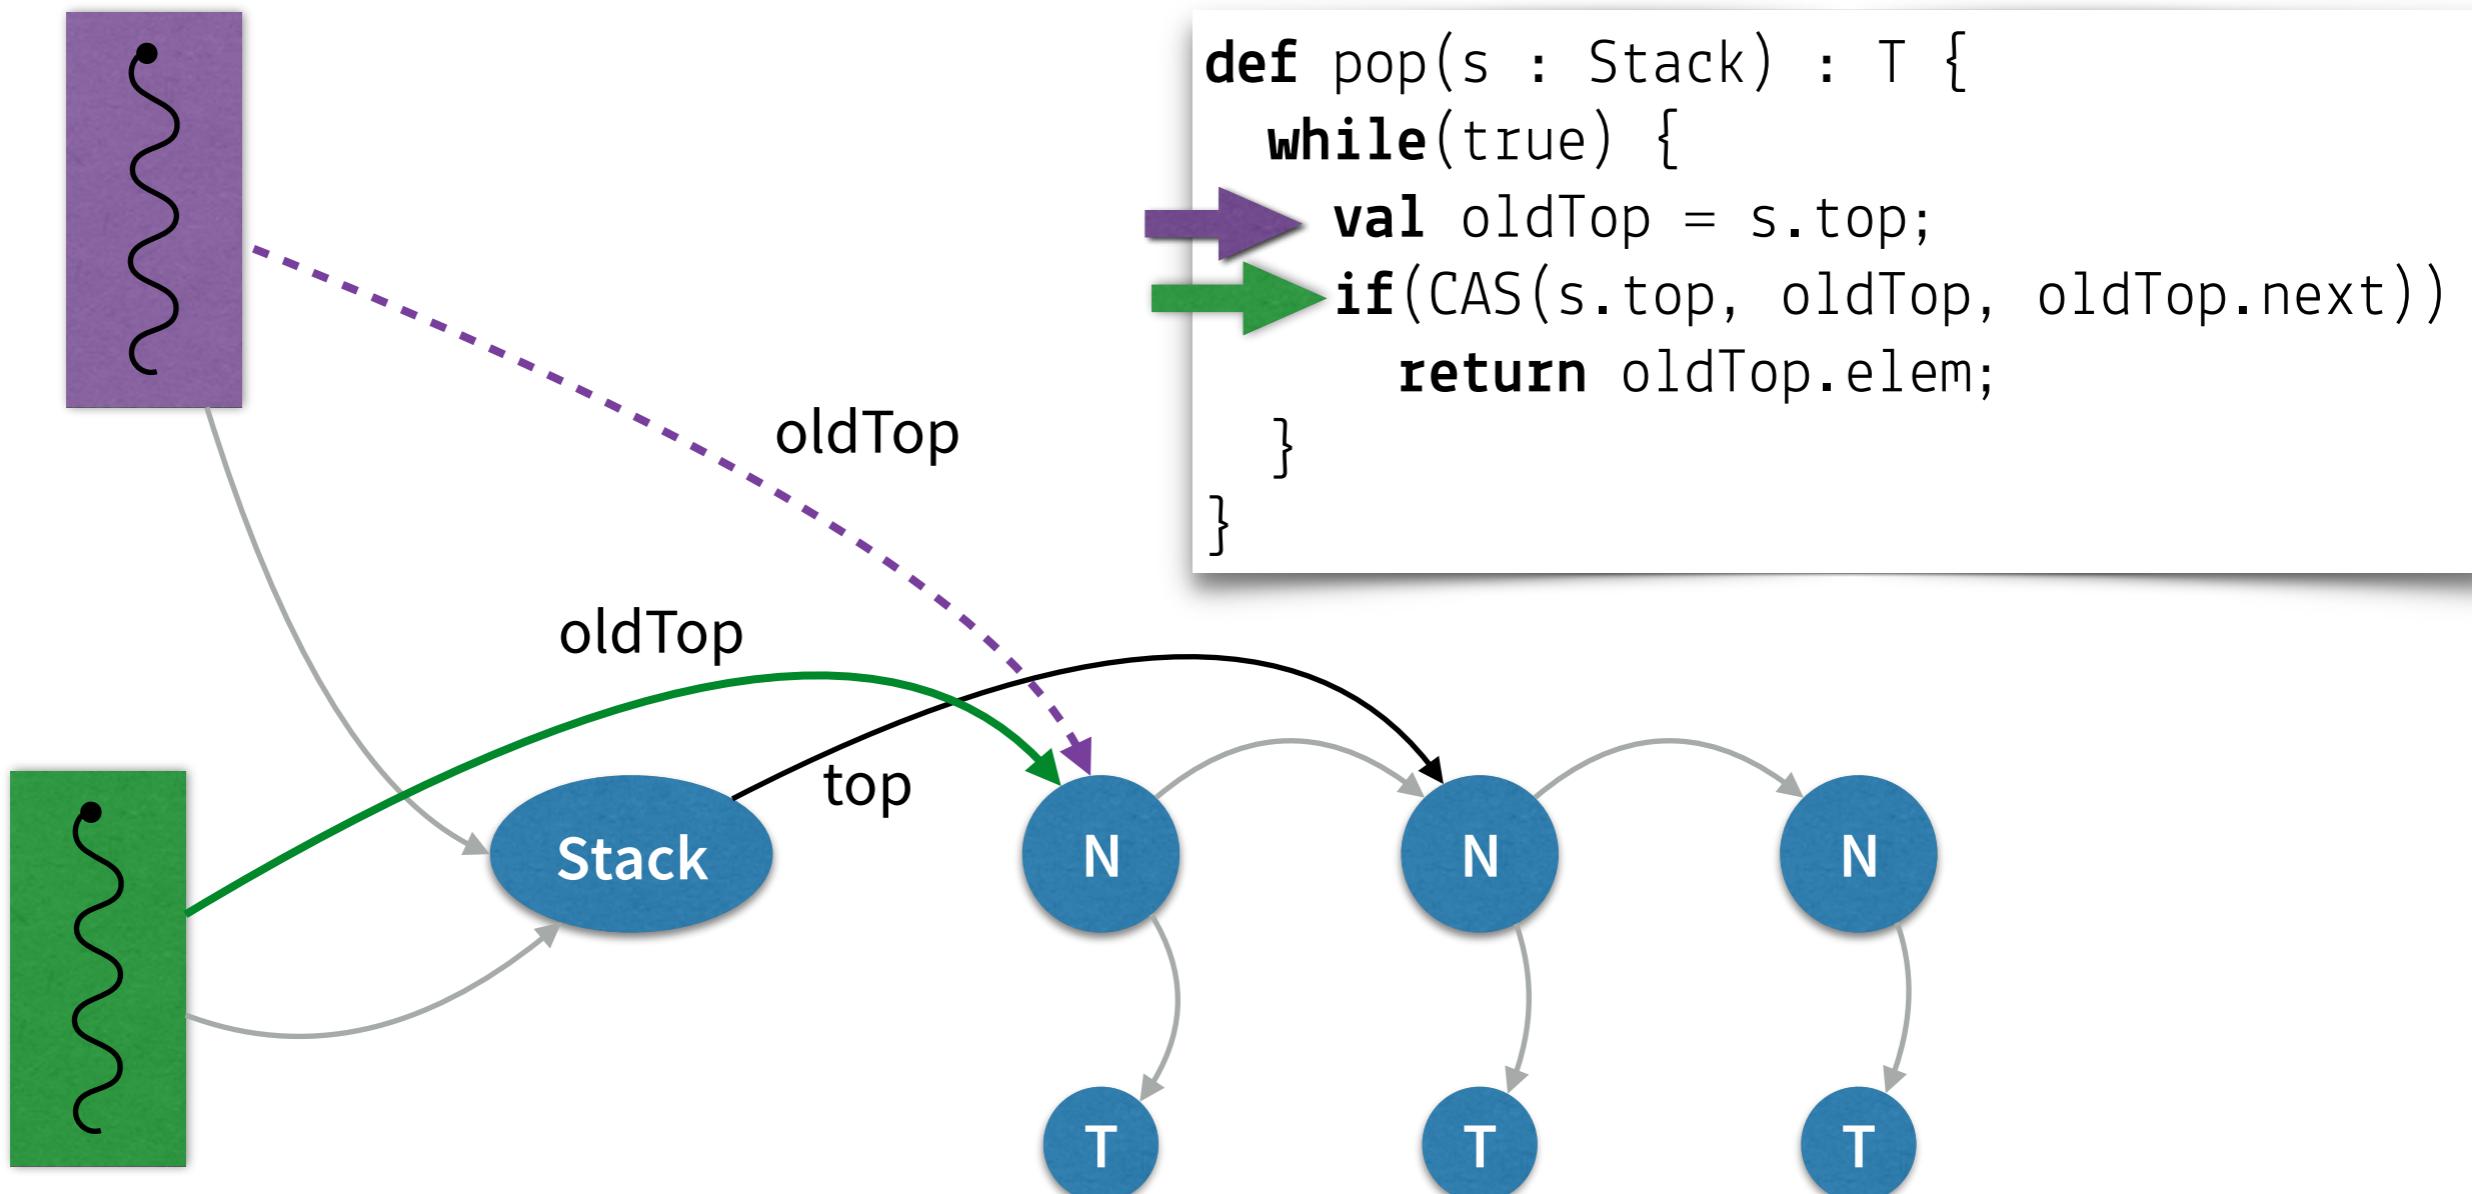
# Example: A Lock-free Stack [Treiber]



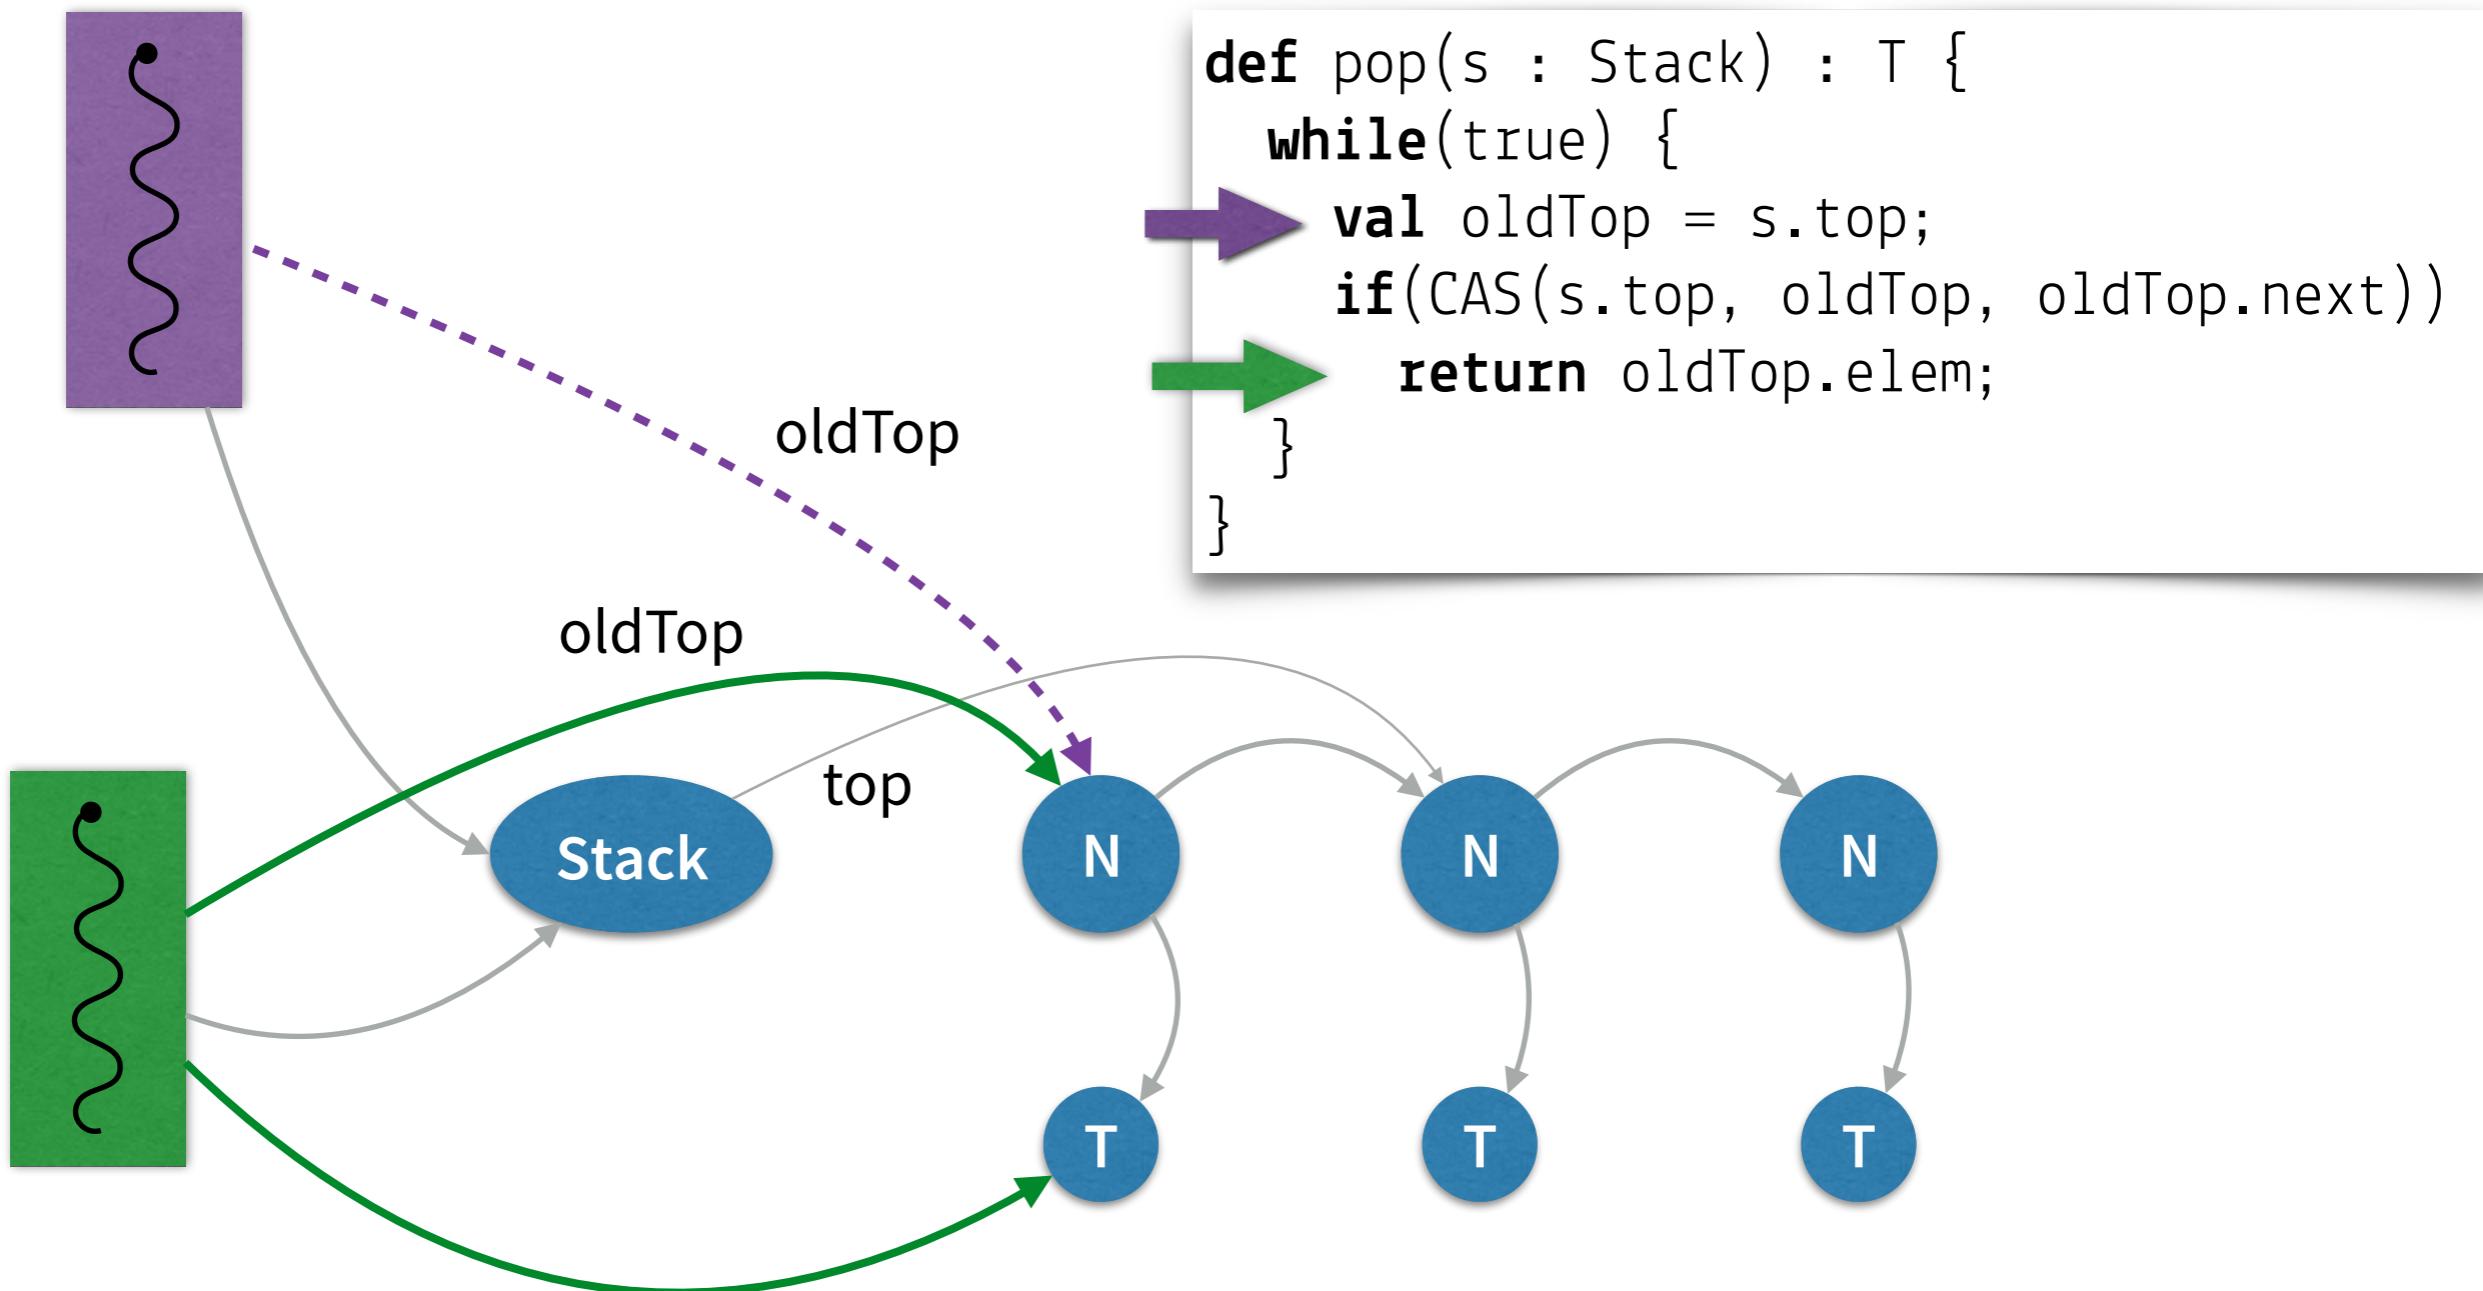
# Example: A Lock-free Stack [Treiber]



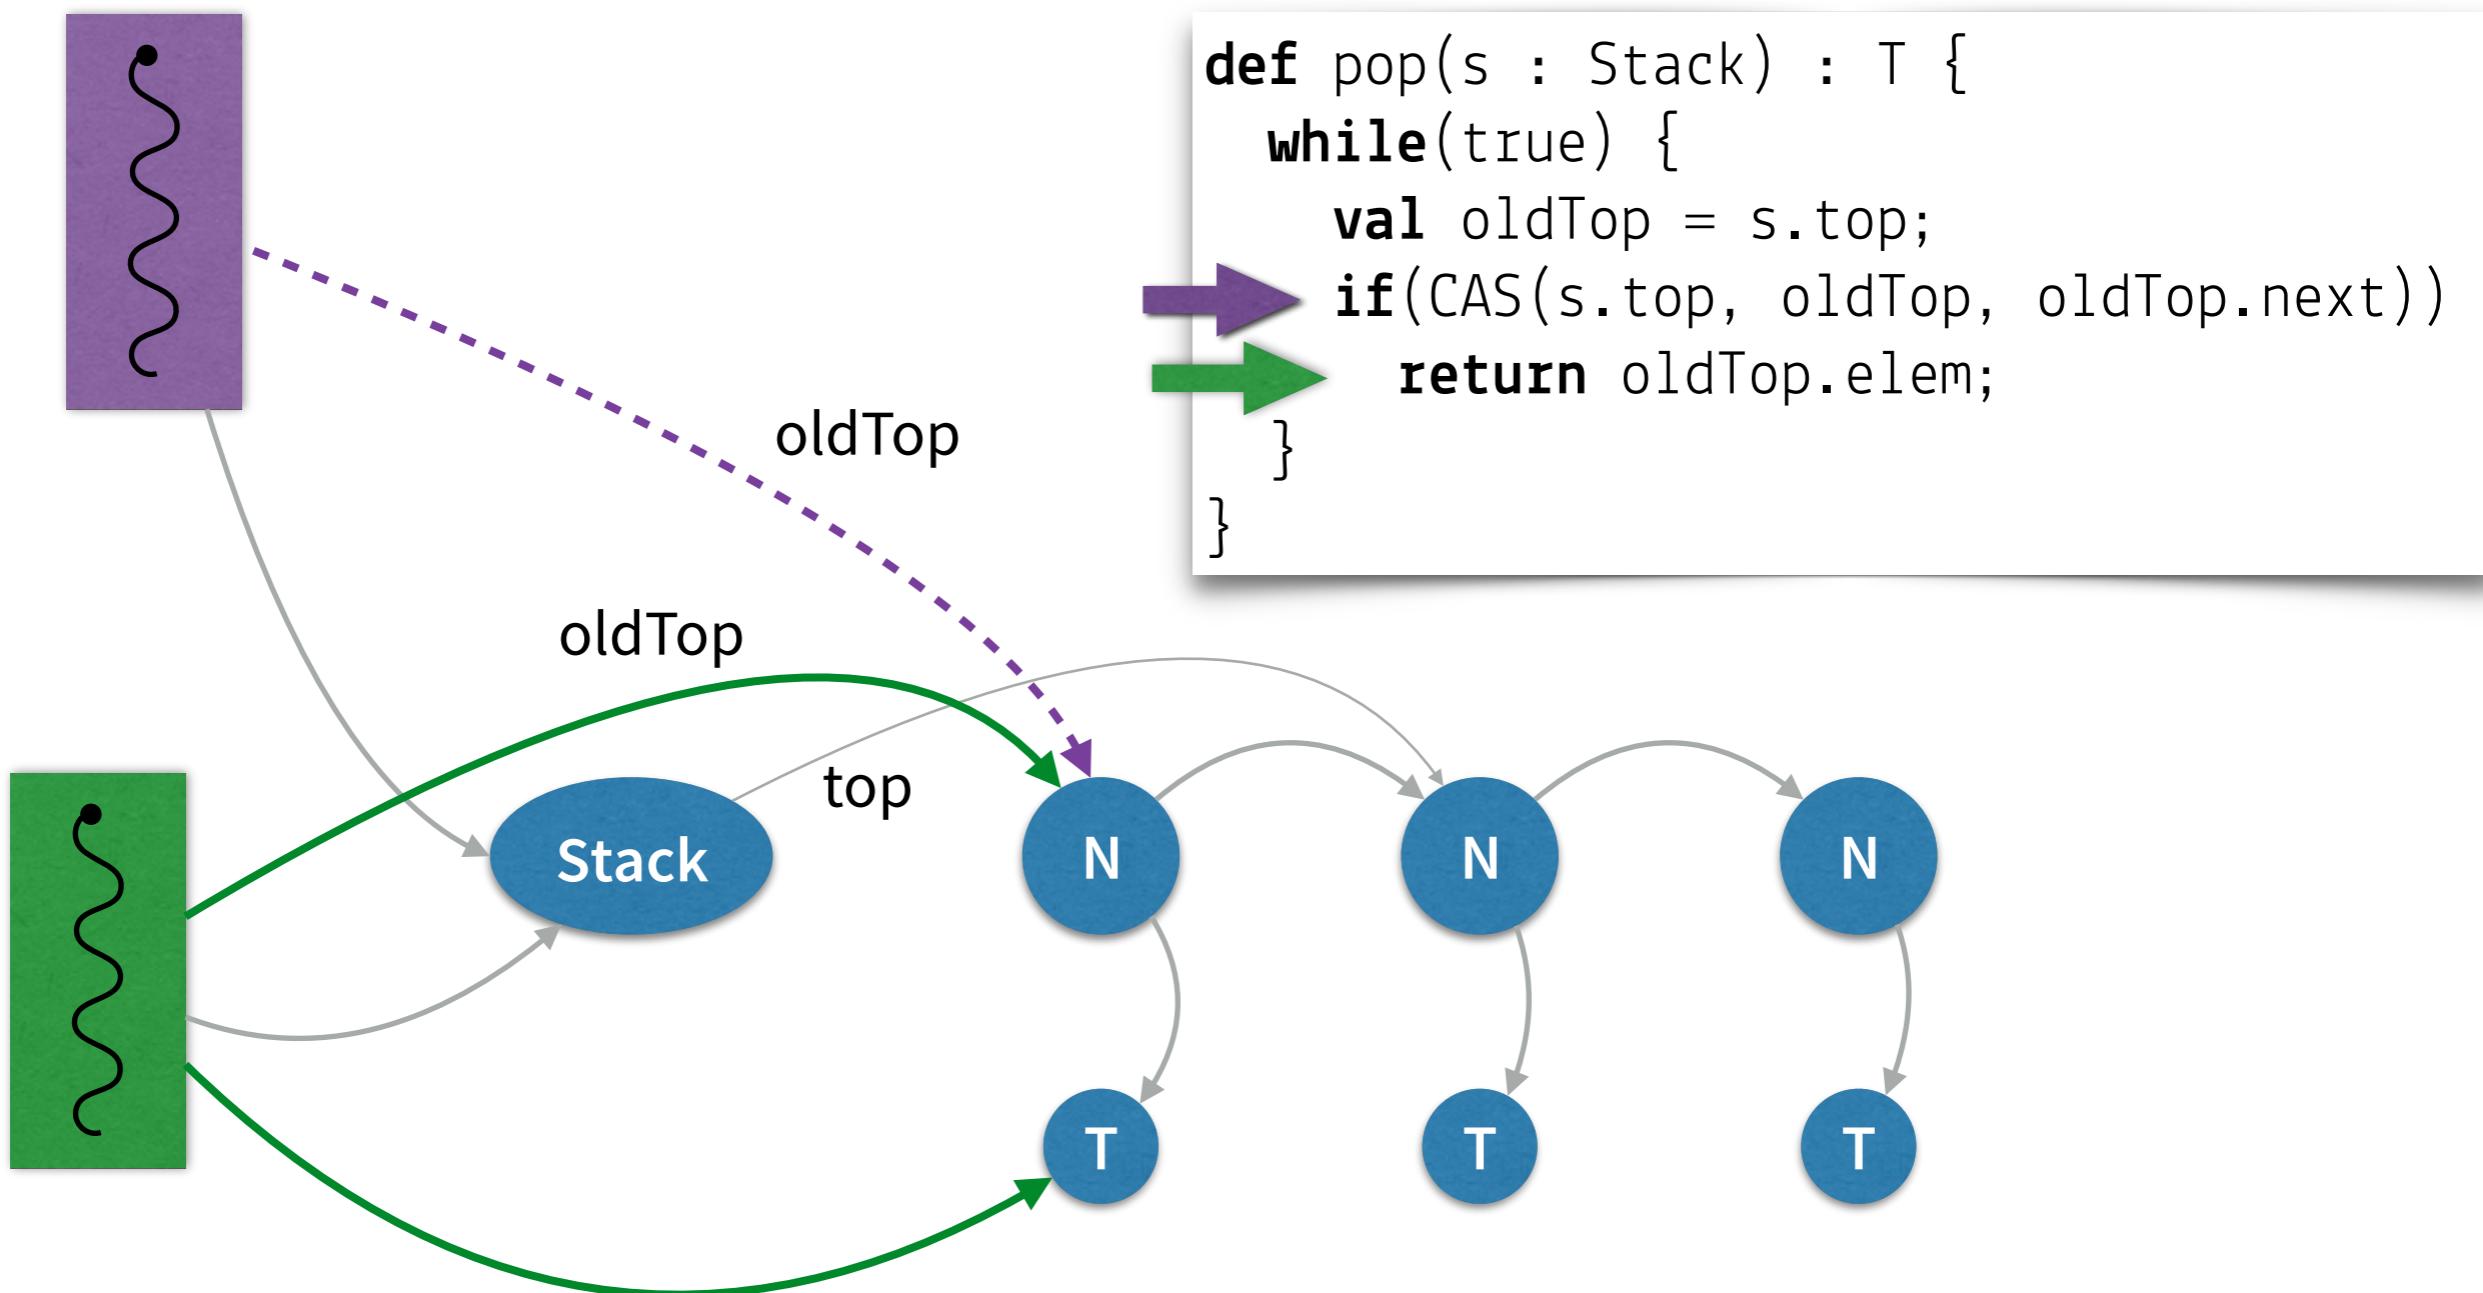
# Example: A Lock-free Stack [Treiber]



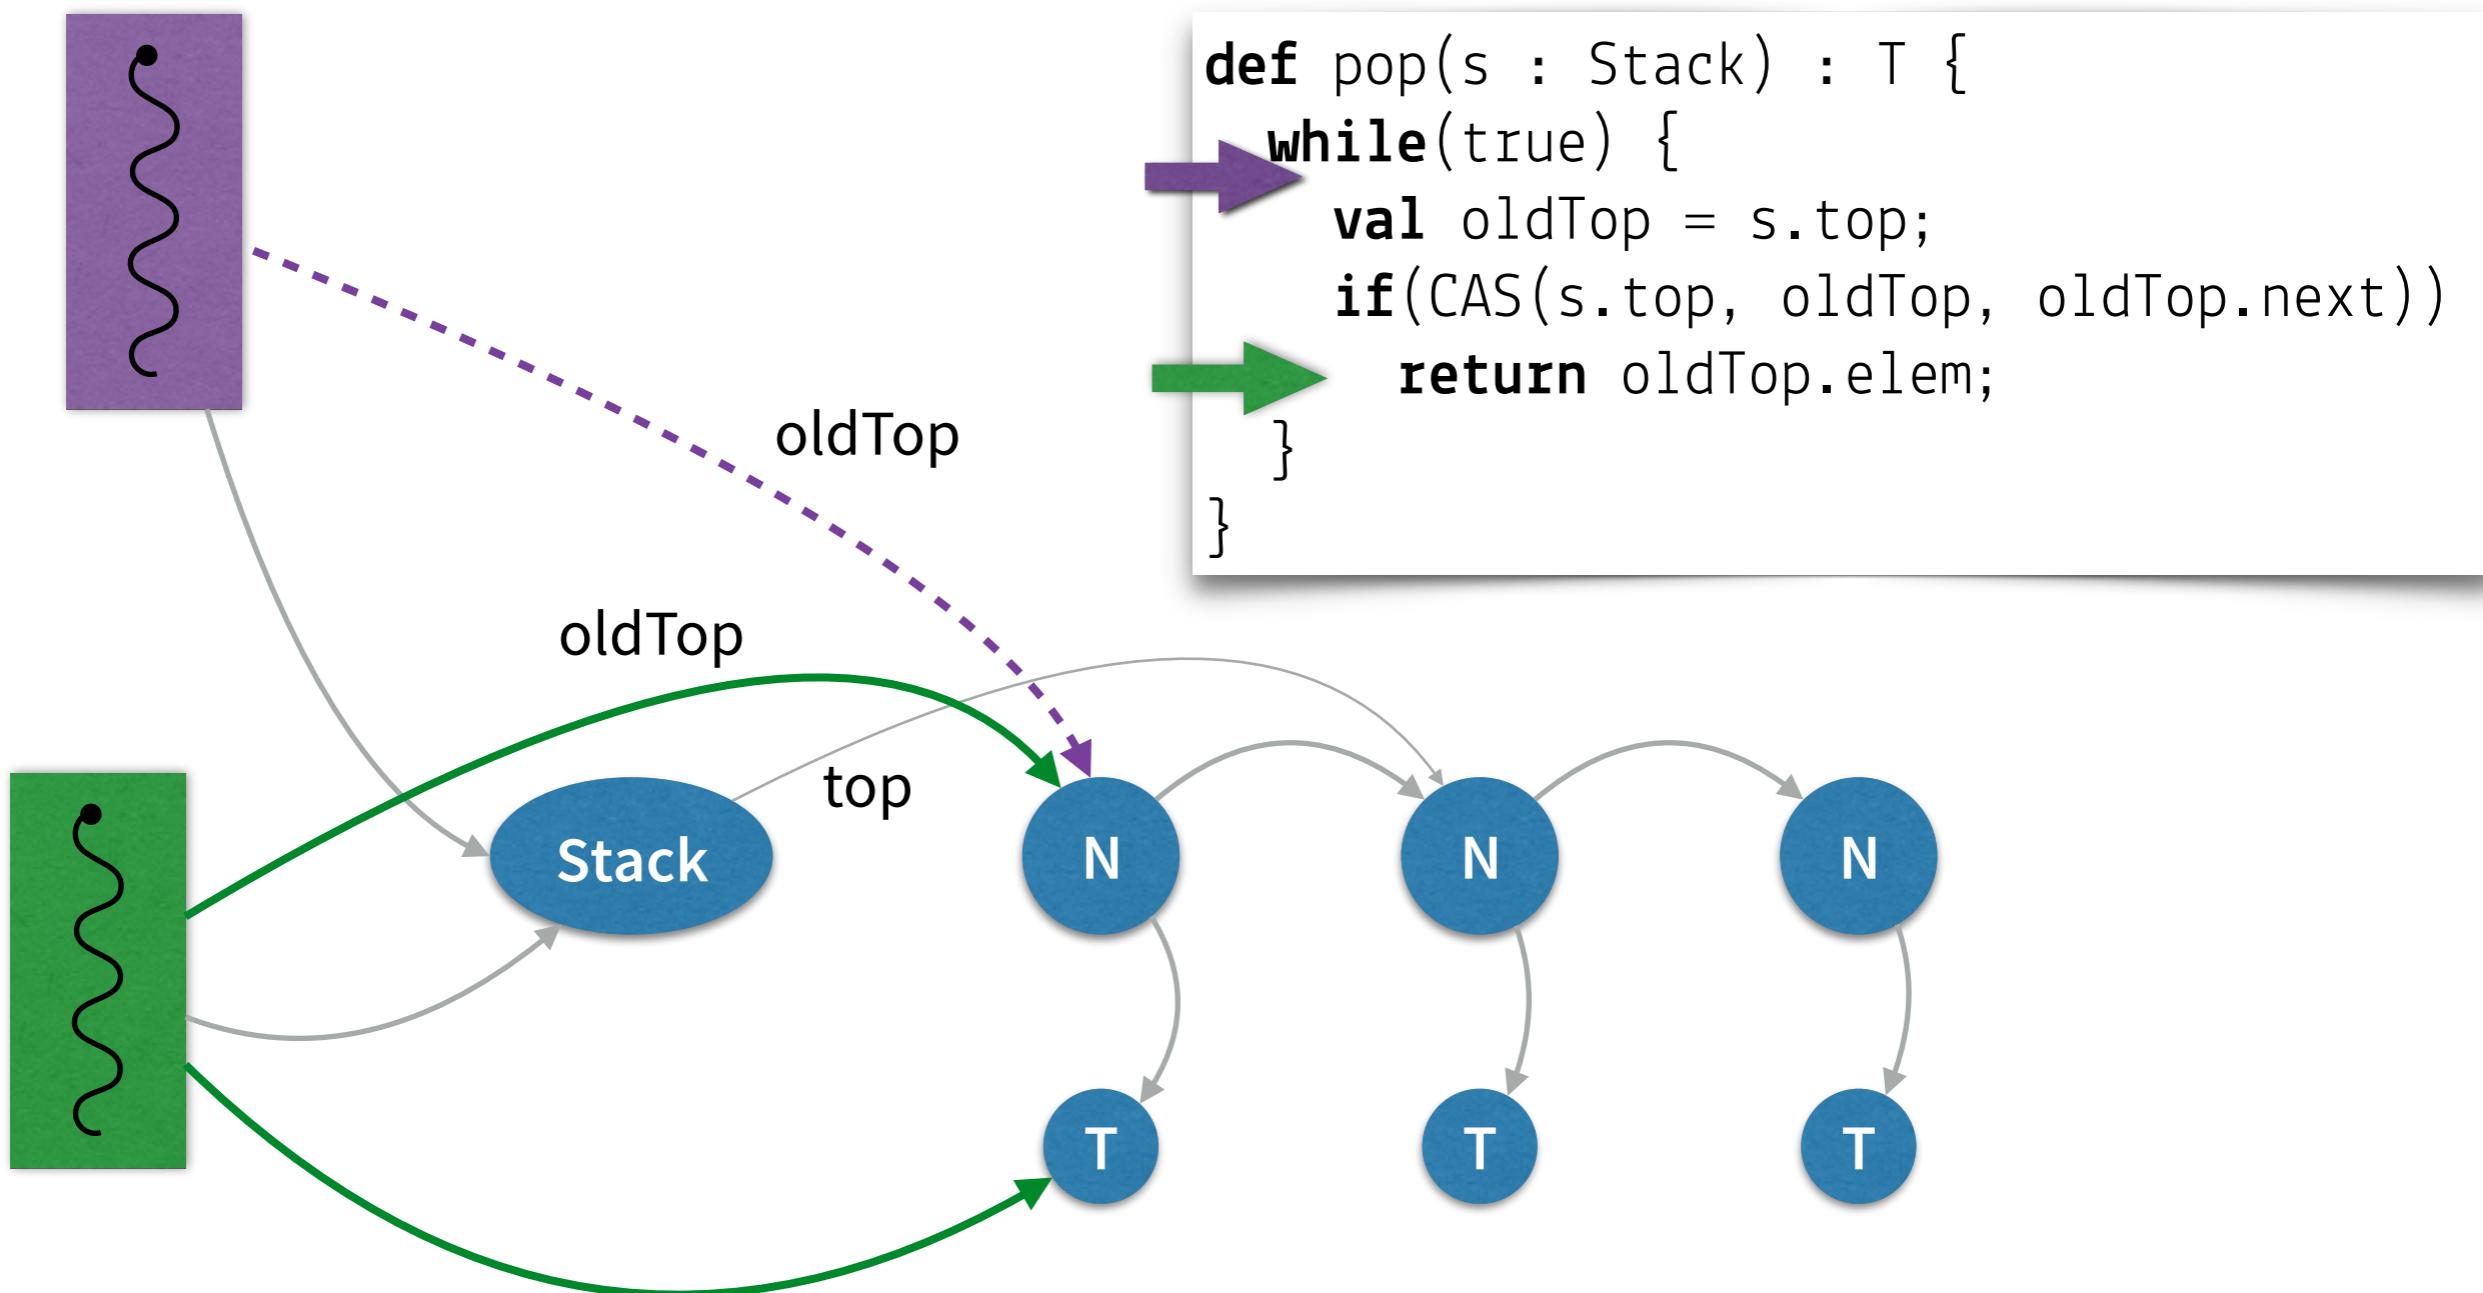
# Example: A Lock-free Stack [Treiber]



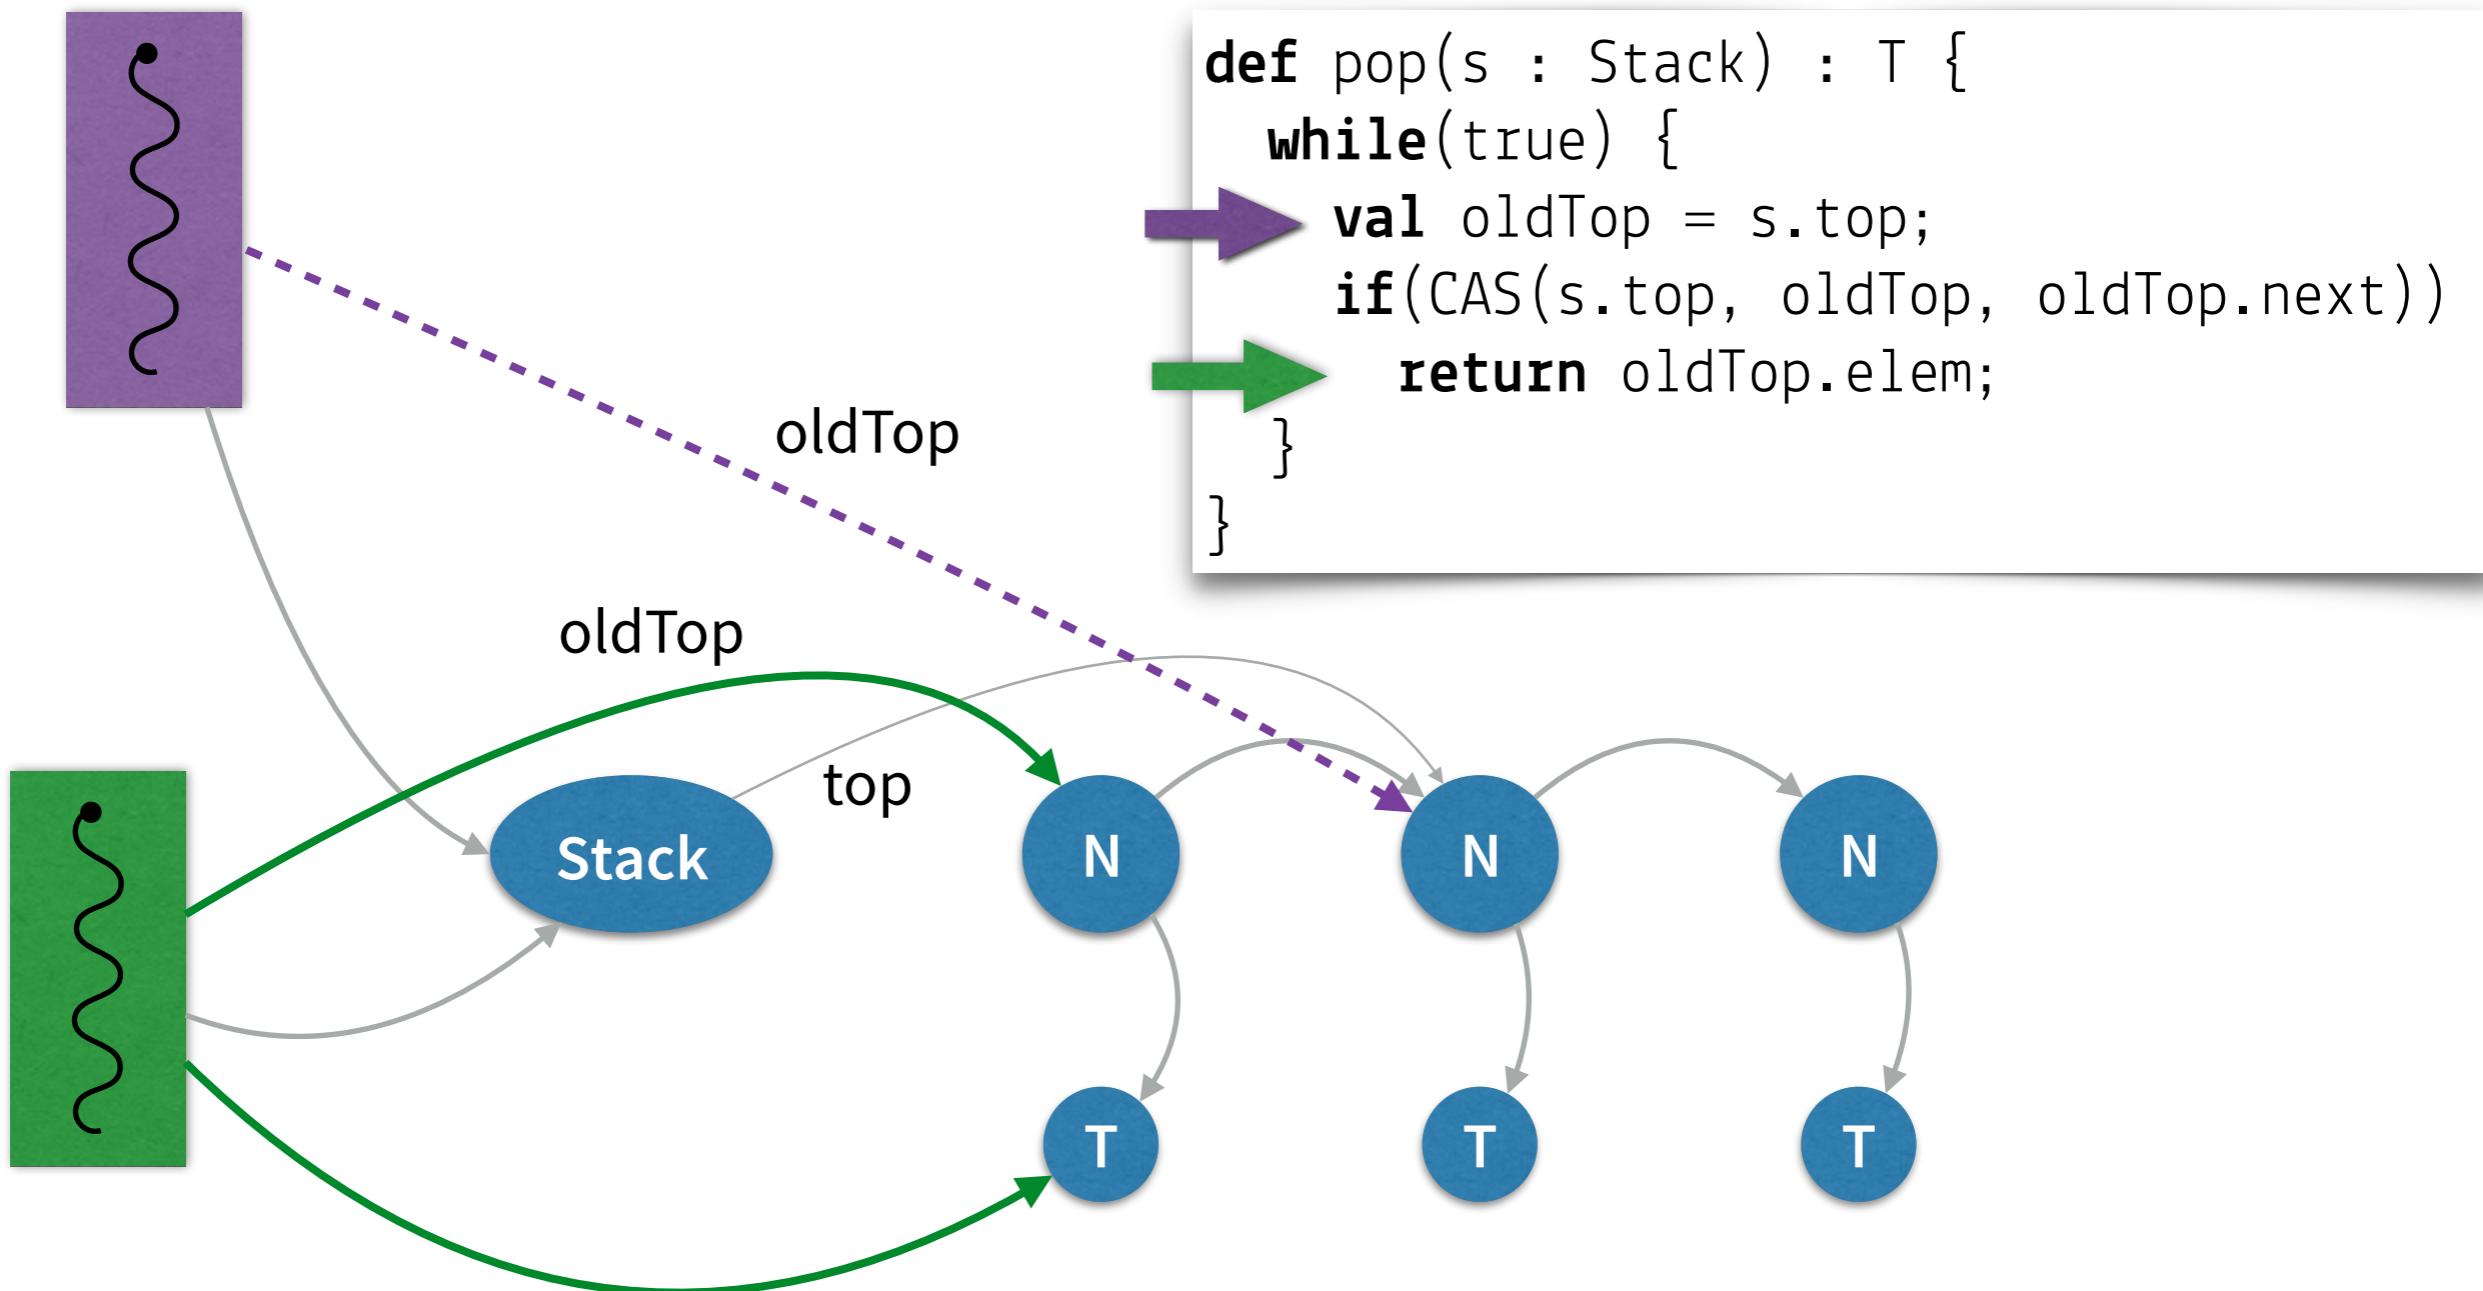
# Example: A Lock-free Stack [Treiber]



# Example: A Lock-free Stack [Treiber]



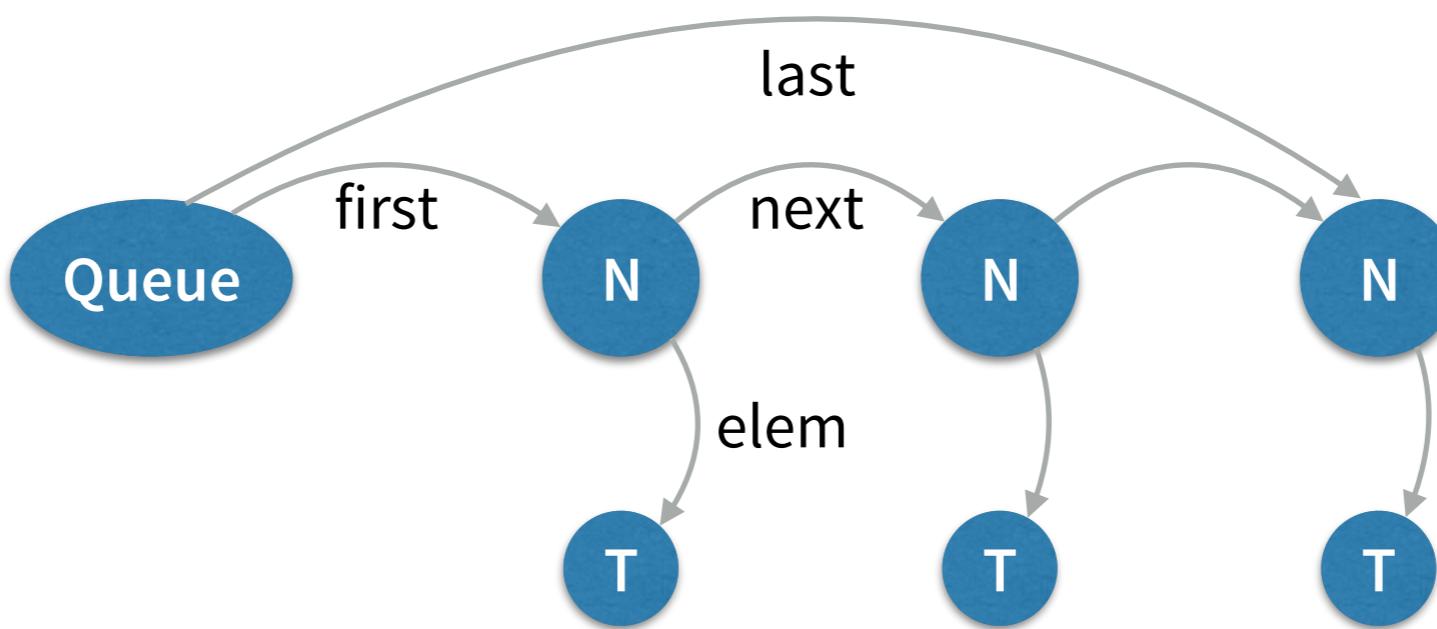
# Example: A Lock-free Stack [Treiber]



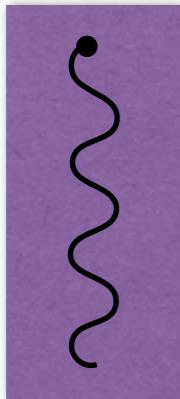
# Example: A Lock-free Queue [Michael & Scott]

```
struct Queue {  
    var first : Node;  
    var last : Node  
}
```

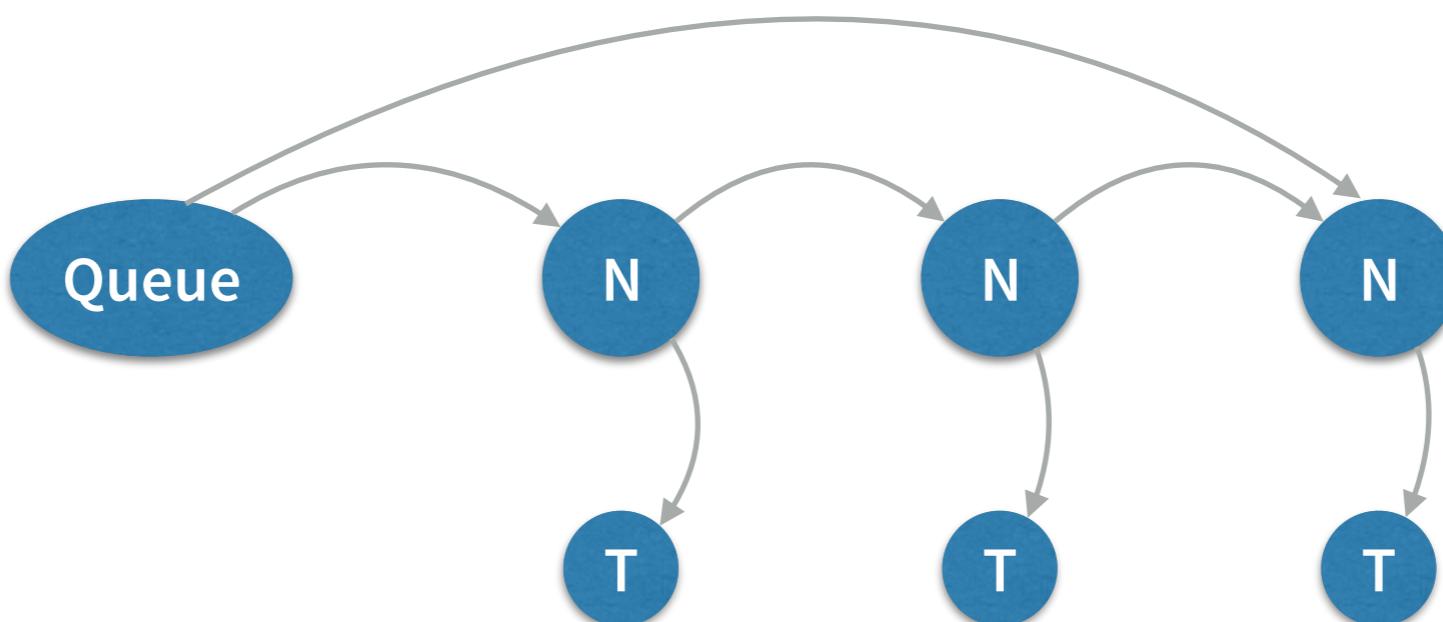
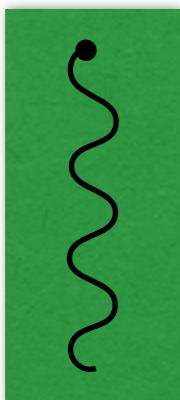
```
struct Node {  
    var elem : T;  
    var next : Node;  
}
```



# Example: A Lock-free Queue [Michael & Scott]

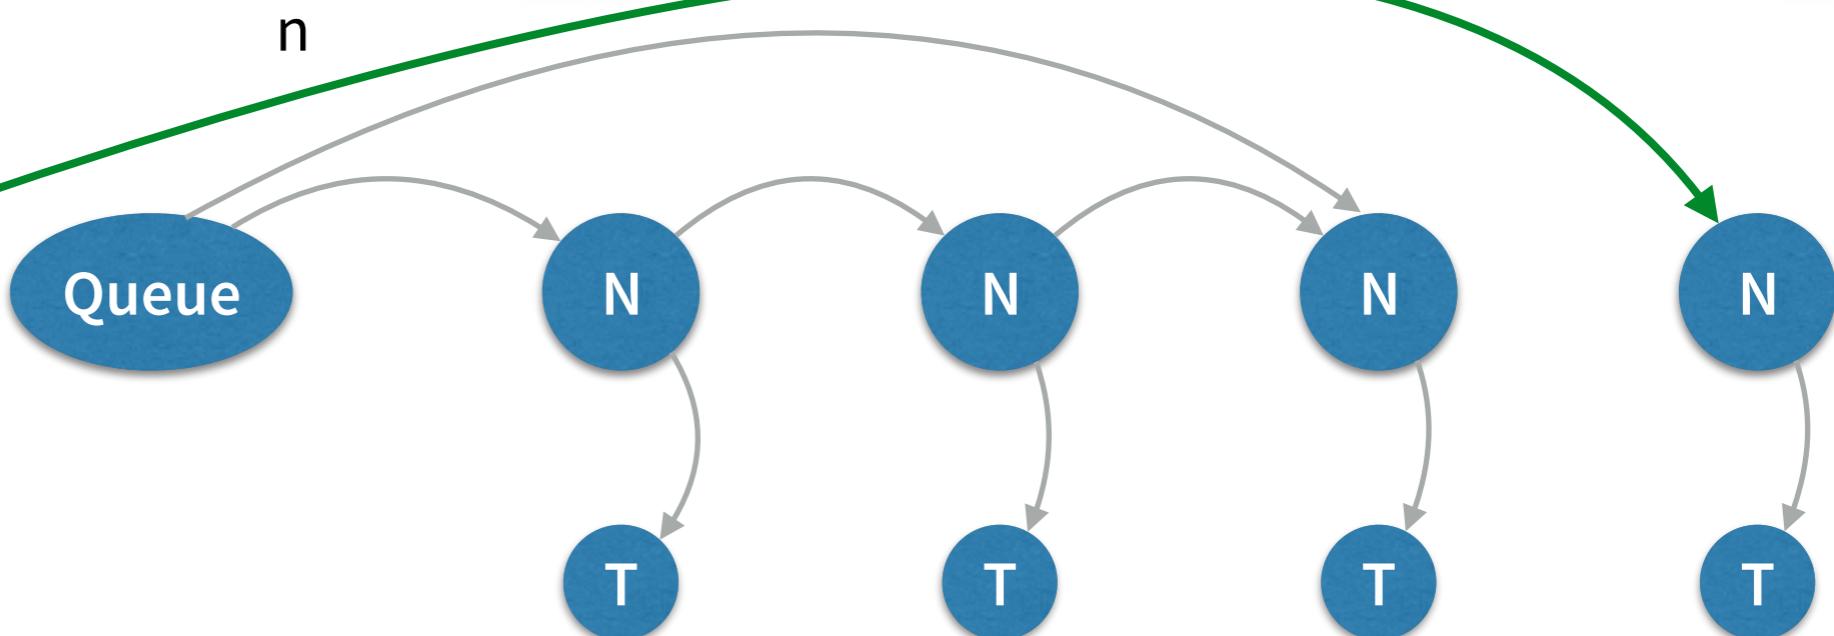
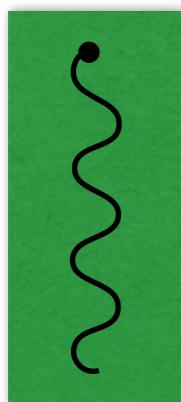
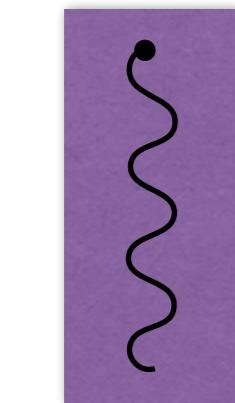


```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    }  
}
```

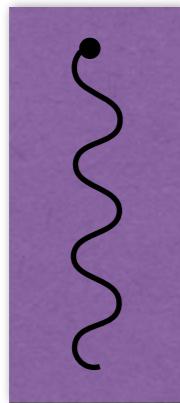


# Example: A Lock-free Queue [Michael & Scott]

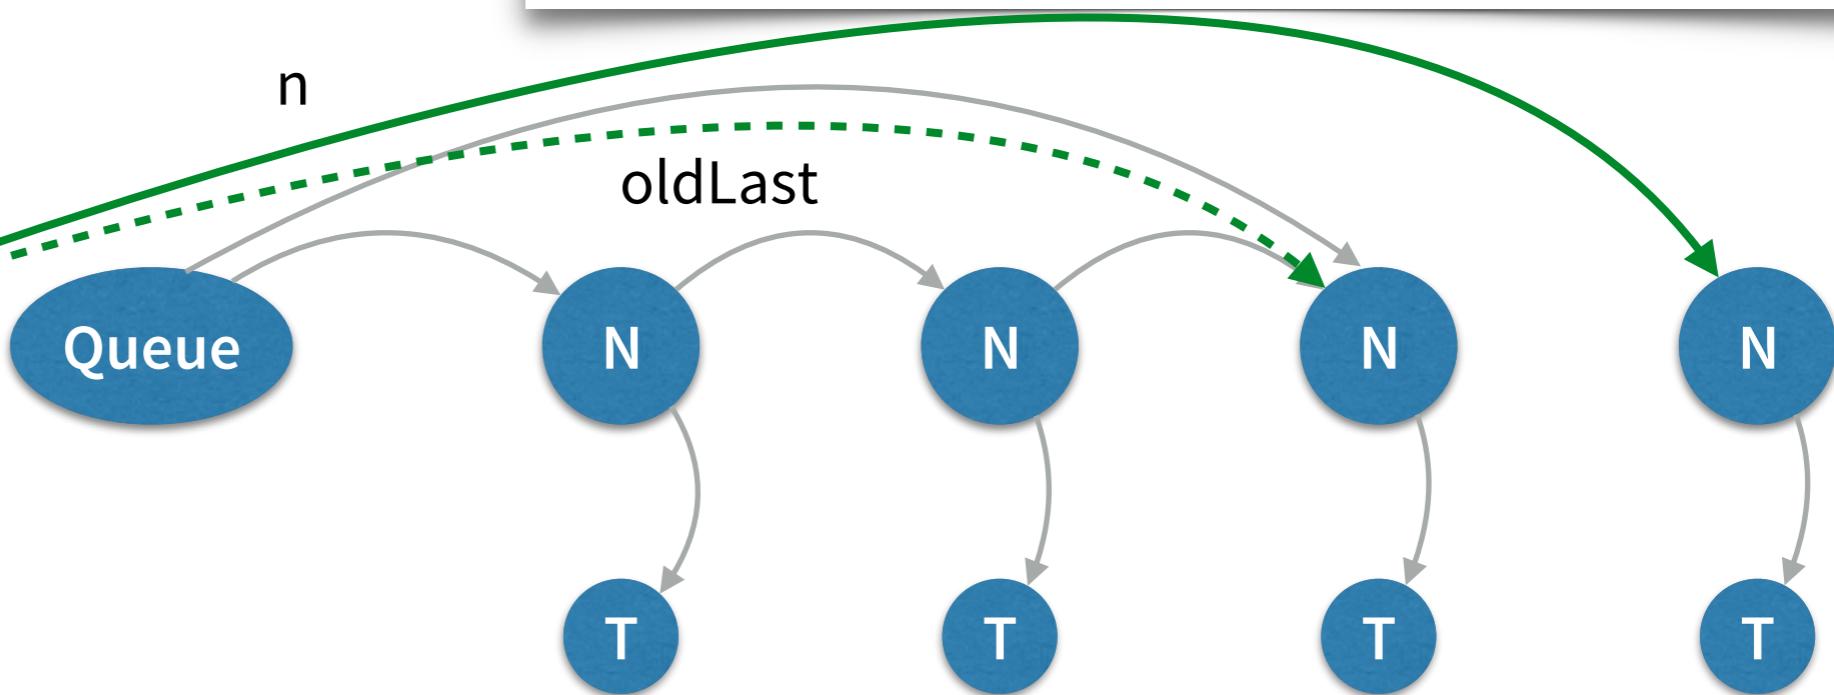
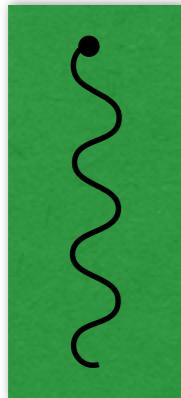
```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    }  
}
```



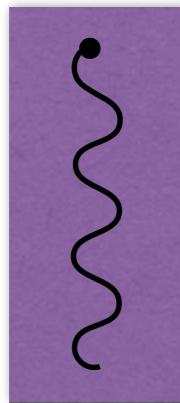
# Example: A Lock-free Queue [Michael & Scott]



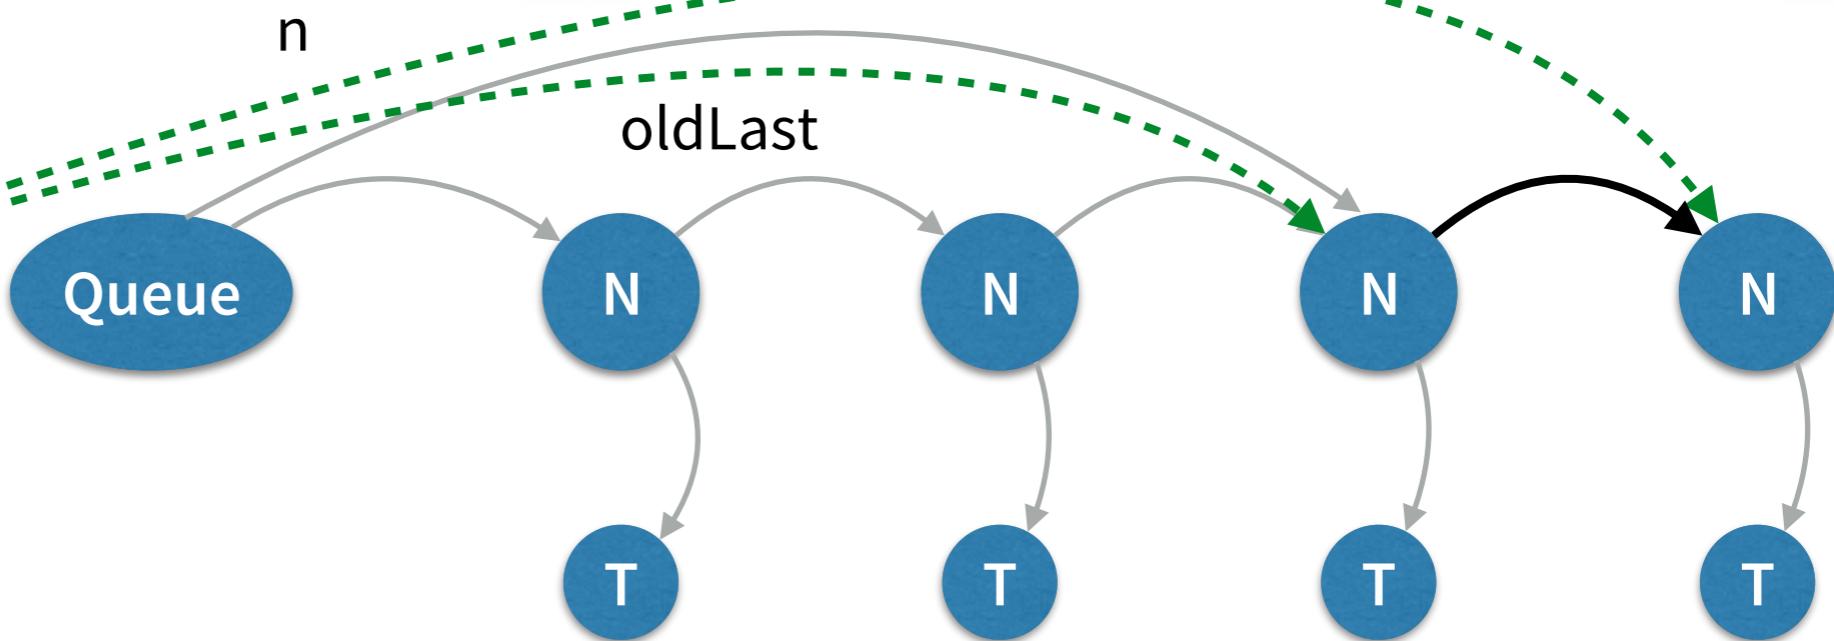
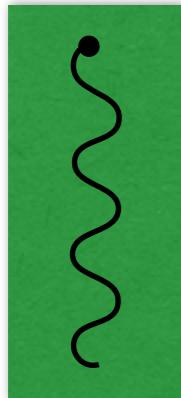
```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    }  
}
```



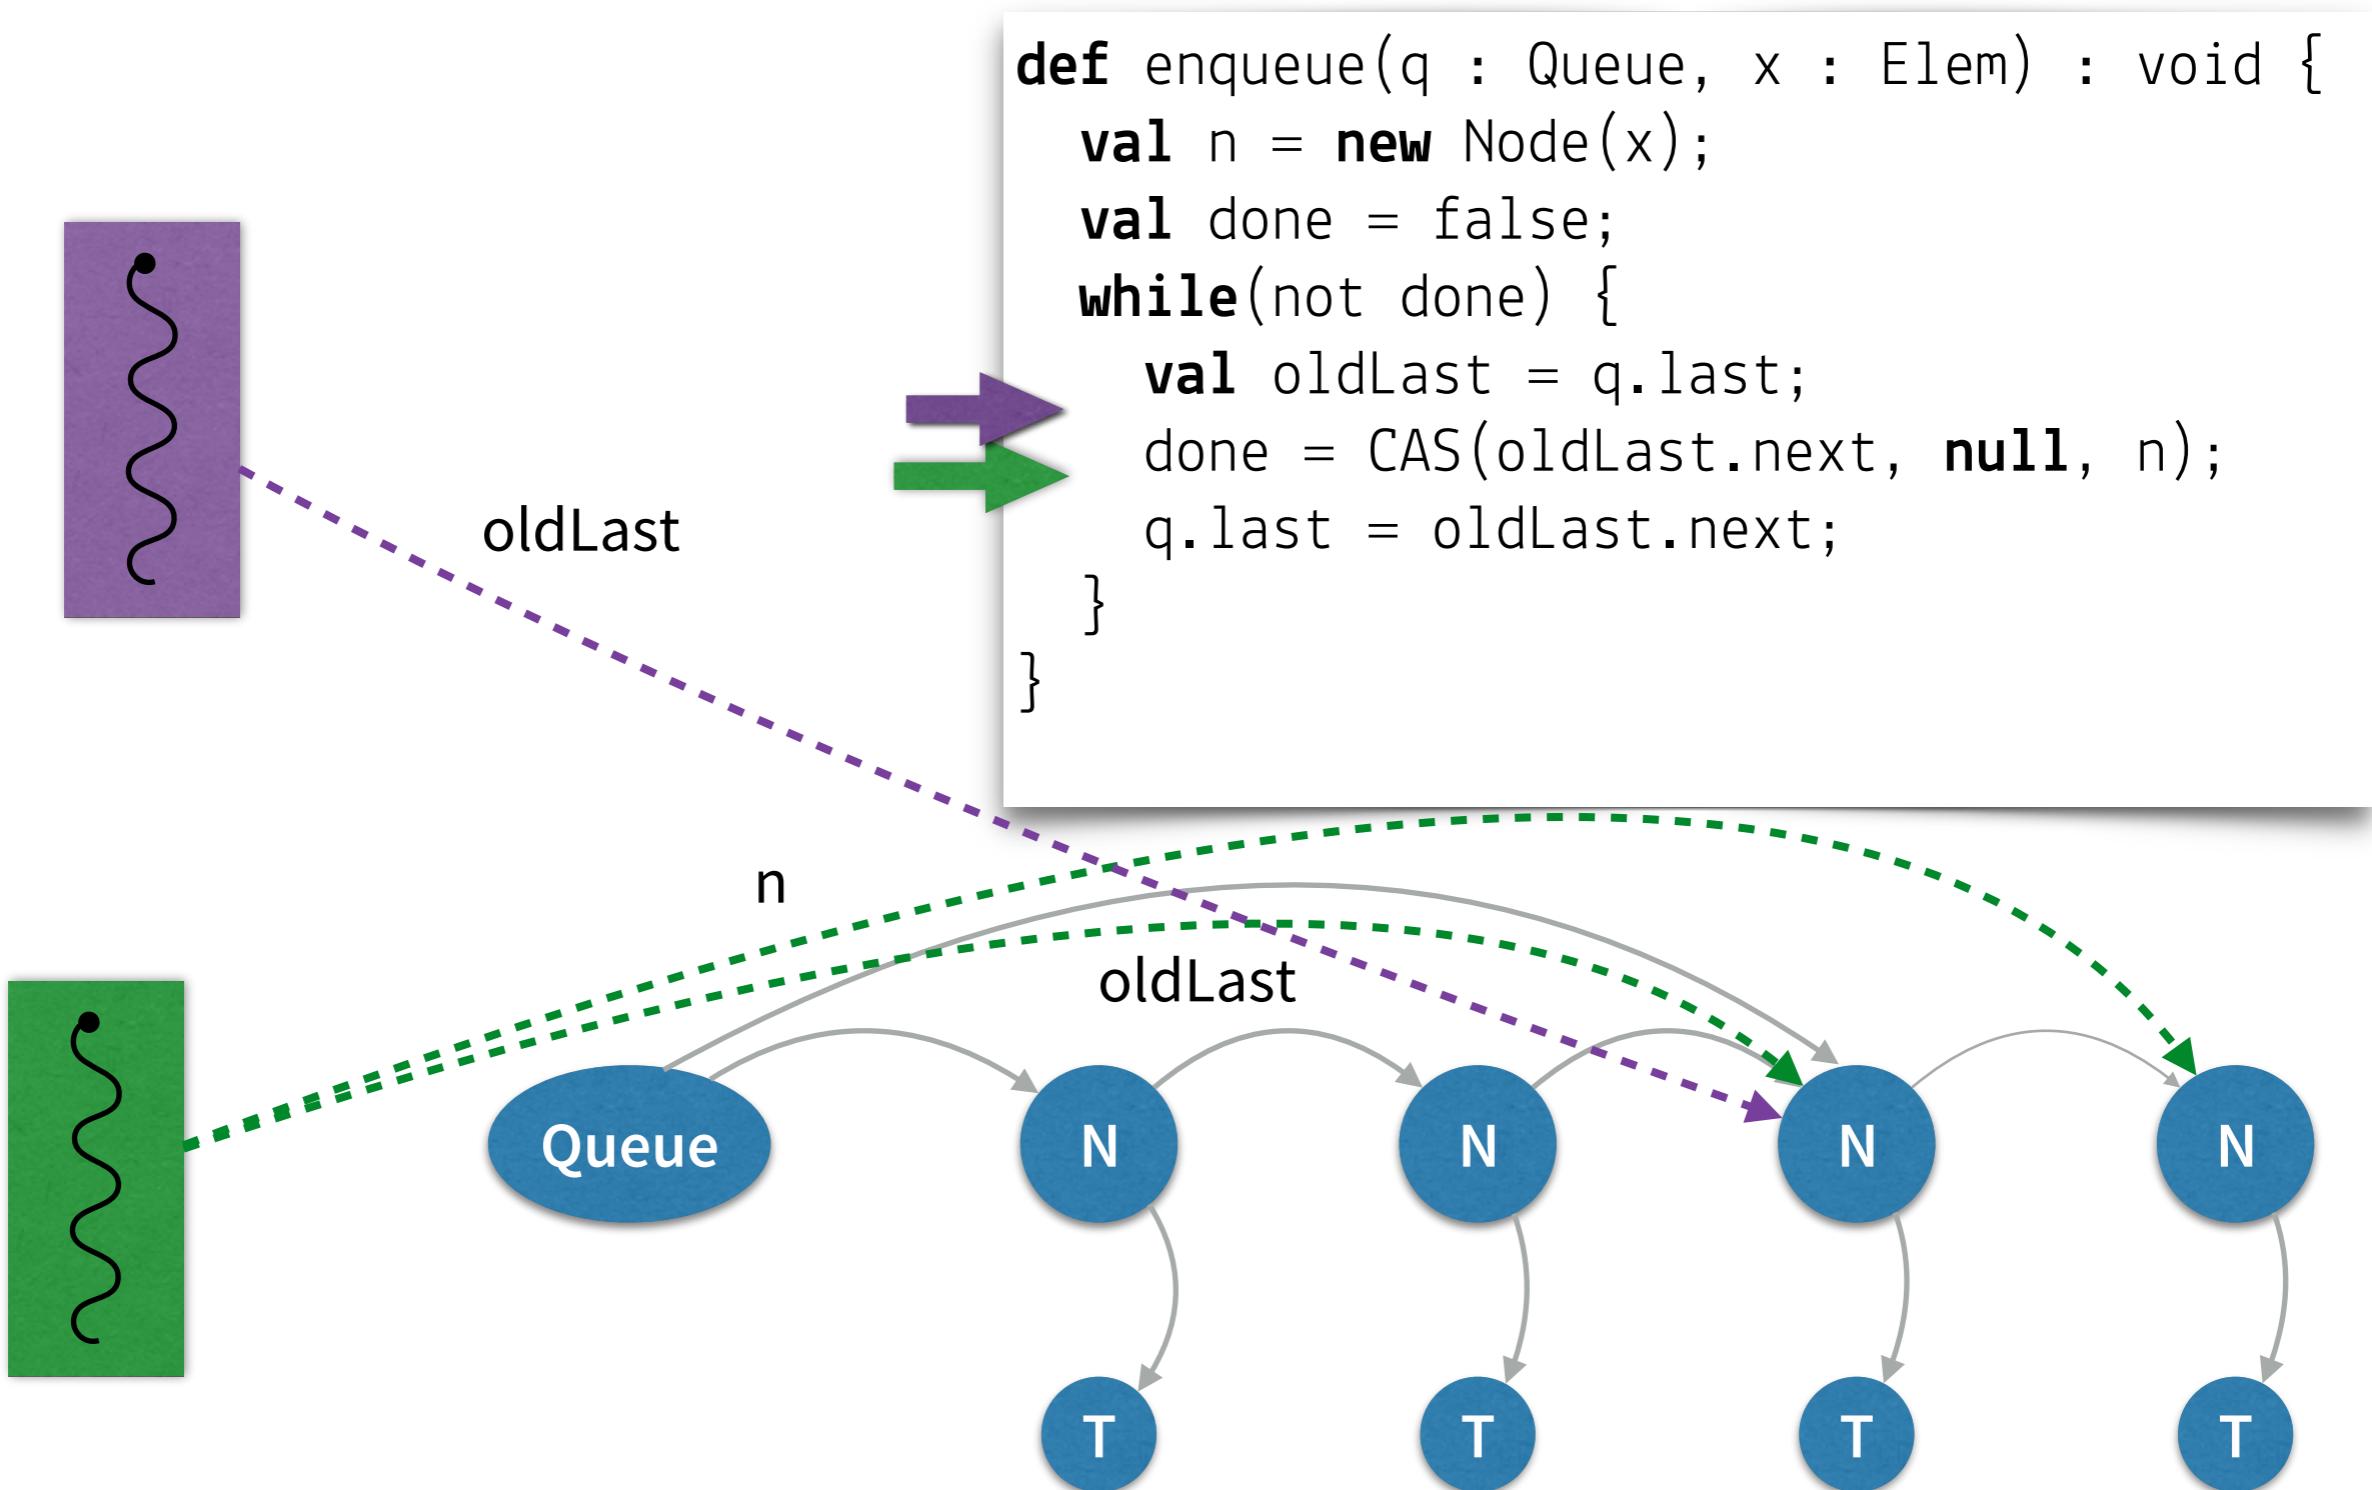
# Example: A Lock-free Queue [Michael & Scott]



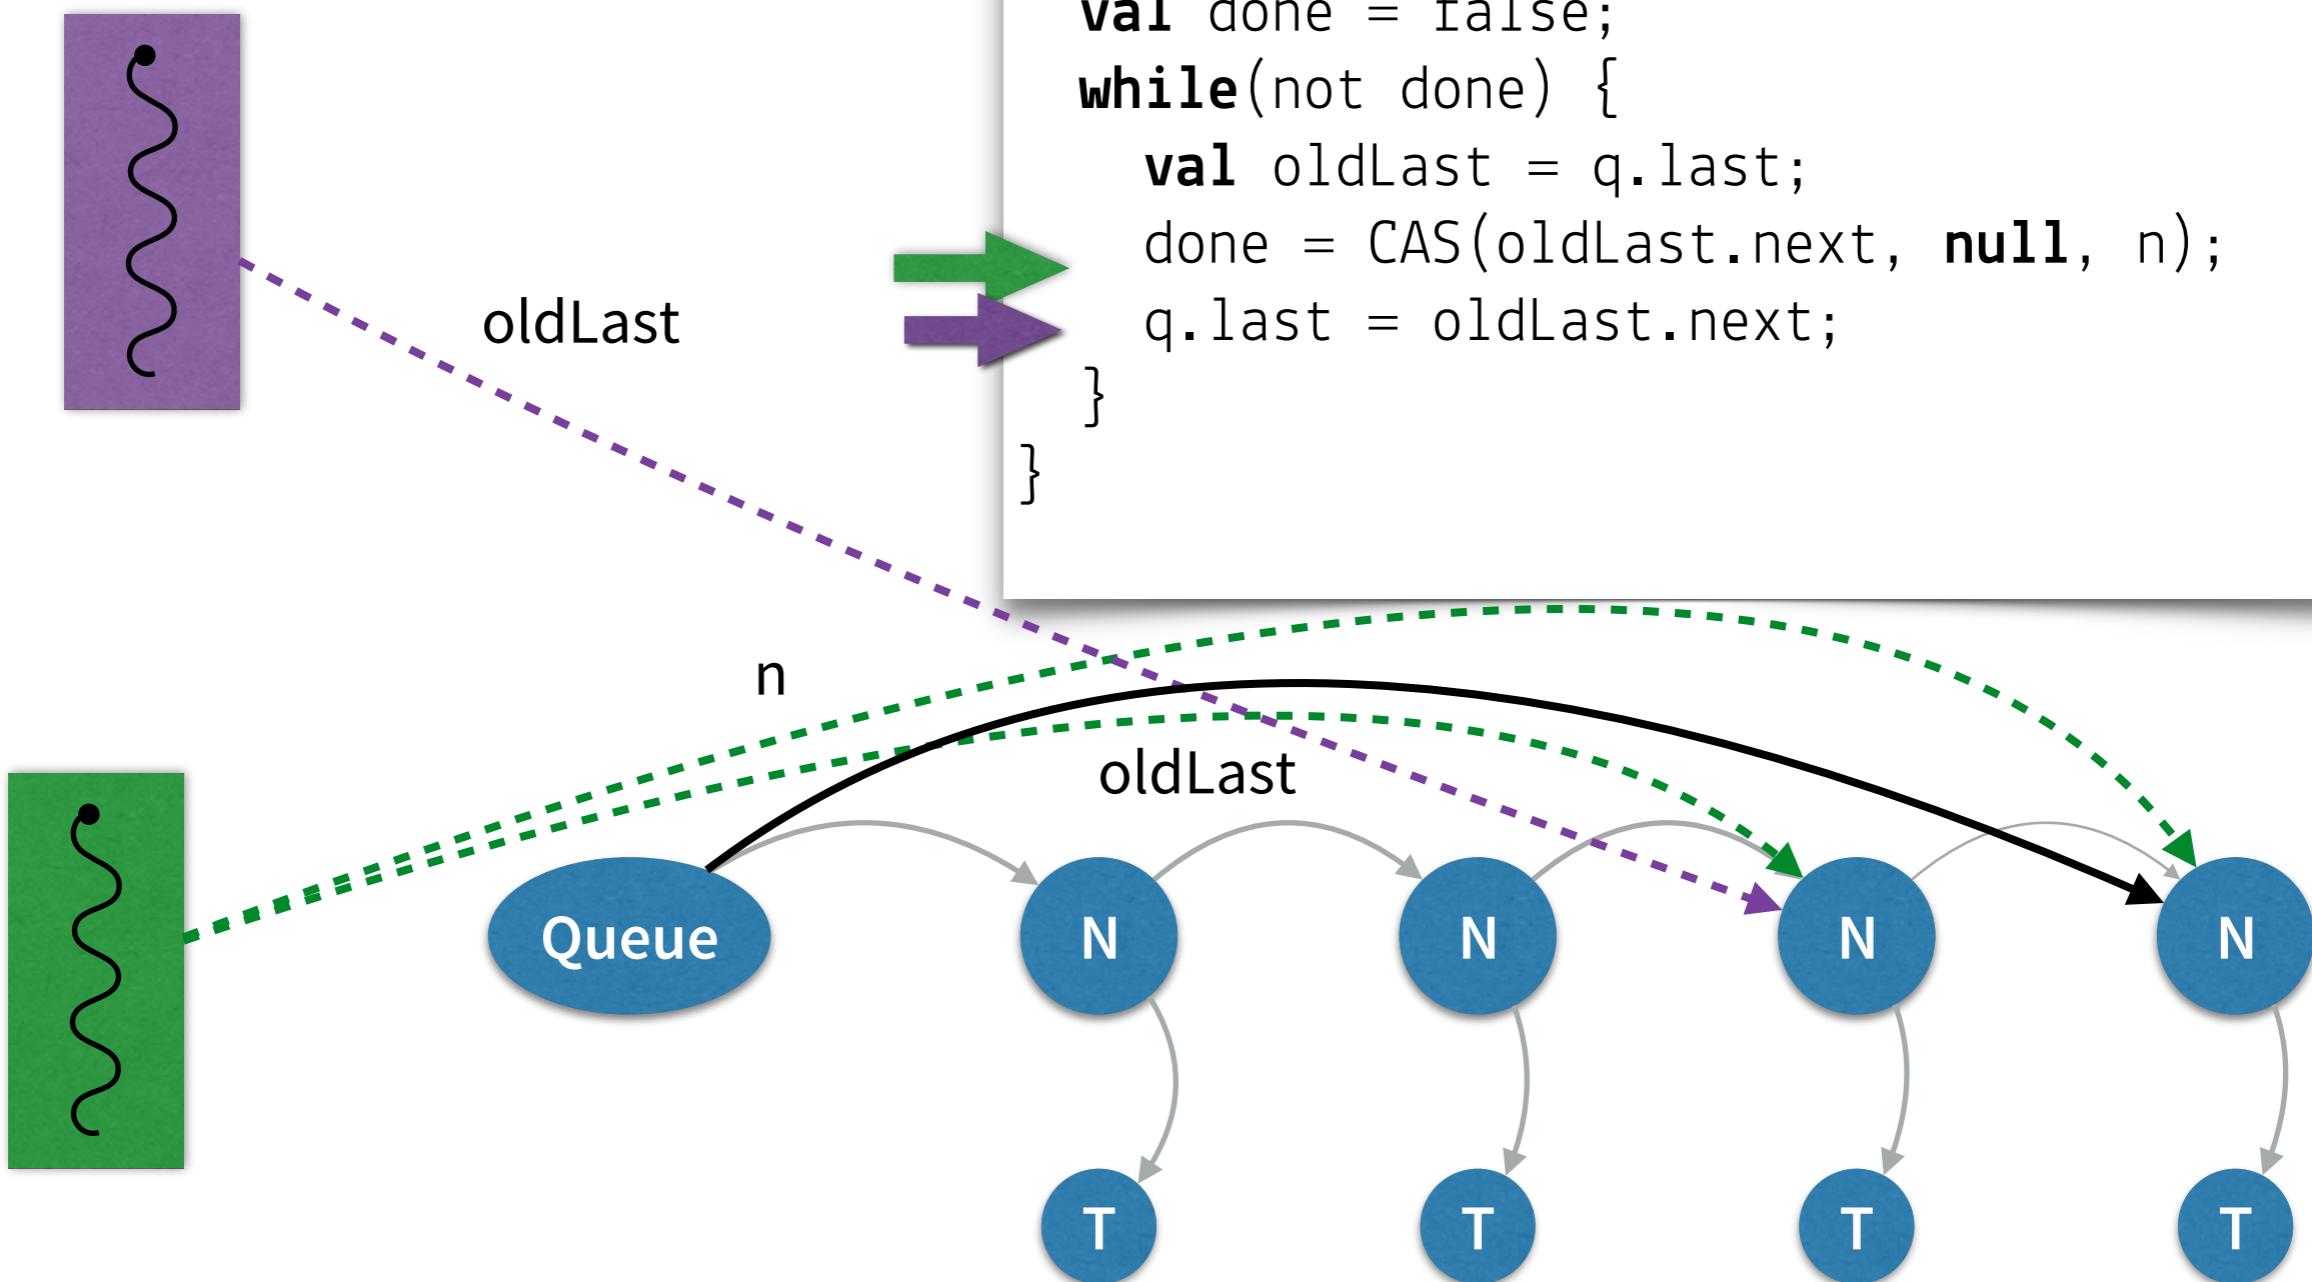
```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    }  
}
```



# Example: A Lock-free Queue [Michael & Scott]



# Example: A Lock-free Queue [Michael & Scott]



# Finding Patterns in Lock-free Programming

- Speculation

- Publication

- Acquisition

```
def dequeue(q : Queue) : T {  
    ...  
}
```

```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    }  
}
```

```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAS(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

```
def push(s : Stack, x : T) : void {  
    ...  
}
```

# Finding Patterns in Lock-free Programming

- Speculation

- Publication

- Acquisition

```
def dequeue(q : Queue) : T {  
    ...  
}
```

```
def enqueue(q : Queue, x : Elem) : void {  
    val n = new Node(x);  
    val done = false;  
    while(not done) {  
        val oldLast = q.last;  
        done = CAS(oldLast.next, null, n);  
        q.last = oldLast.next;  
    } }  
Field 'last' is subject to races and  
must be written to with a CAS.
```

```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAS(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

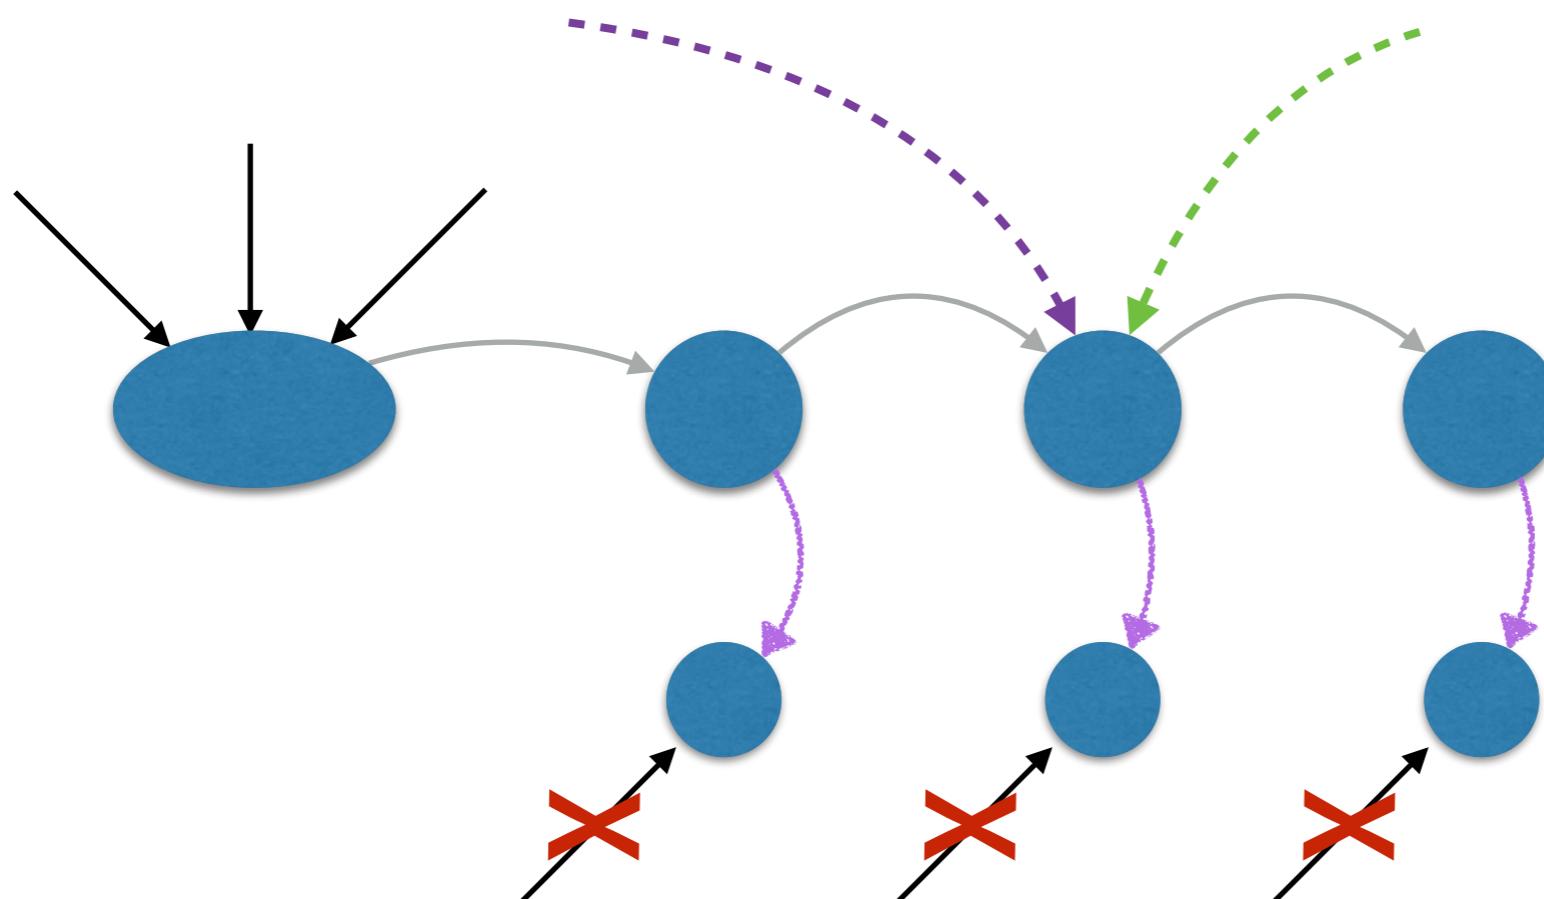
```
def push(s : Stack, x : T) : void {  
    ...  
}
```

Our idea:  
Give these concepts  
meaningful types!

# Why Types?

---

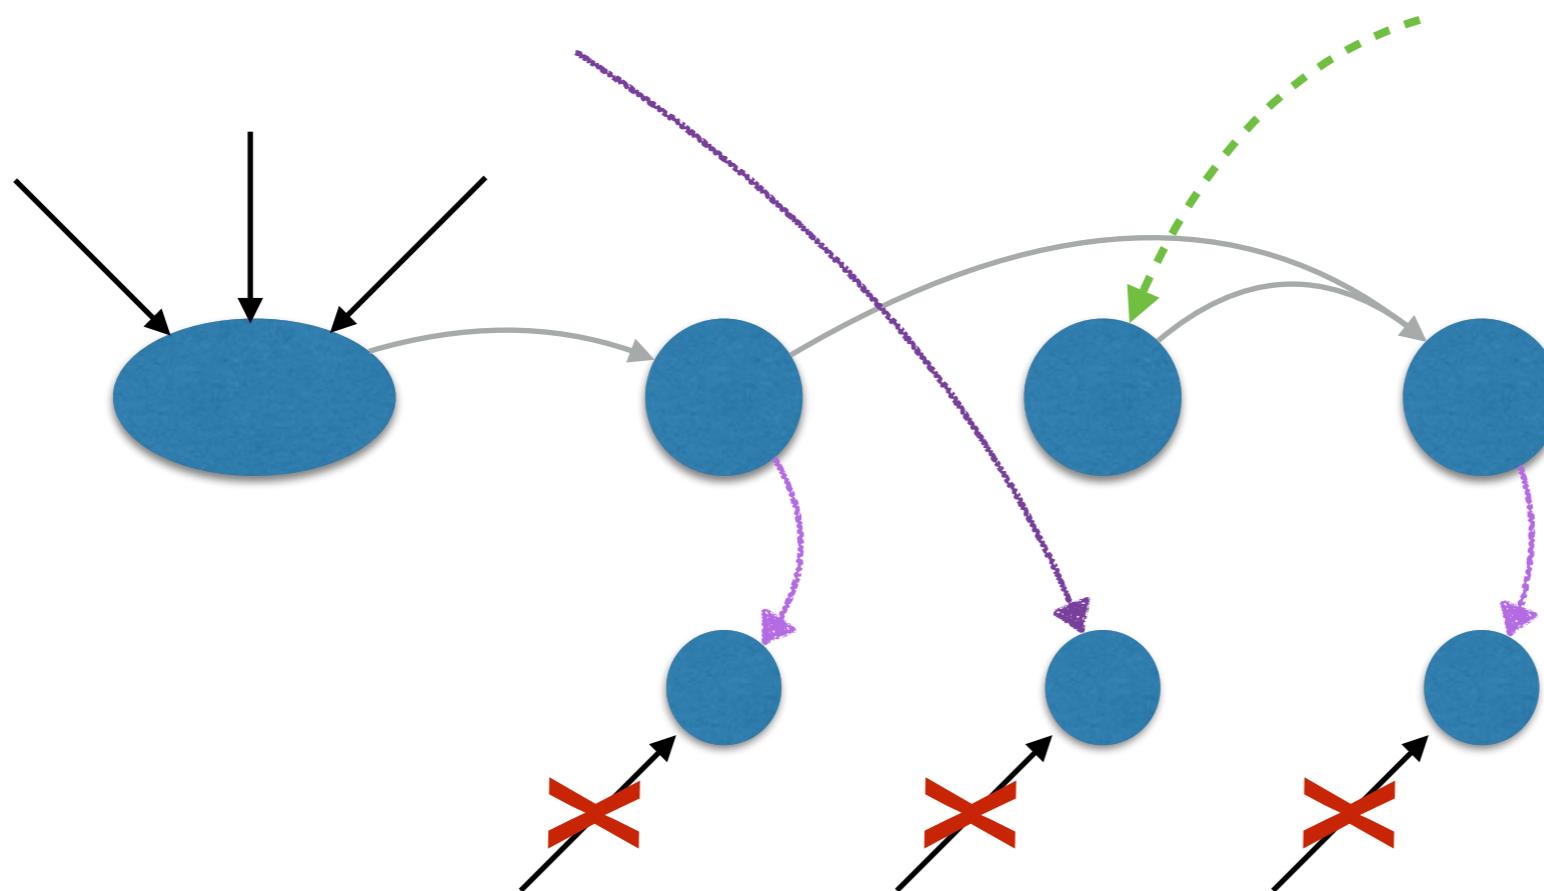
- Why not model checking, concurrent separation logic, etc.?
- Modularity: no need for inter-procedural analysis
- Language expressivity: allow concurrent data structures storing linear references



# Why Types?

---

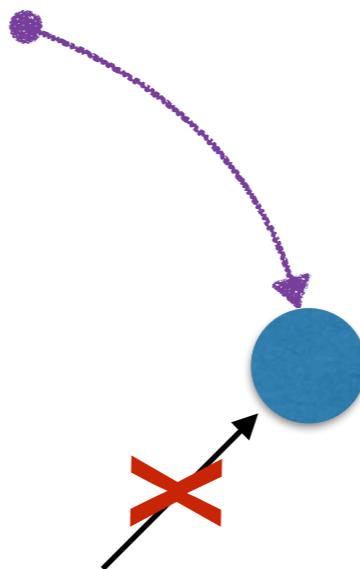
- Why not model checking, concurrent separation logic, etc.?
- Modularity: no need for inter-procedural analysis
- Language expressivity: allow concurrent data structures storing linear references



# Linear References for Ownership Transfer

---

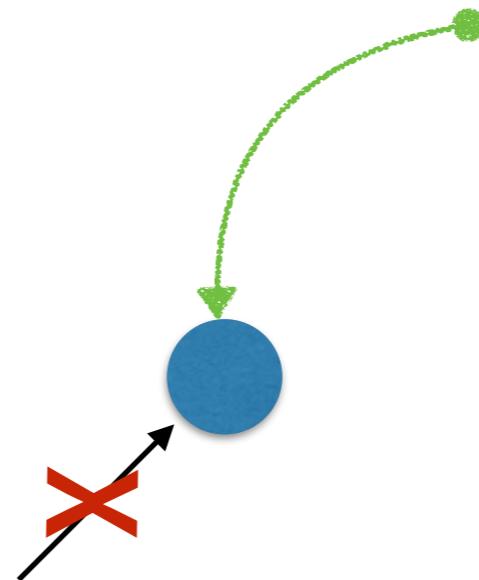
- A linear reference is the only reference to an object
- Prohibits shared access to objects
- Makes ownership explicit



# Linear References for Ownership Transfer

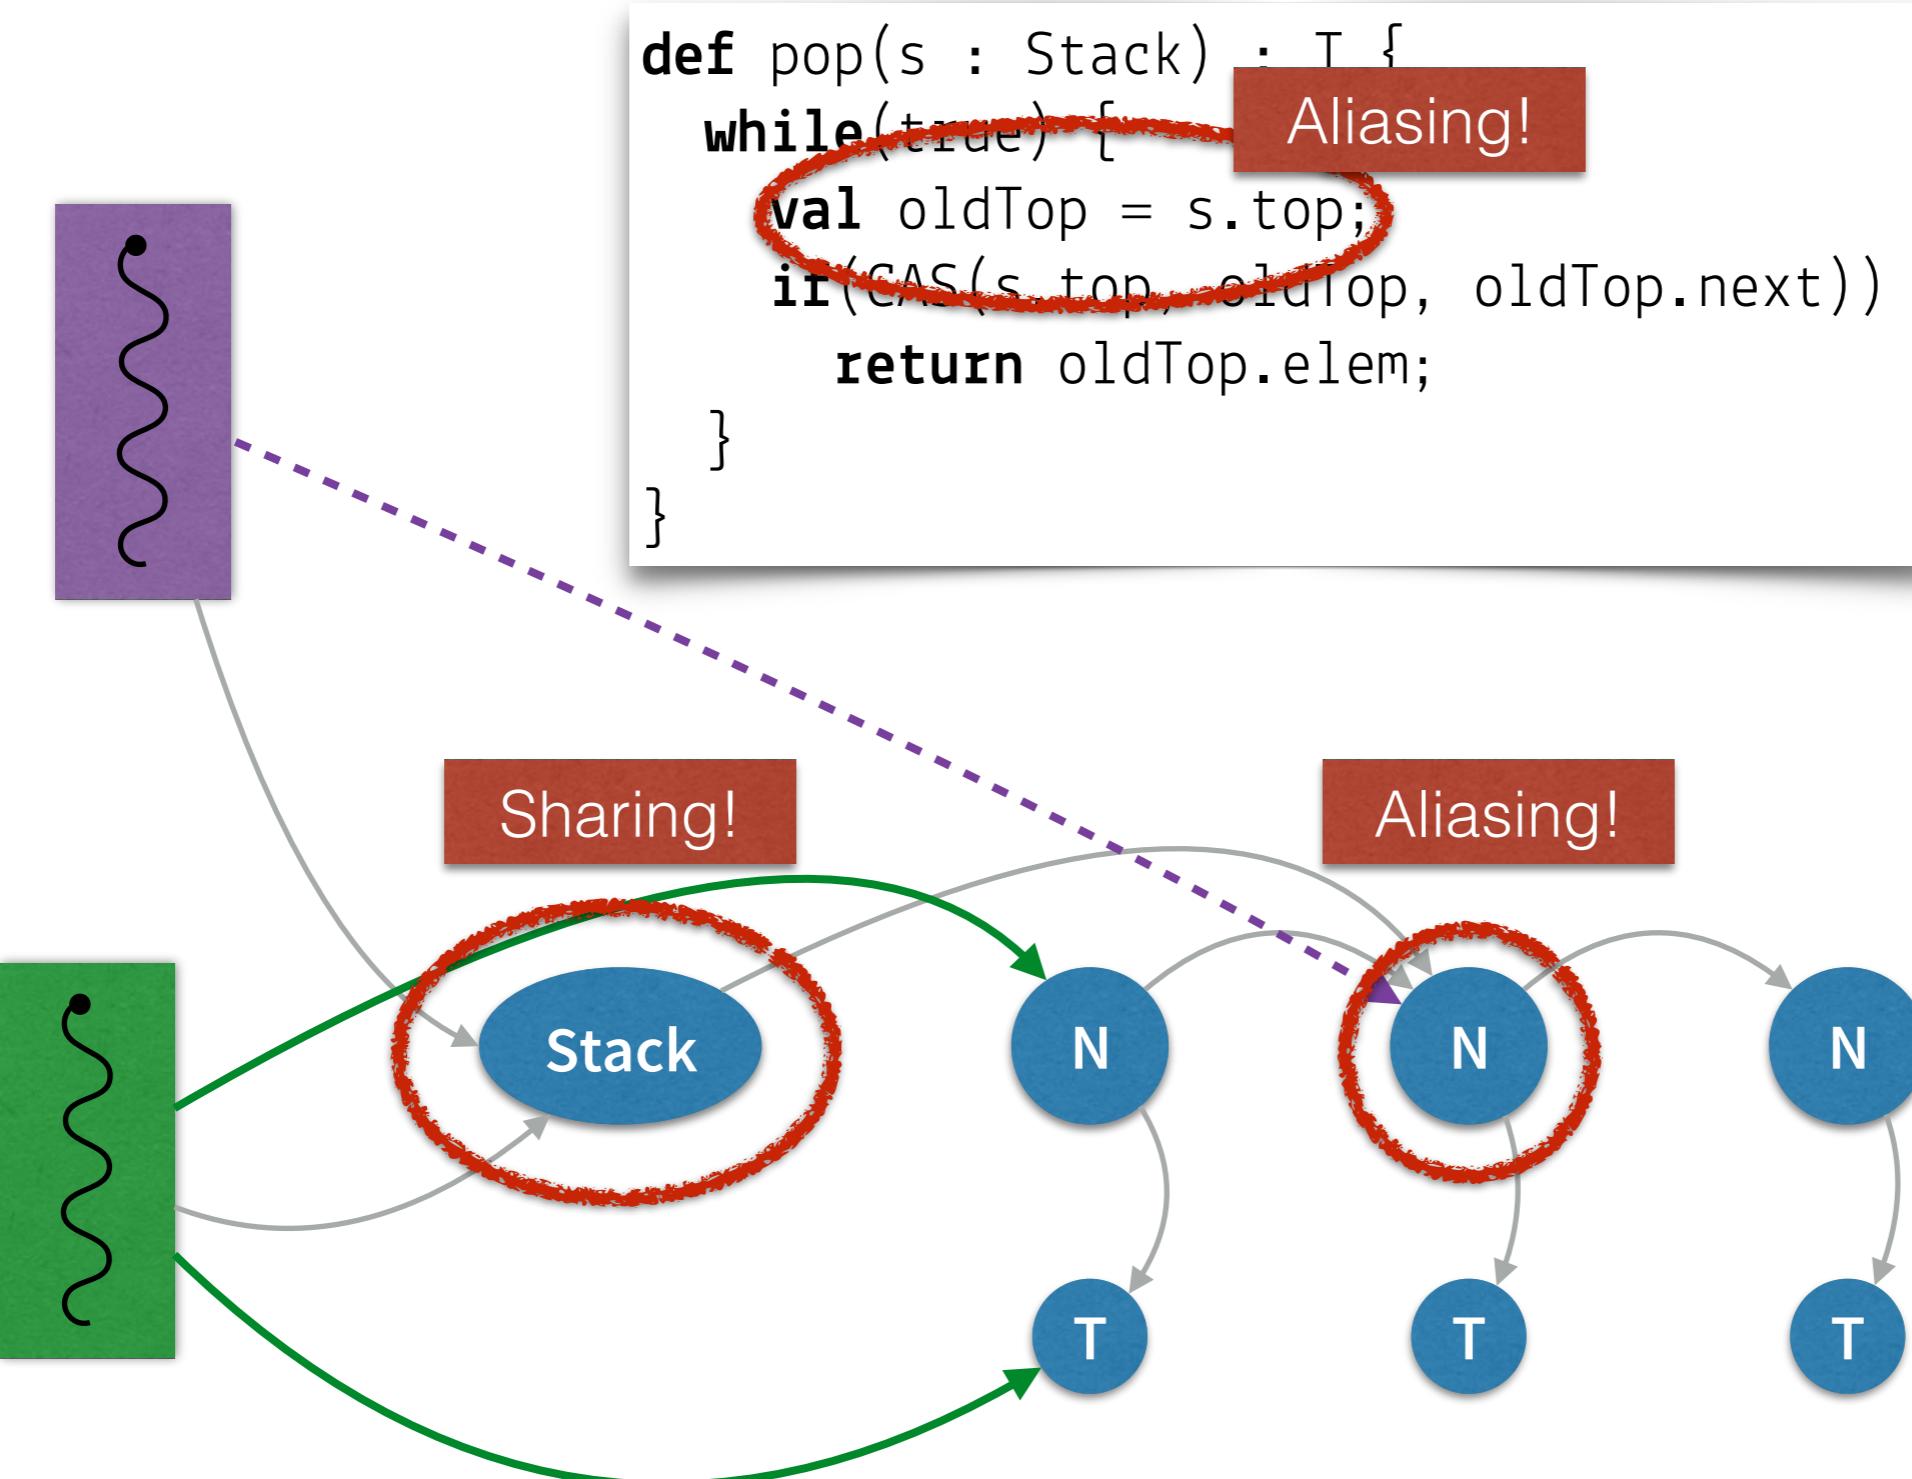
---

- A linear reference is the only reference to an object
- Prohibits shared access to objects
- Makes ownership explicit



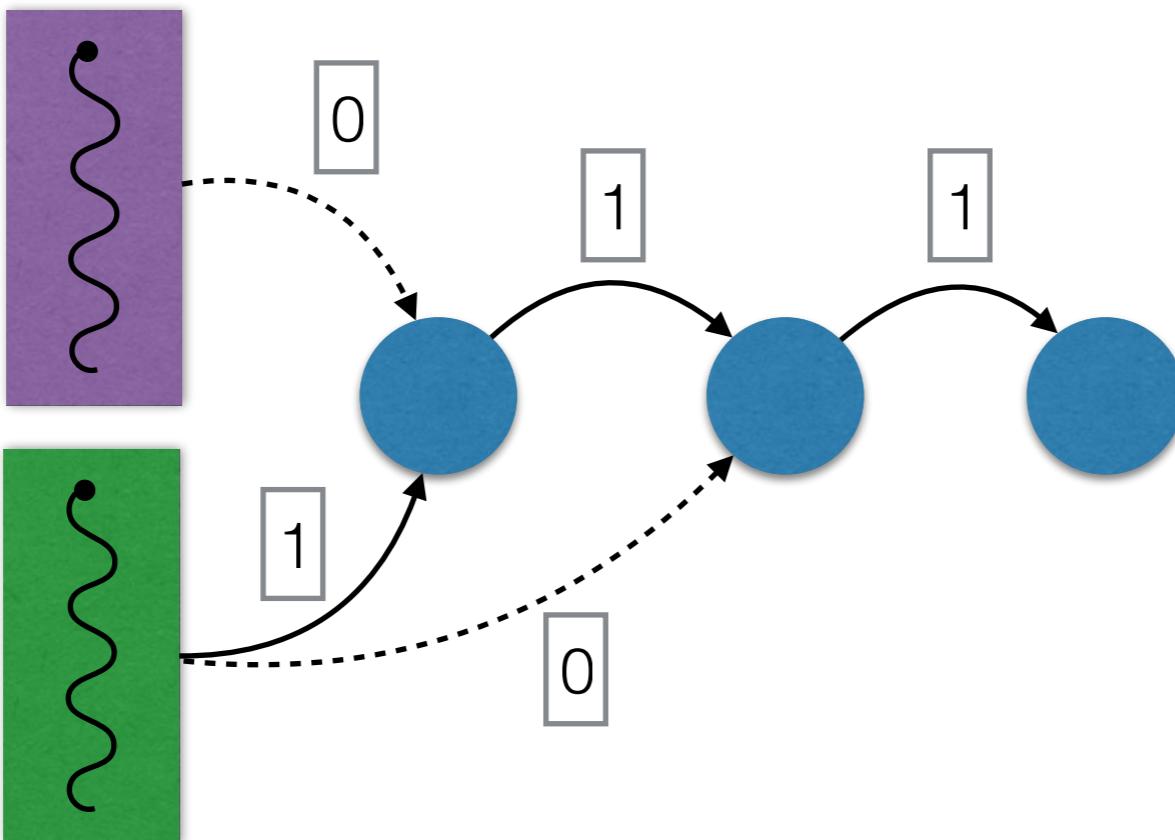
- A correct lock-free data-structure can guarantee exclusive ownership of an element
  - Can we use linear references to capture this?

# Linear References and Lock-free Programming



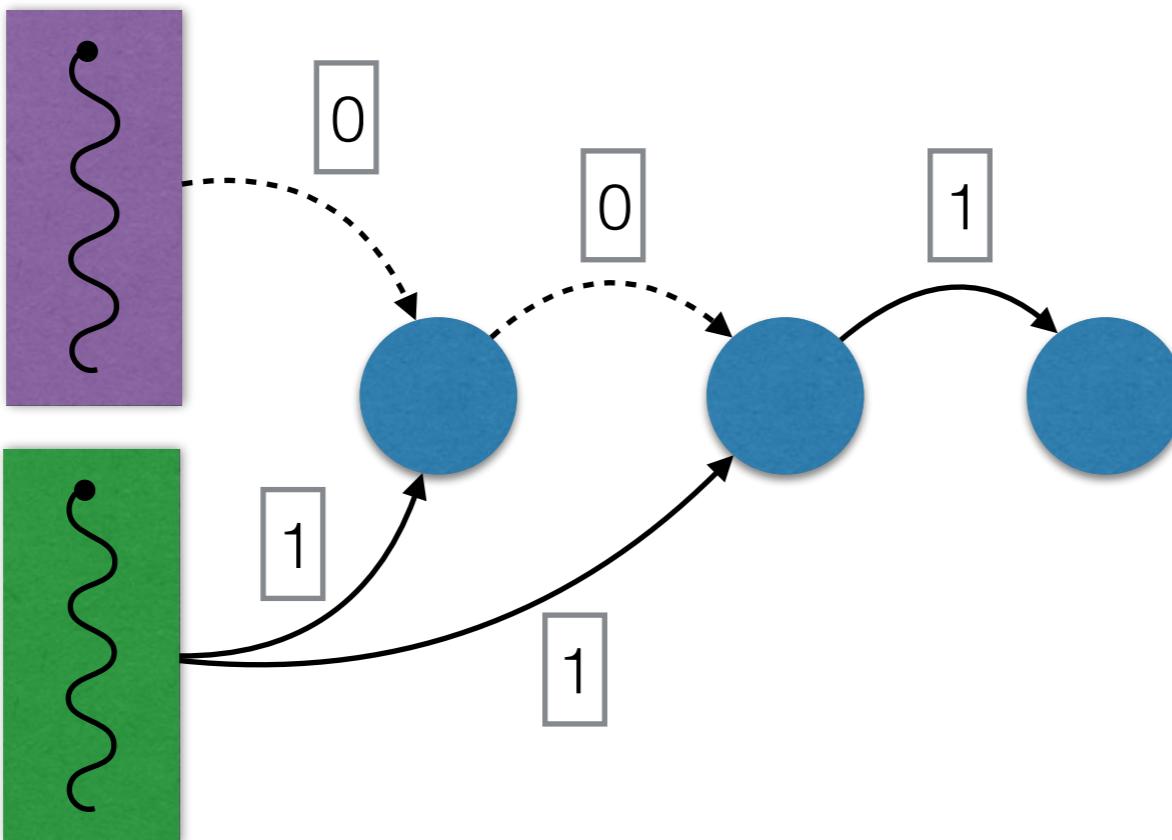
# Relaxed Linear References

- Separate aliasing from ownership
  - Unbounded aliasing, linear ownership
  - Ownership can be *atomically* transferred between aliases



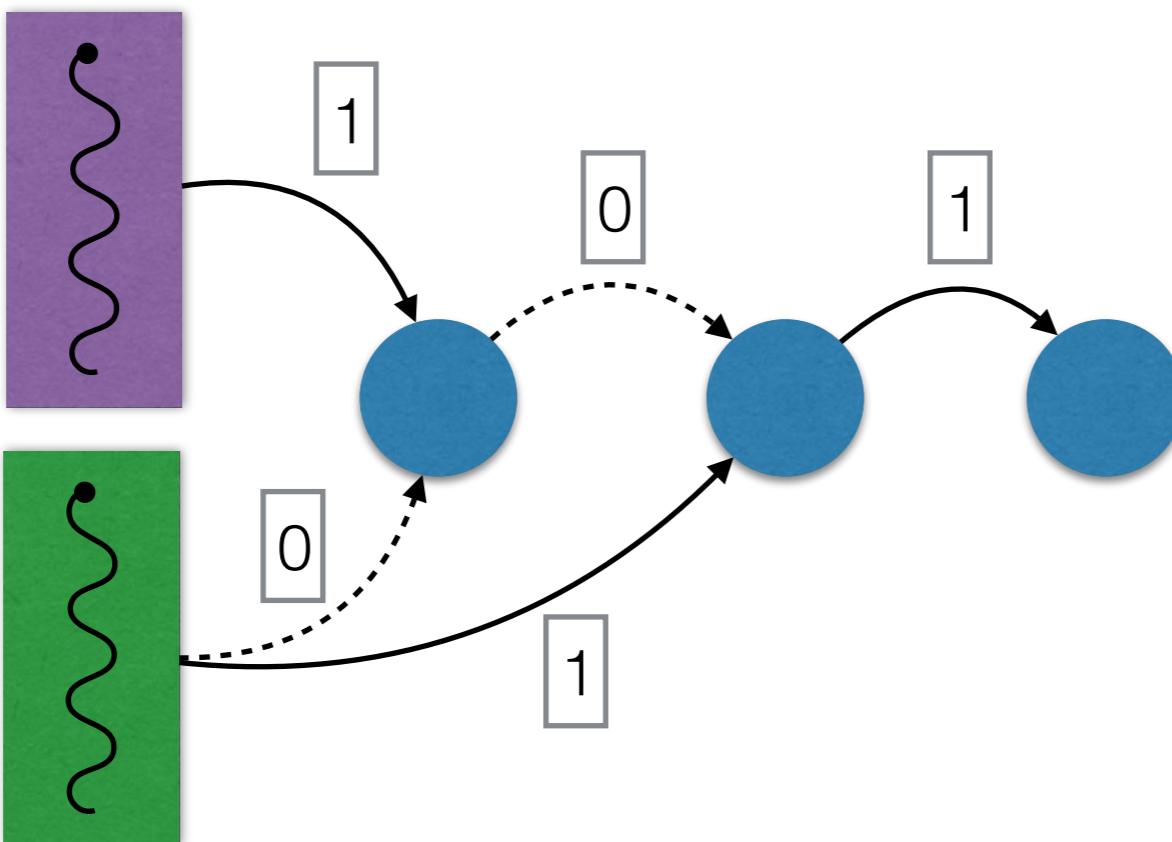
# Relaxed Linear References

- Separate aliasing from ownership
  - Unbounded aliasing, linear ownership
  - Ownership can be *atomically* transferred between aliases



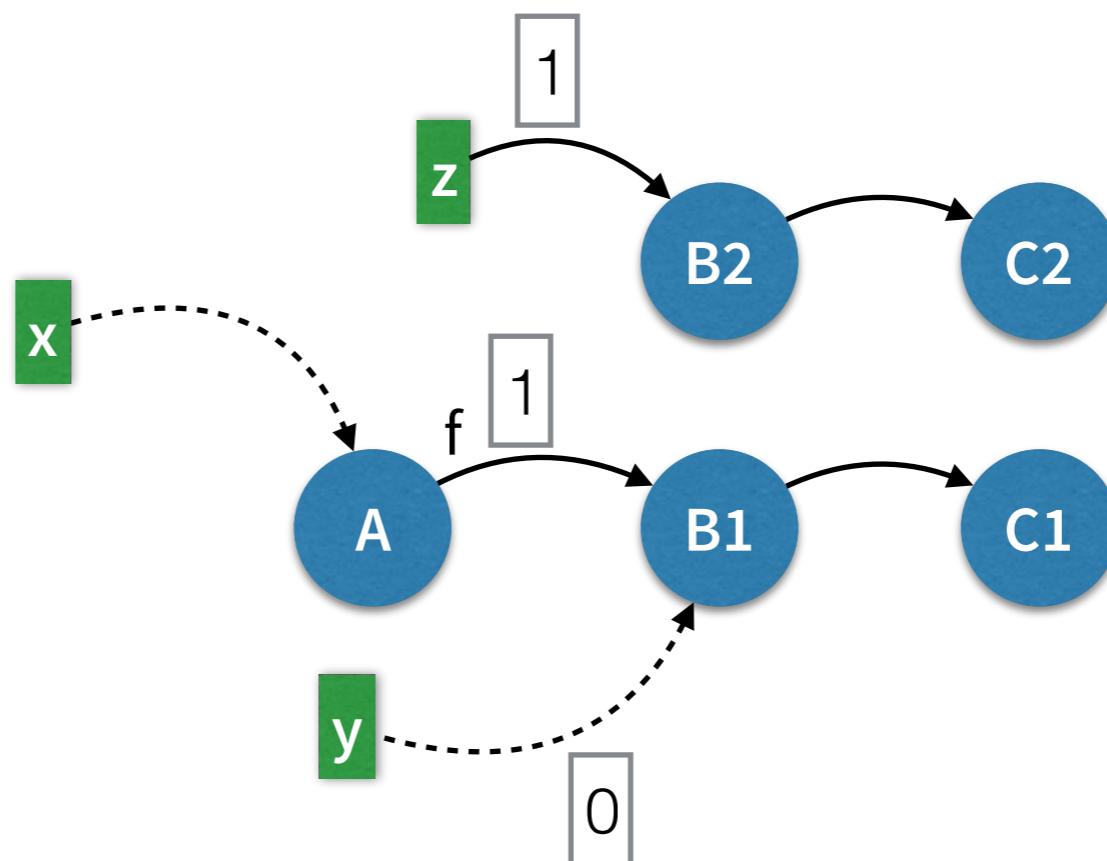
# Relaxed Linear References

- Separate aliasing from ownership
  - Unbounded aliasing, linear ownership
  - Ownership can be *atomically* transferred between aliases



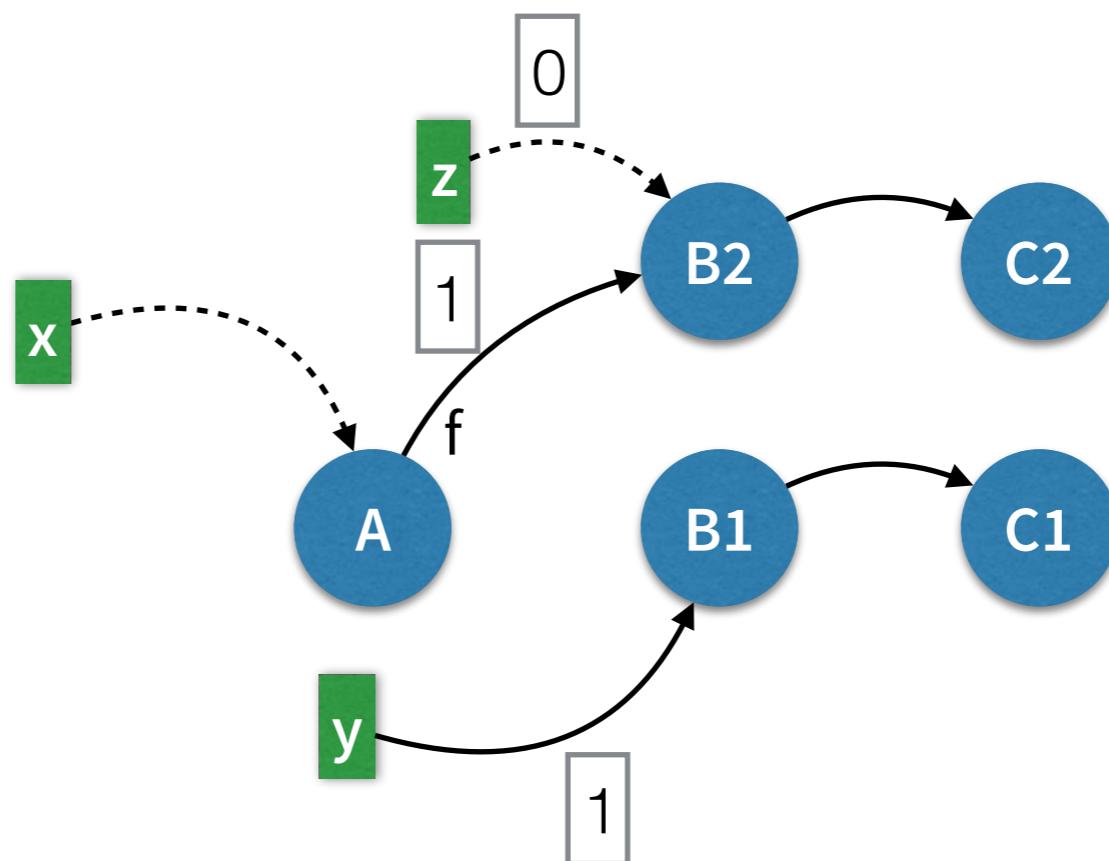
# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) – "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "



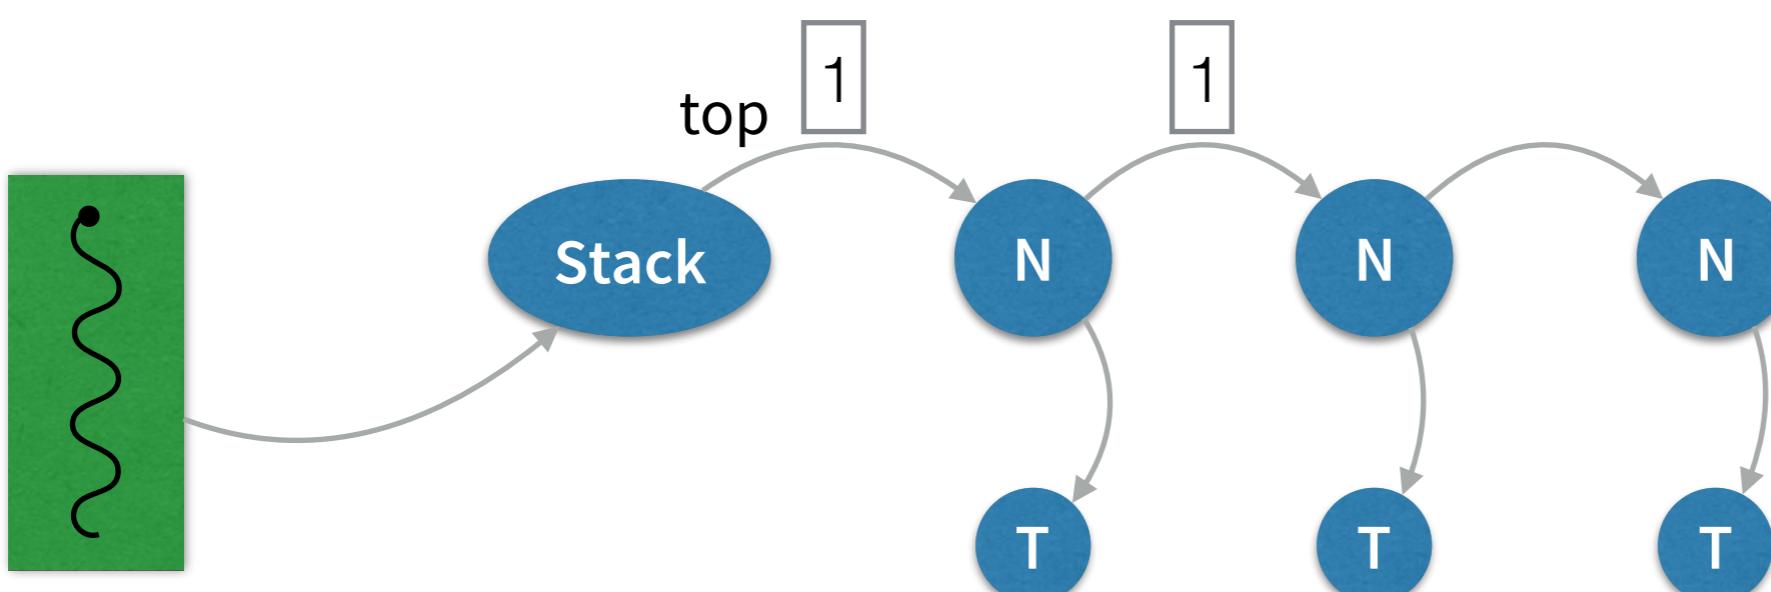
# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) – "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "



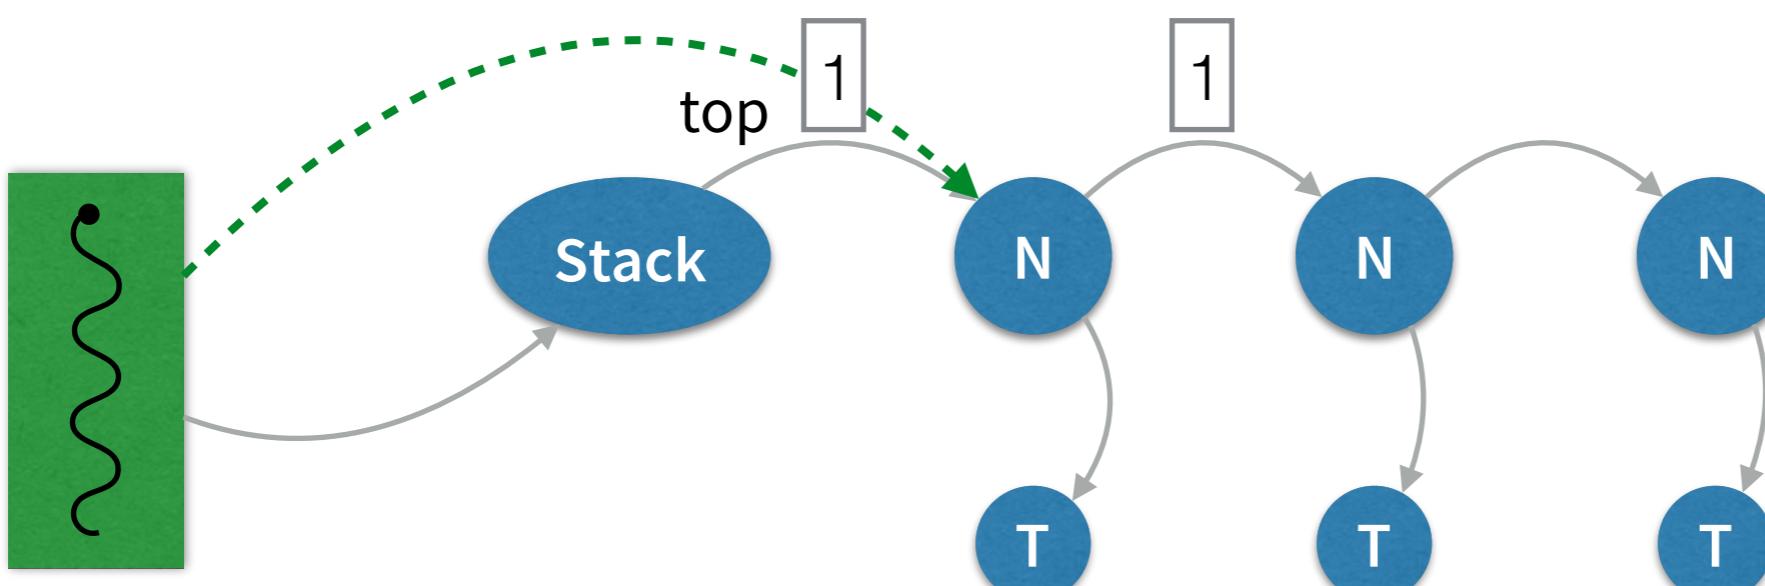
# Ownership Transfer in Action

```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```



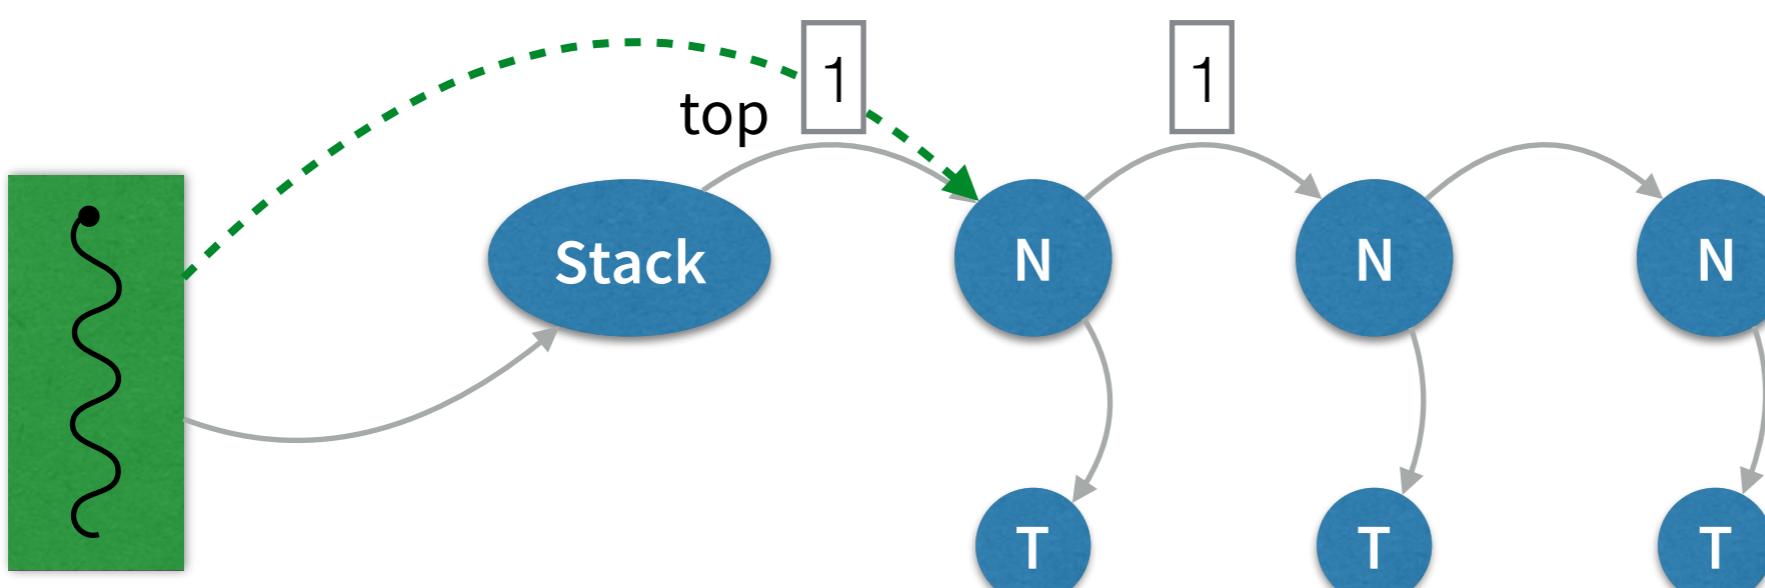
# Ownership Transfer in Action

```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

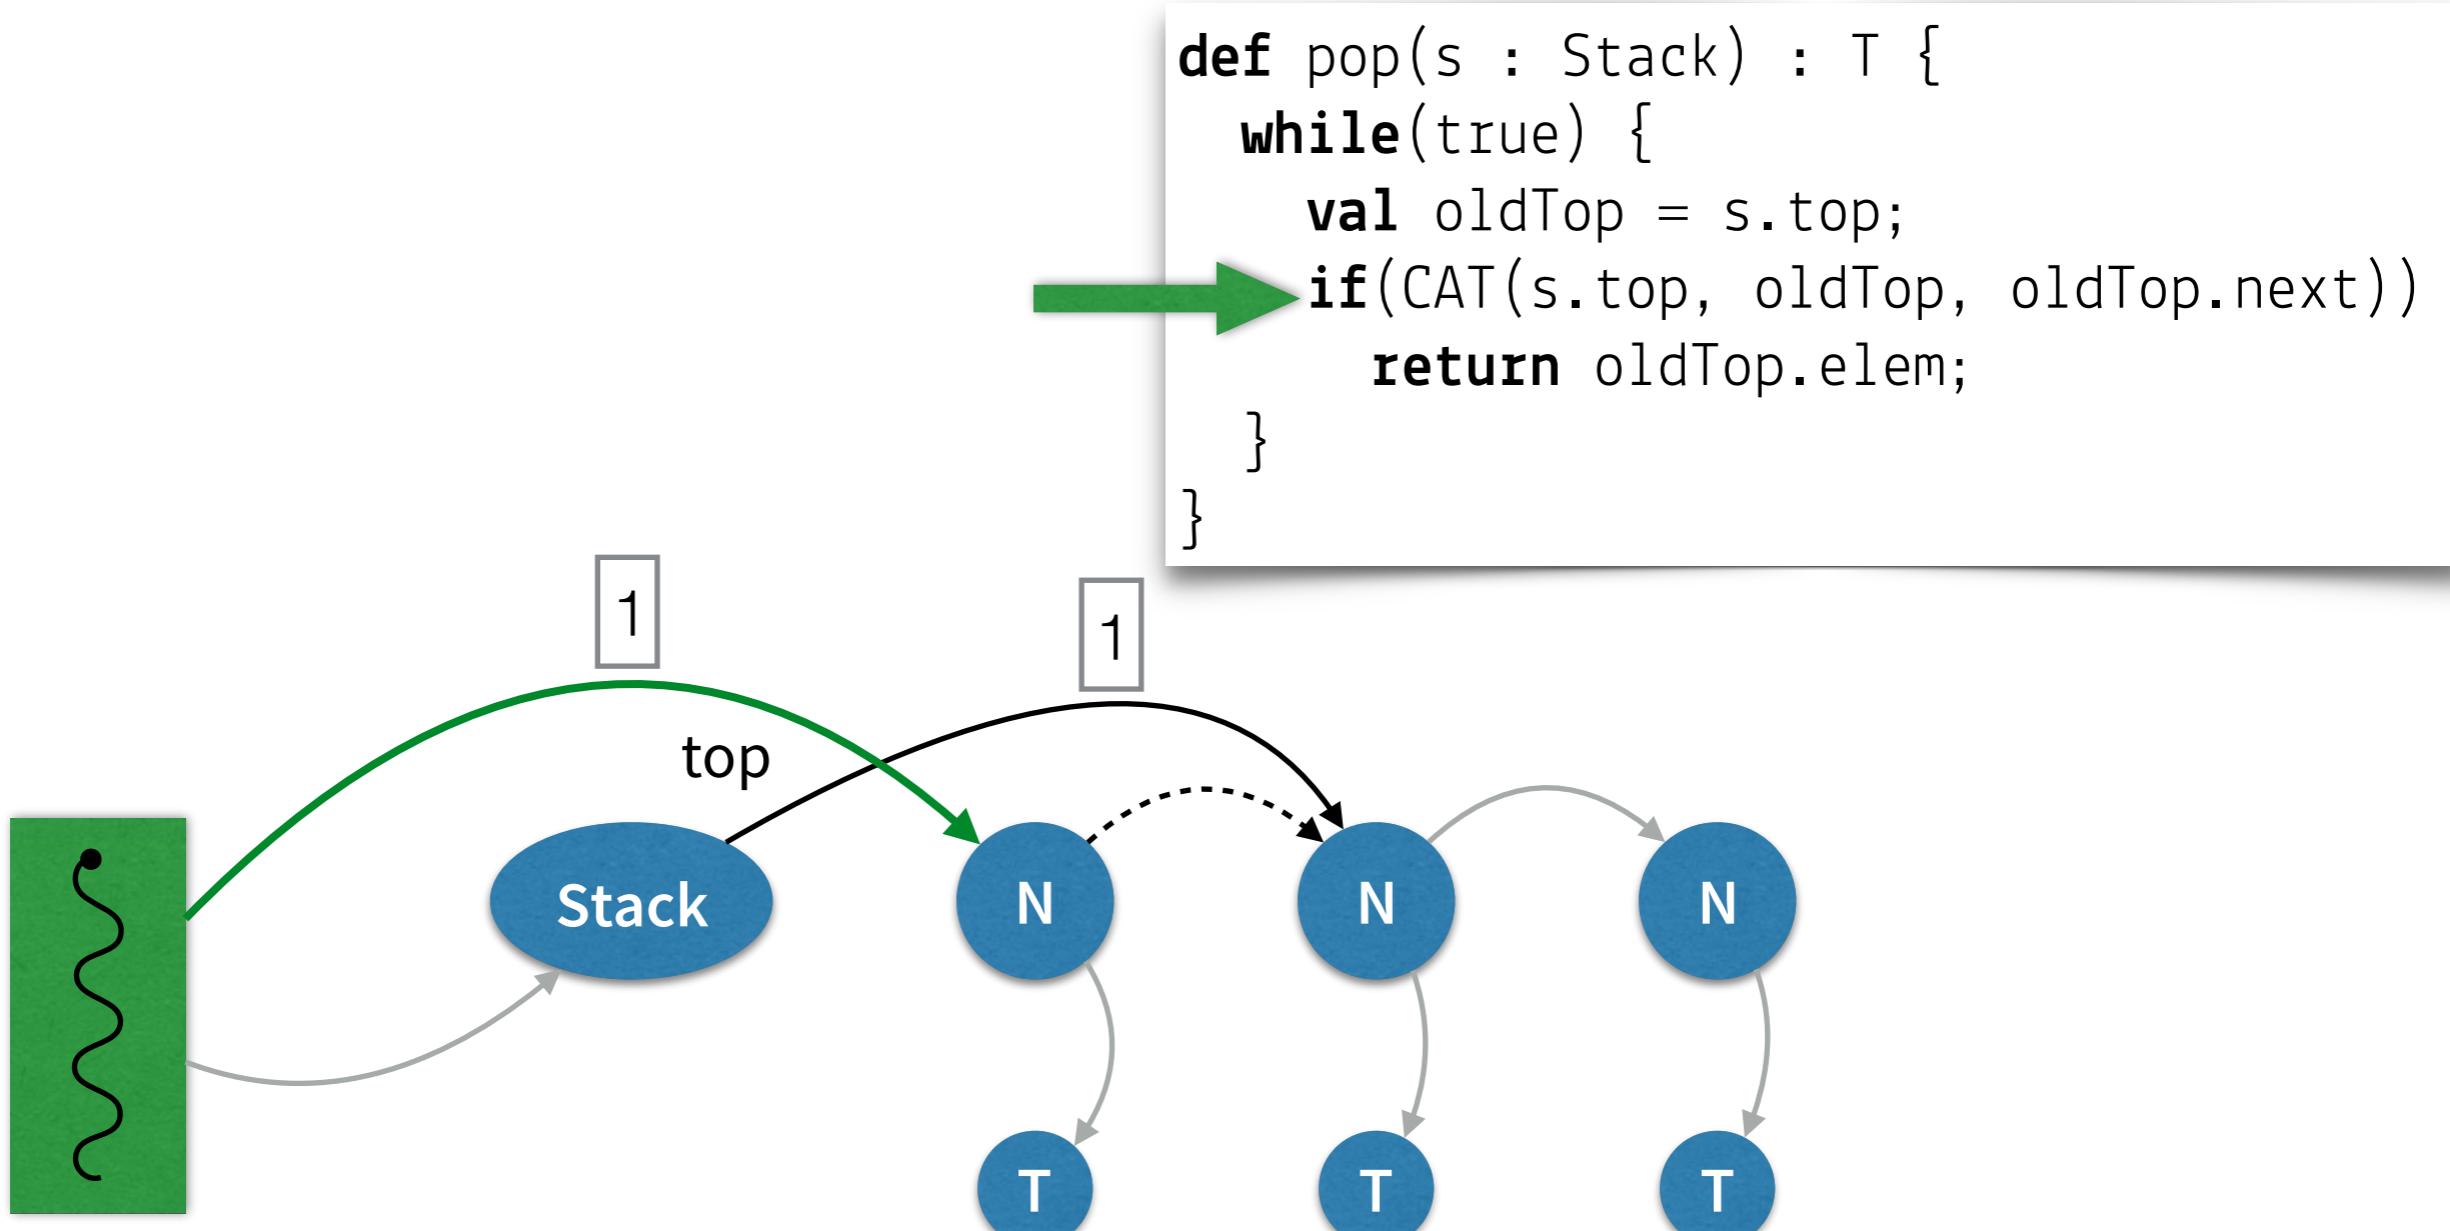


# Ownership Transfer in Action

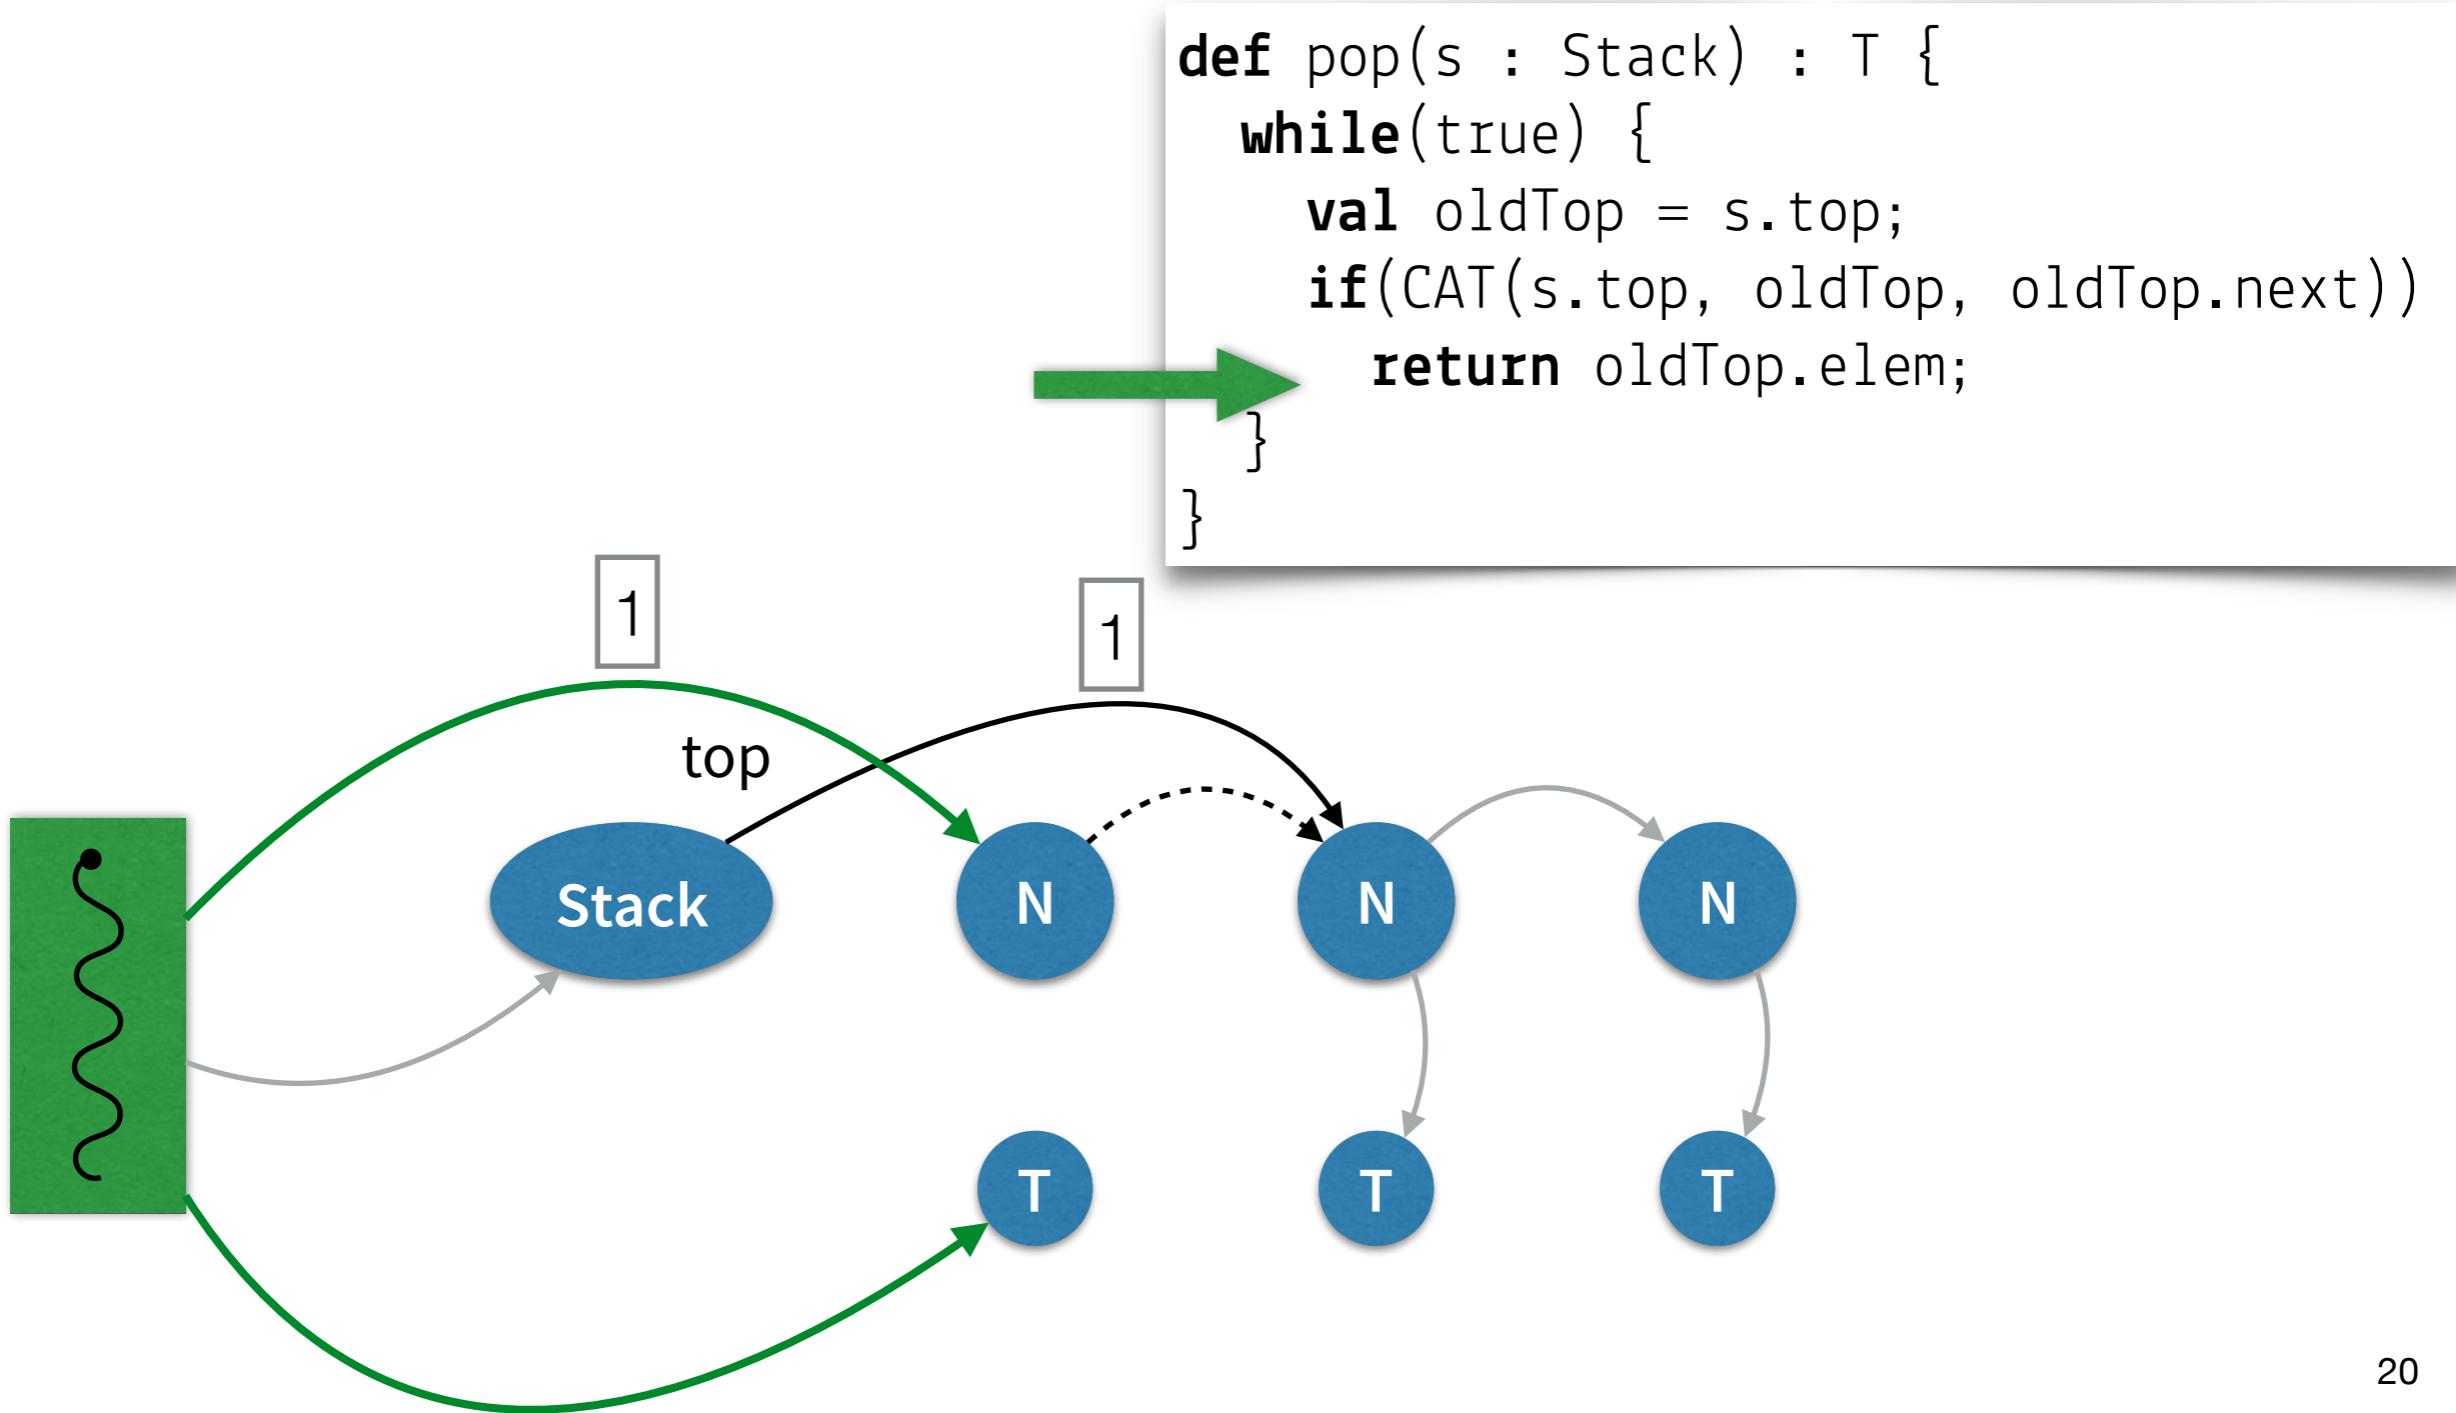
```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```



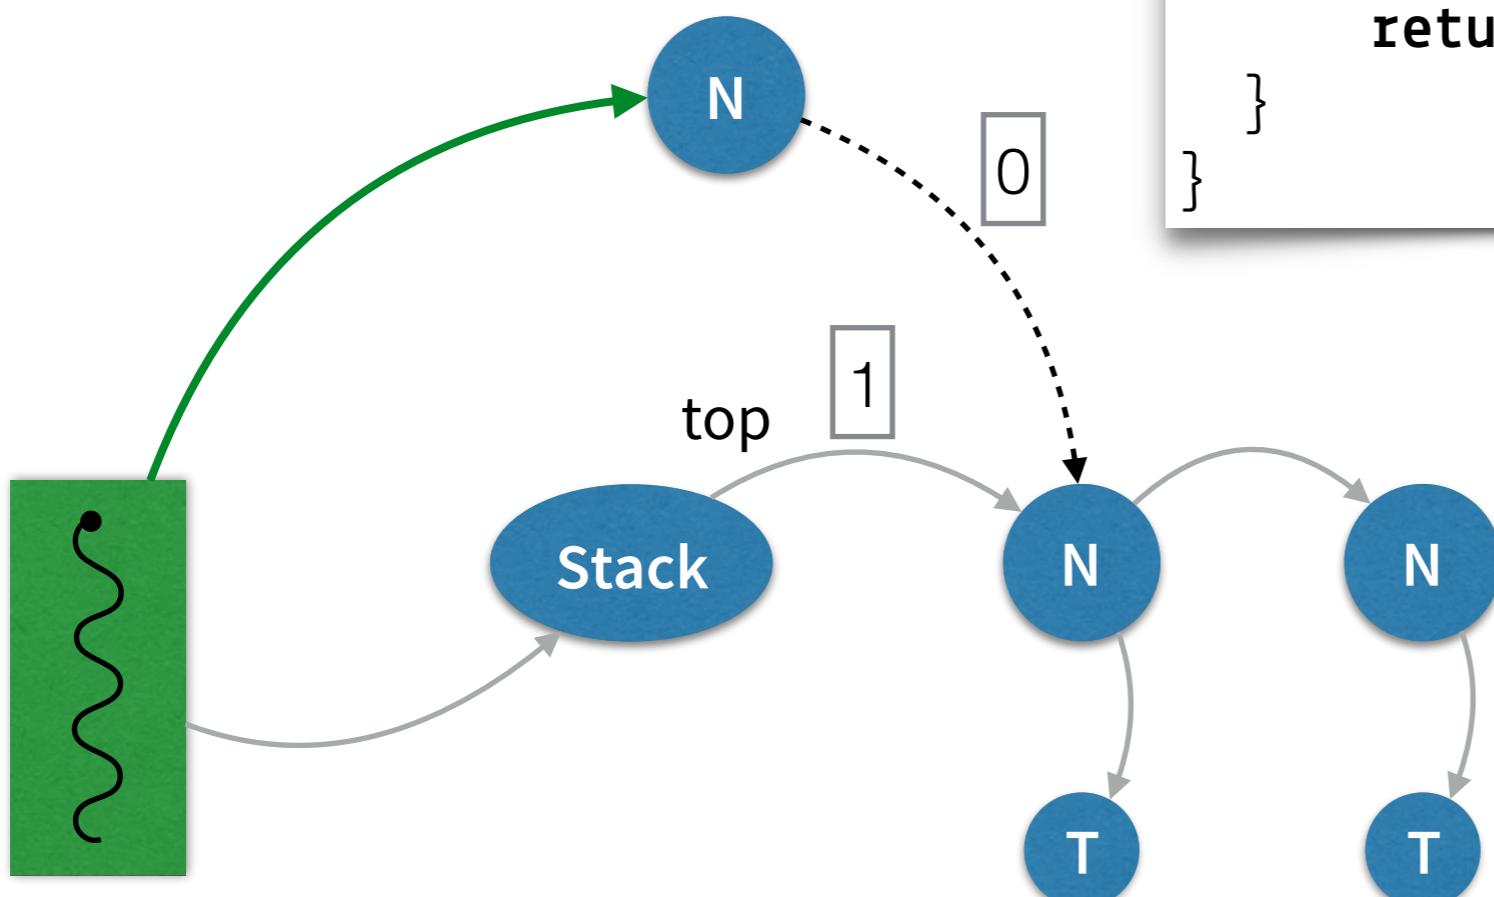
# Ownership Transfer in Action



# Ownership Transfer in Action



# Ownership Transfer in Action



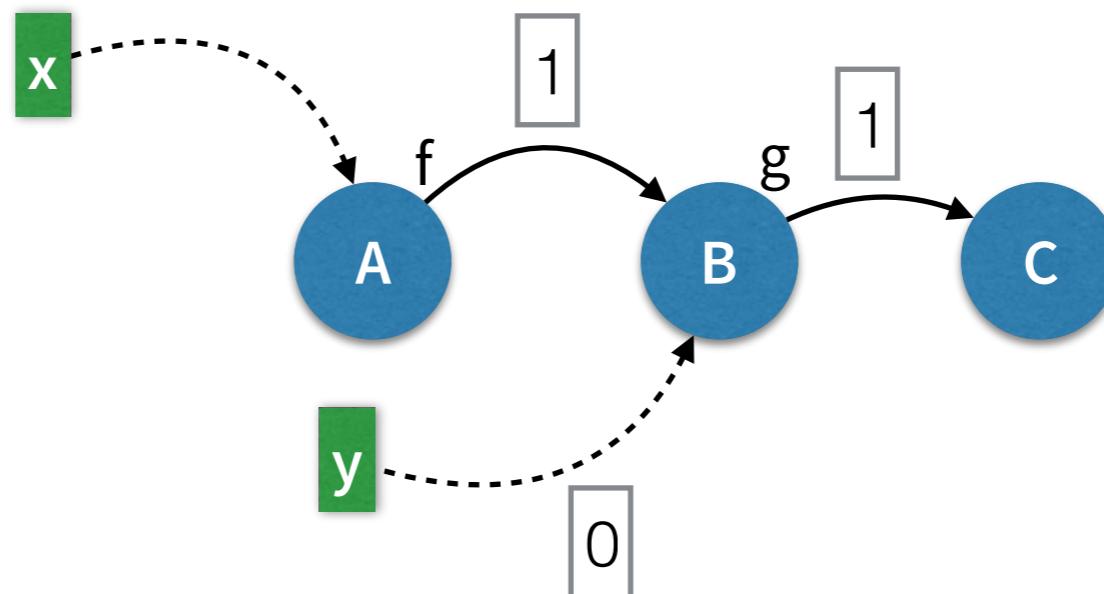
```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "

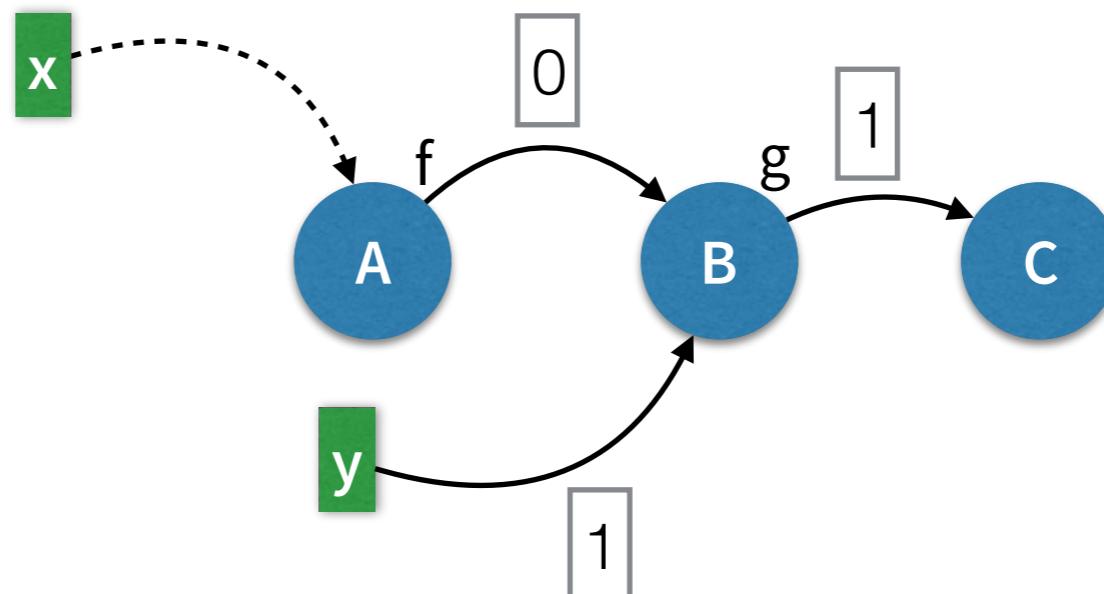


# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "

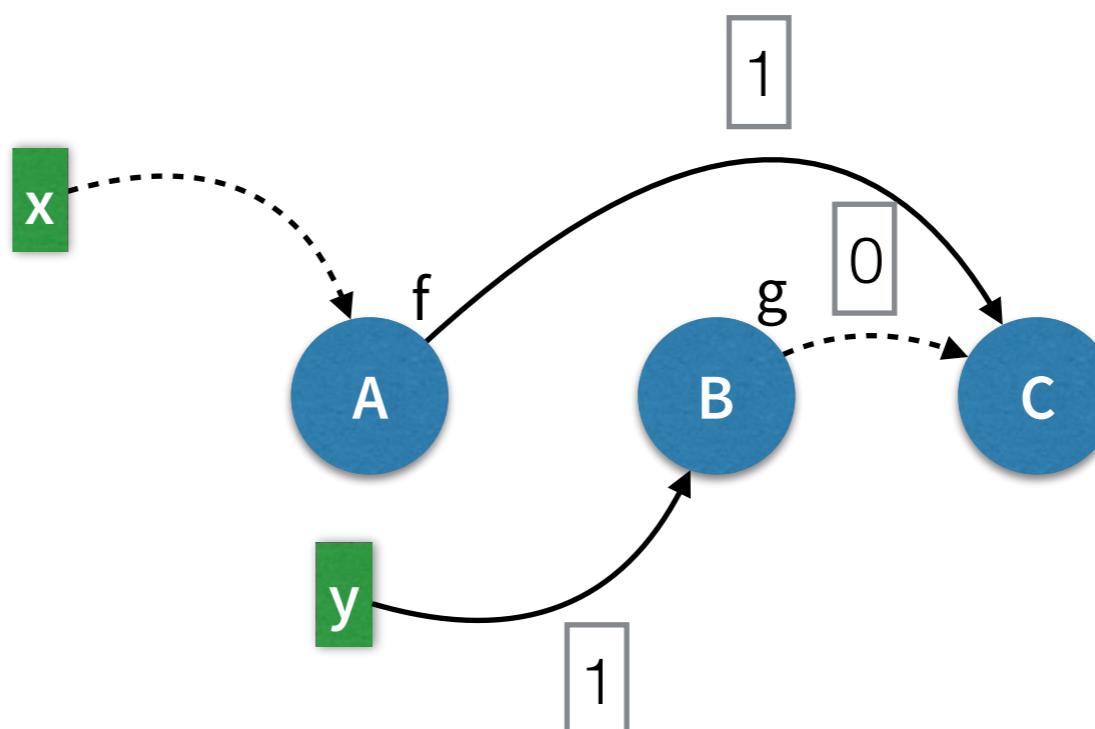


# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "



# CAT Atomic Ownership Transfer with ~~CAS~~

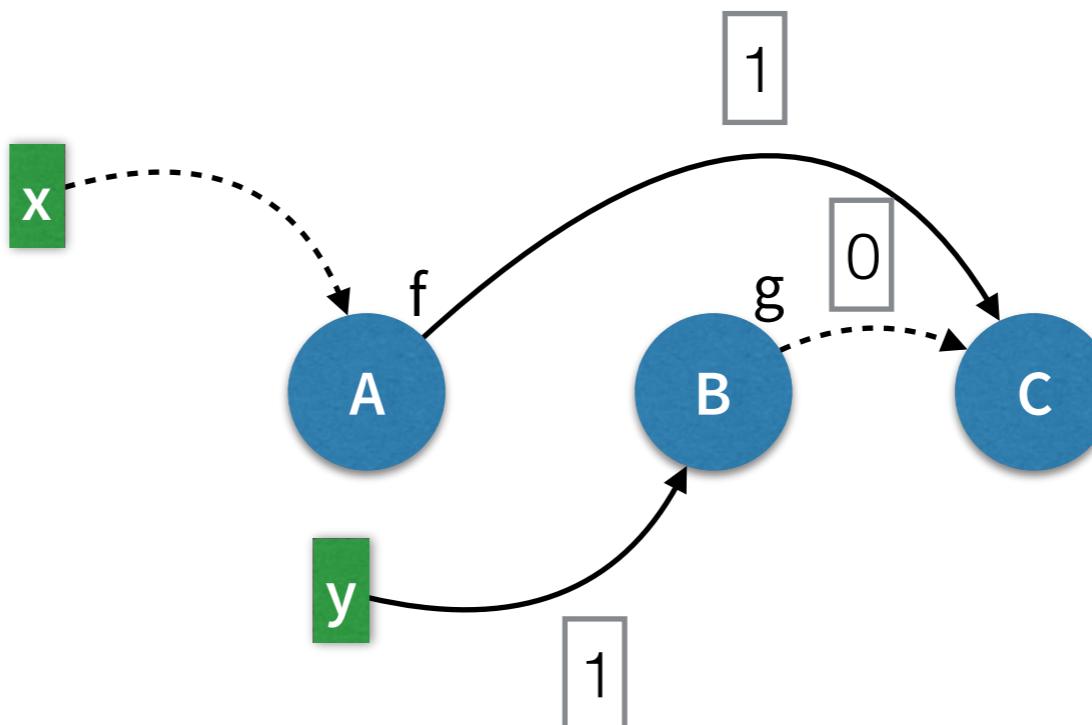
CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "

Link into a chain of objects:

CAT( $x.f$ ,  $y.g$ ,  $y$ ) — "Transfer ownership of  $x.f$  to  $y.g$  and of  $y$  to  $x.f$ "



# CAT Atomic Ownership Transfer with ~~CAS~~

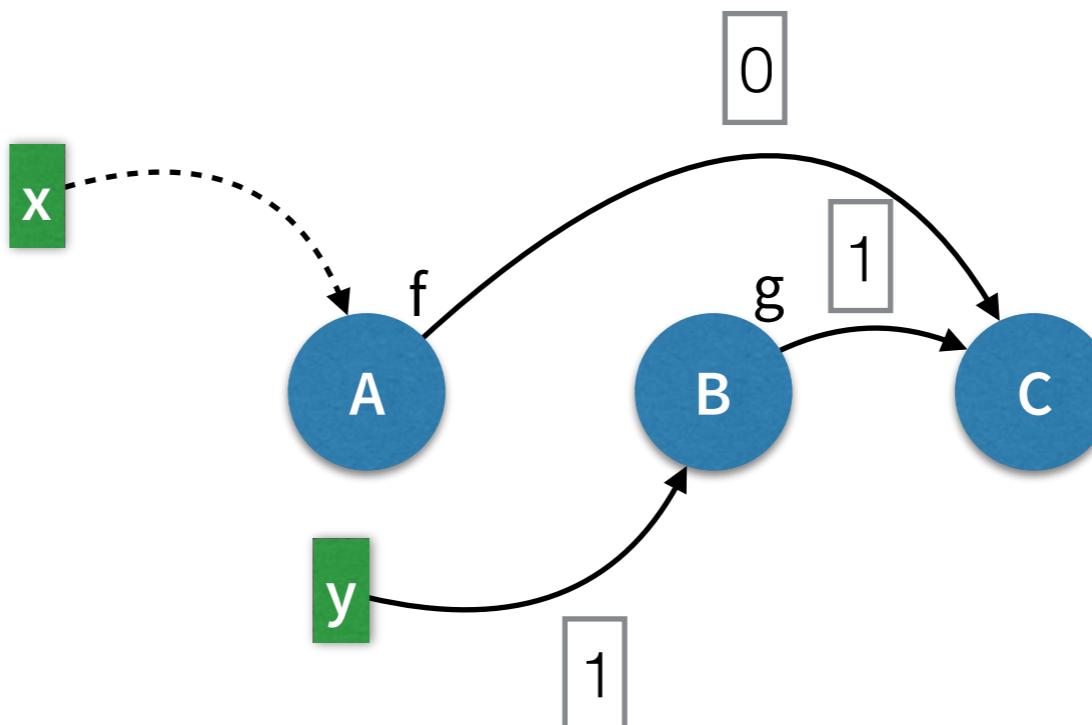
CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "

Link into a chain of objects:

CAT( $x.f$ ,  $y.g$ ,  $y$ ) — "Transfer ownership of  $x.f$  to  $y.g$  and of  $y$  to  $x.f$ "



# CAT Atomic Ownership Transfer with ~~CAS~~

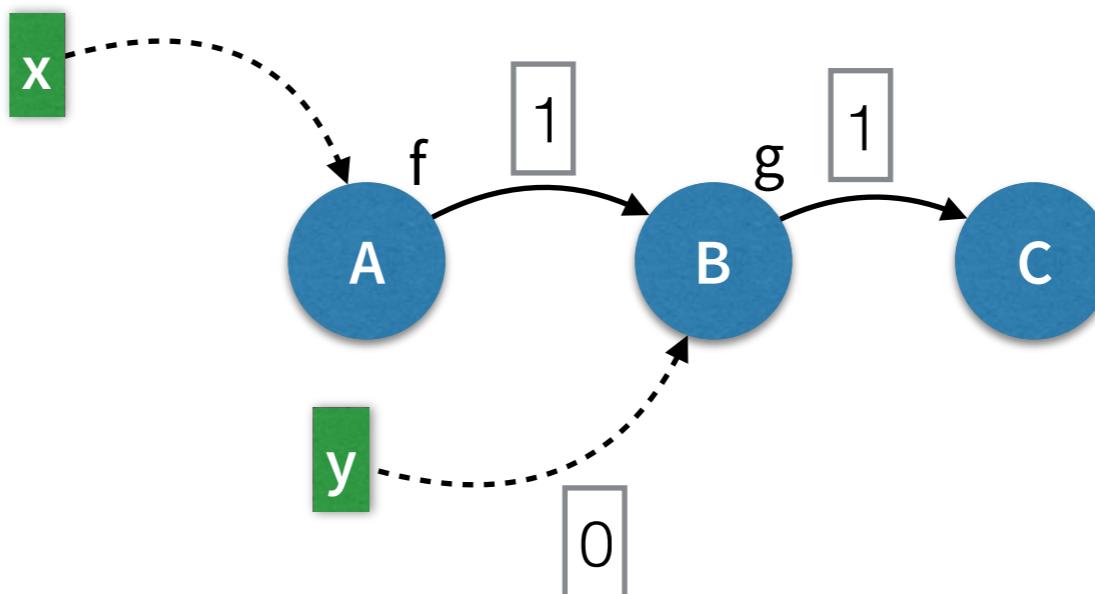
CAT(x.f, y, z) — "Transfer ownership of x.f to y and of z to x.f"

Unlink from a chain of objects:

CAT(x.f, y, y.g) — "Transfer ownership of x.f to y and of y.g to x.f"

Link into a chain of objects:

CAT(x.f, y.g, y) — "Transfer ownership of x.f to y.g and of y to x.f"



# CAT Atomic Ownership Transfer with ~~CAS~~

CAT( $x.f$ ,  $y$ ,  $z$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $z$  to  $x.f$ "

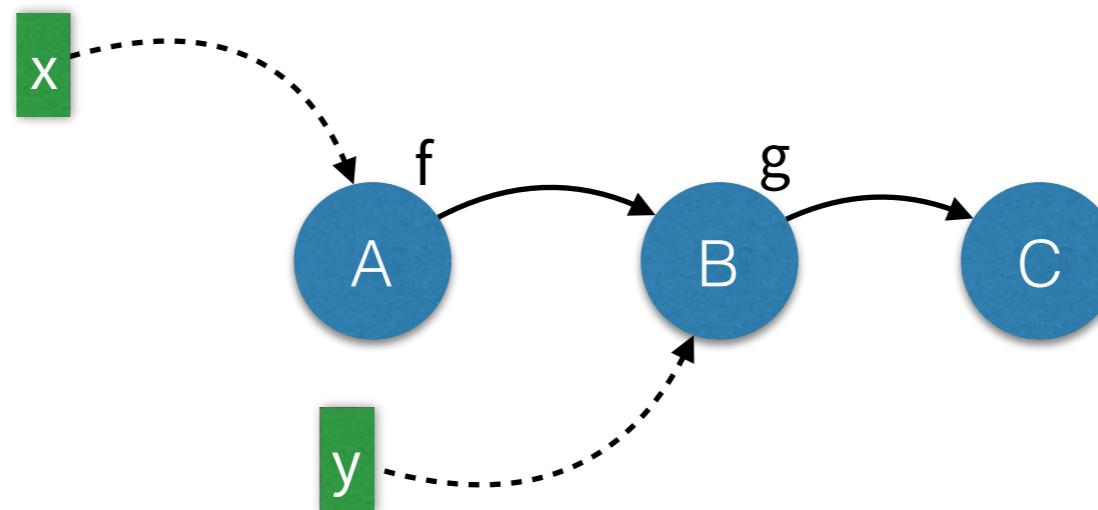
Unlink from a chain of objects:

CAT( $x.f$ ,  $y$ ,  $y.g$ ) — "Transfer ownership of  $x.f$  to  $y$  and of  $y.g$  to  $x.f$ "

Link into a chain of objects:

CAT( $x.f$ ,  $y.g$ ,  $y$ ) — "Transfer ownership of  $x.f$  to  $y.g$  and of  $y$  to  $x.f$ "

Three kinds of CATs



# Linear Ownership with Lock-free CAT



# LOLCAT



# Ownership in LOLCAT

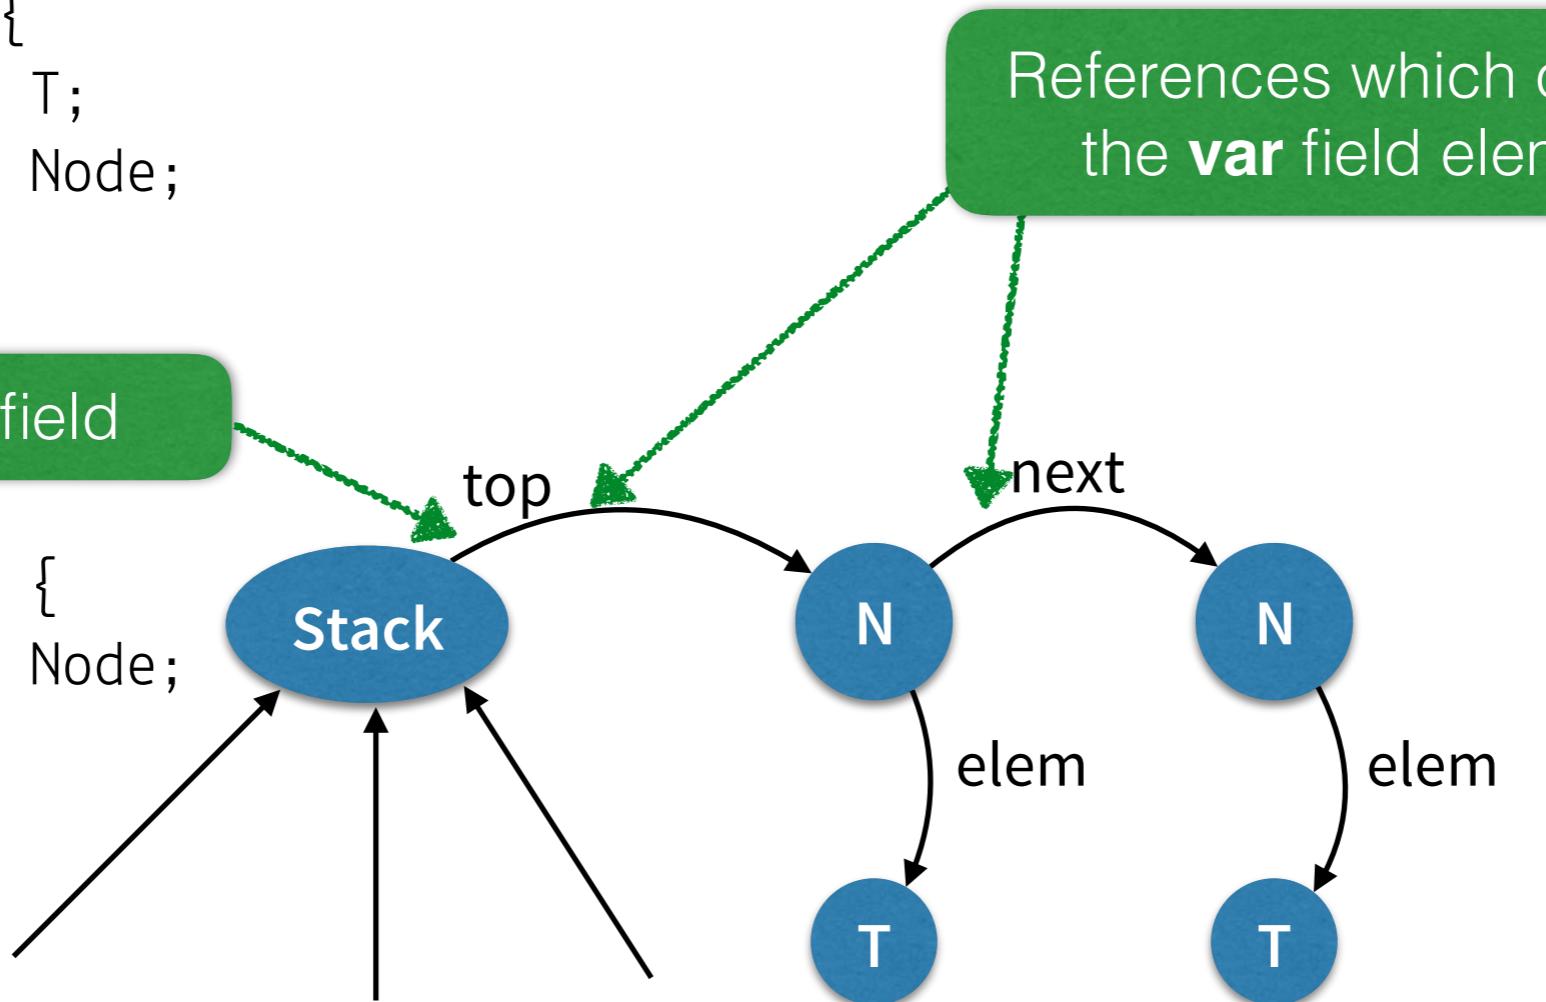
- References do not own objects, but rather the (mutable) fields of the objects

```
struct Node {  
    var elem : T;  
    val next : Node;  
}
```

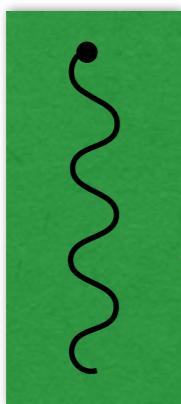
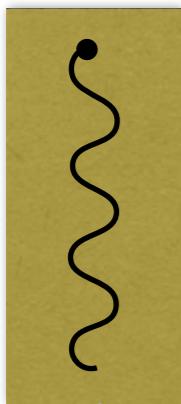
Contended field

```
struct Stack {  
    spec top : Node;  
}
```

References which own  
the **var** field elem



# Ownership Tracked using Flow-sensitive Types



top

Stack

N

N

N

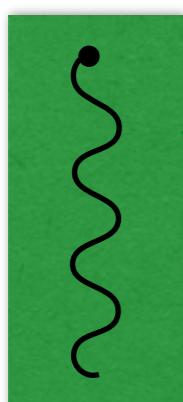
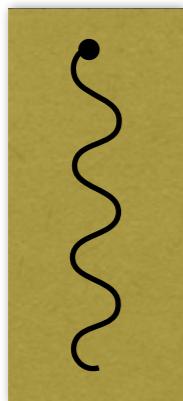
T

T

T

```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

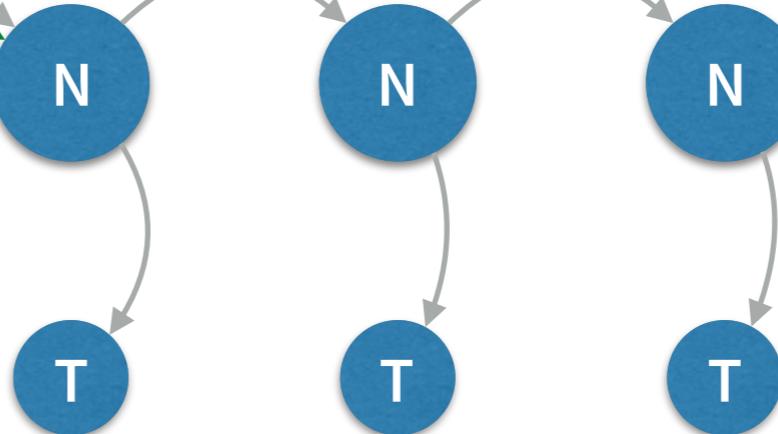
# Ownership Tracked using Flow-sensitive Types



Node | elem

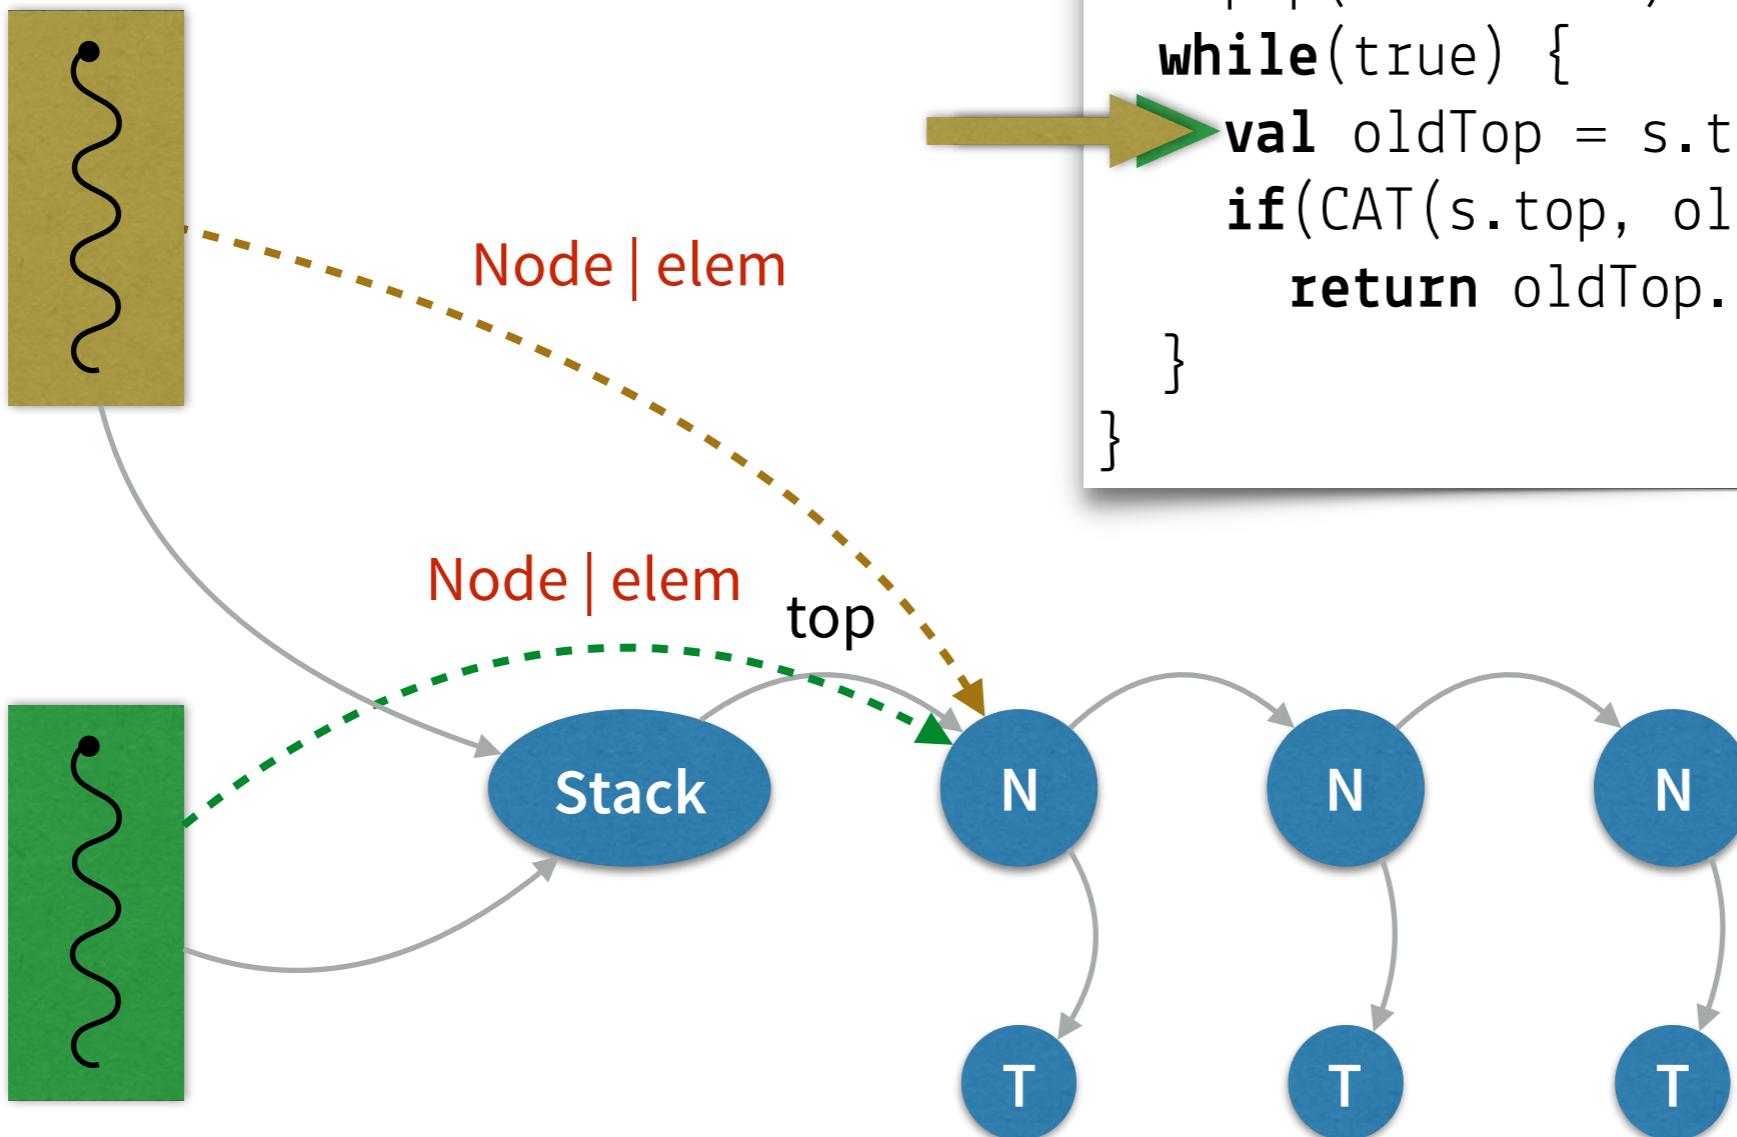
top

Stack

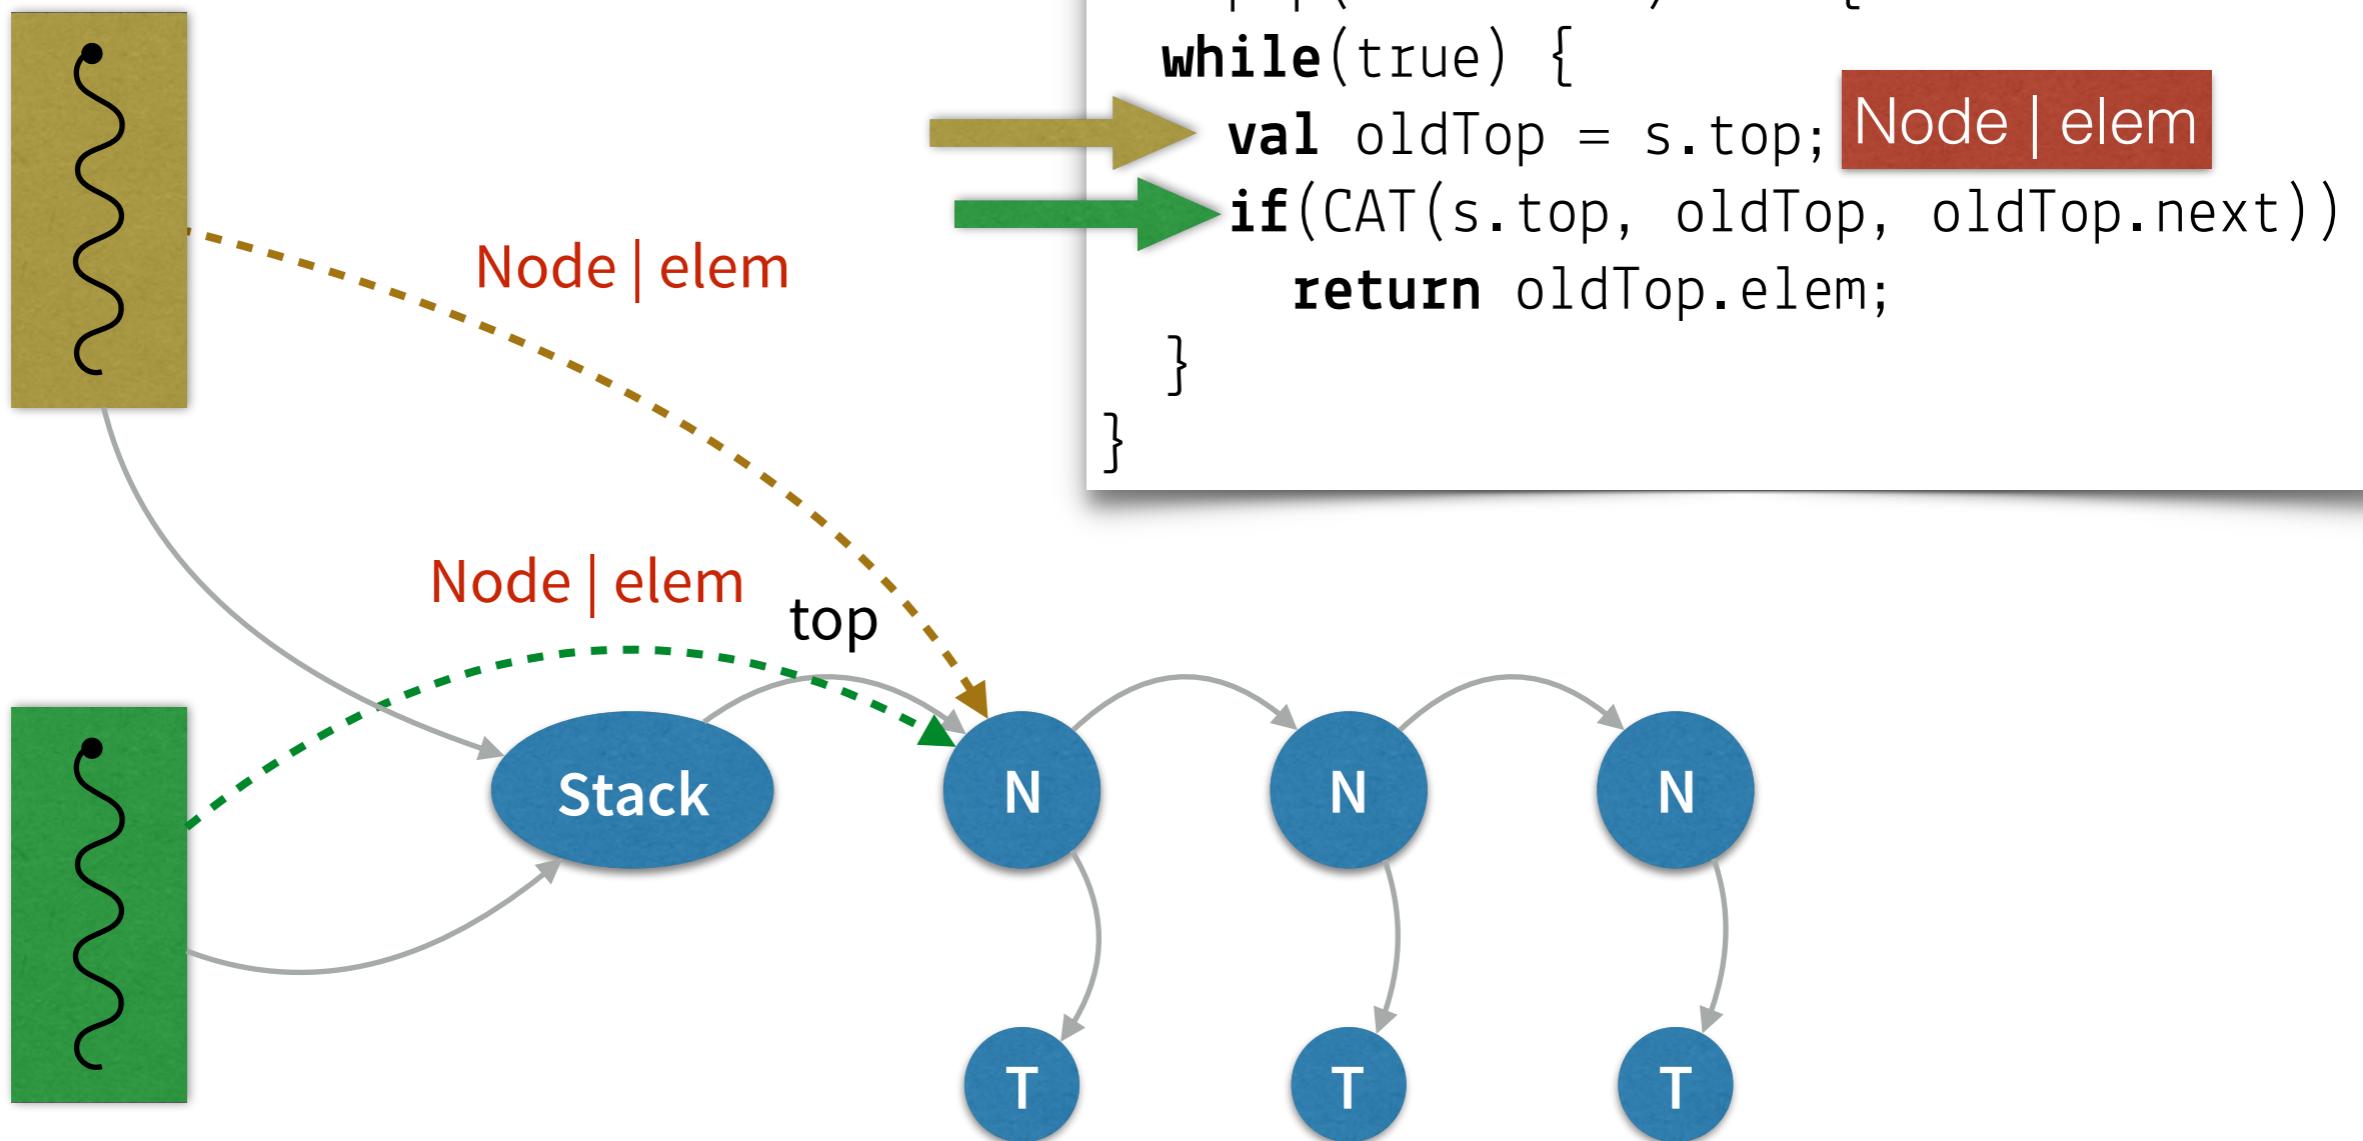


```
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top; Node | elem  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}
```

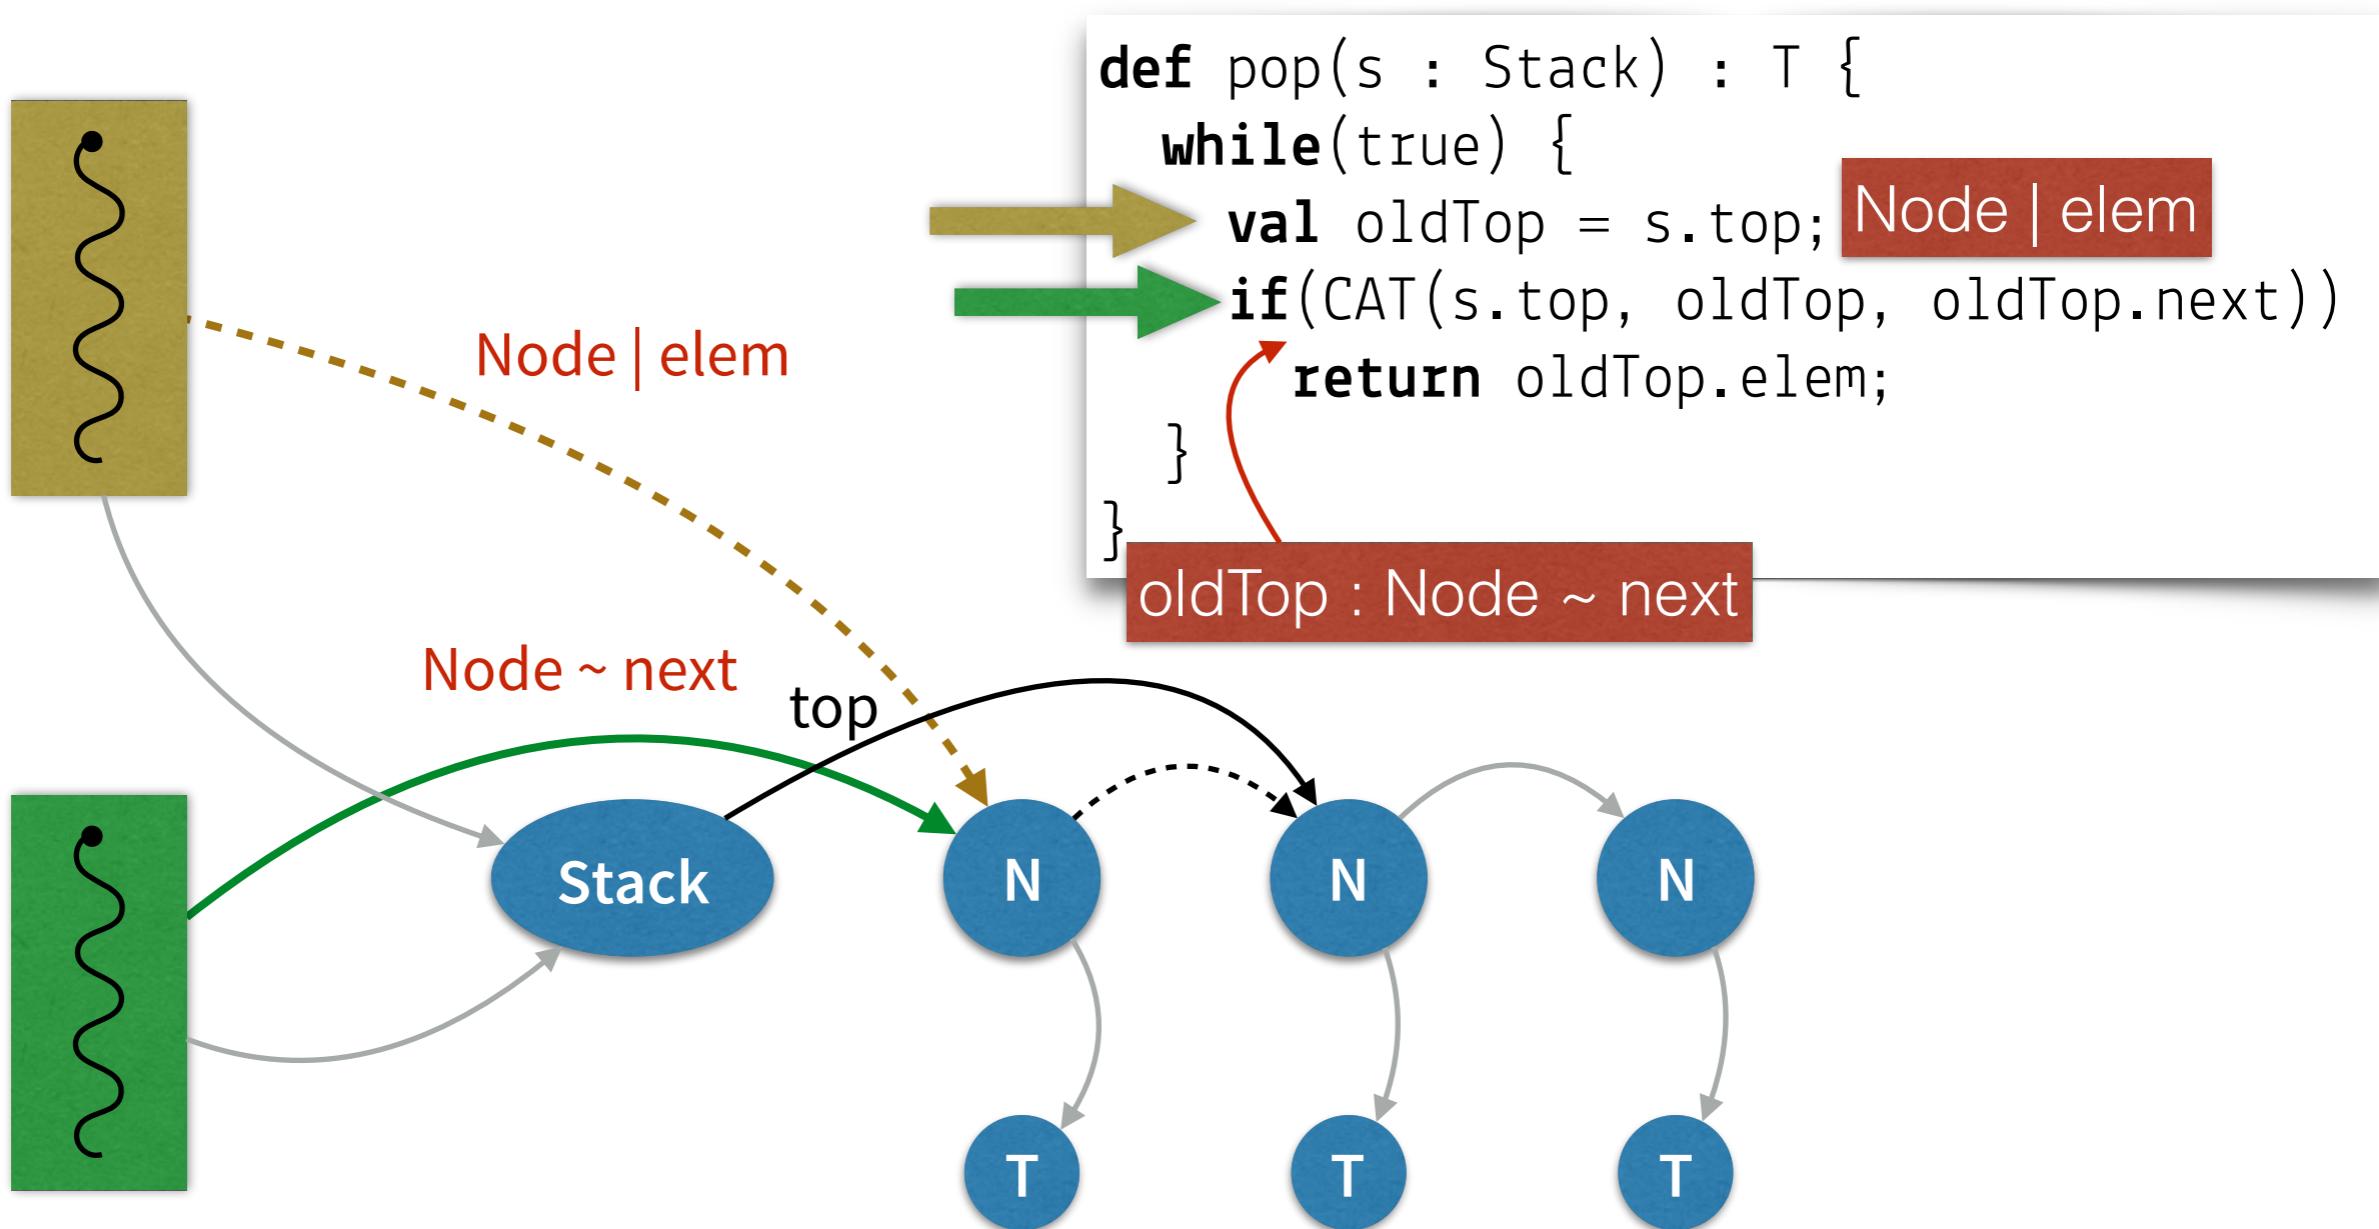
# Ownership Tracked using Flow-sensitive Types



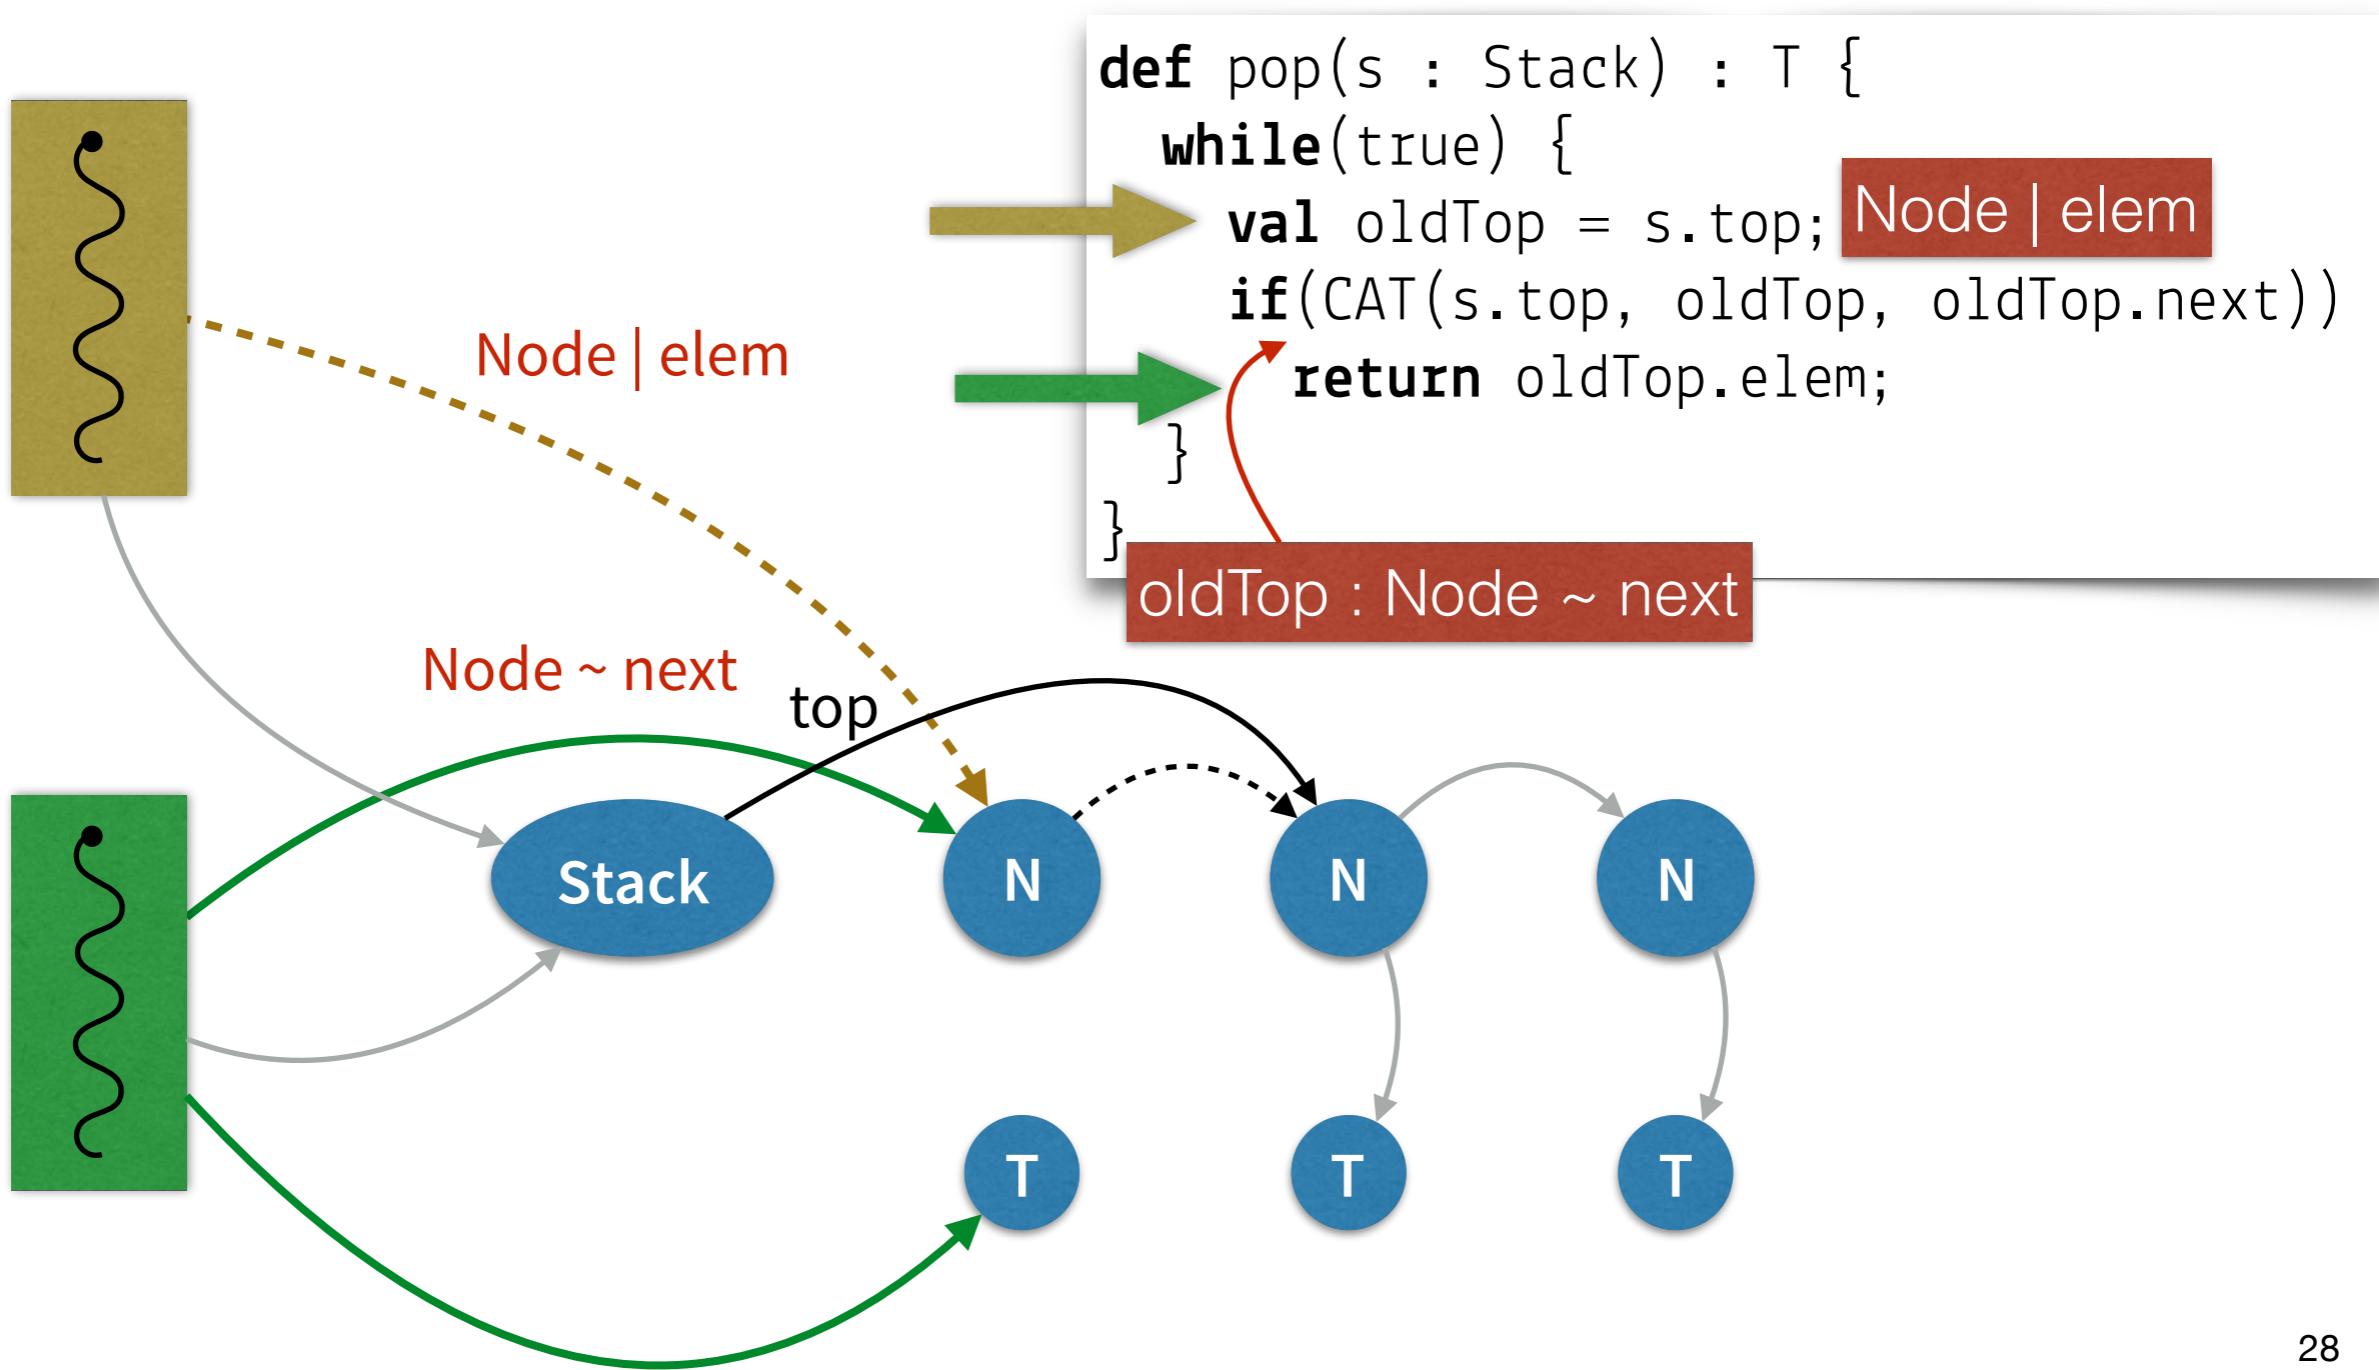
# Ownership Tracked using Flow-sensitive Types



# Ownership Tracked using Flow-sensitive Types



# Ownership Tracked using Flow-sensitive Types

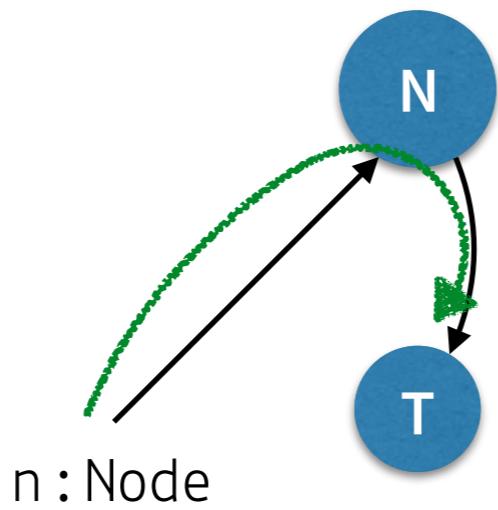


# Restricted References

---

- A reference of type Node can access the full node

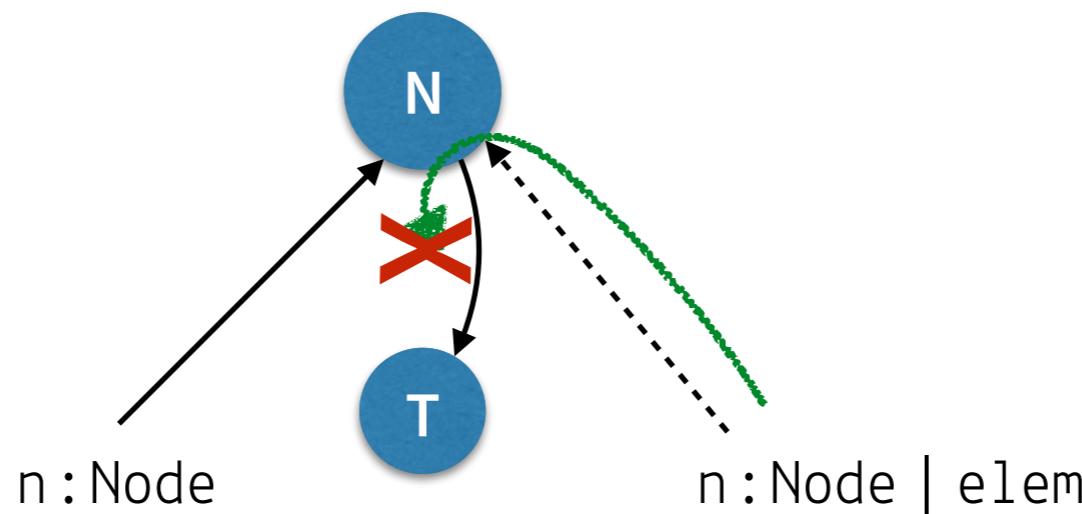
```
struct Node {  
    var elem : T;  
    val next : Node;  
}
```



# Restricted References

- A reference of type Node can access the full node
- A reference of type Node | elem cannot access the var field elem

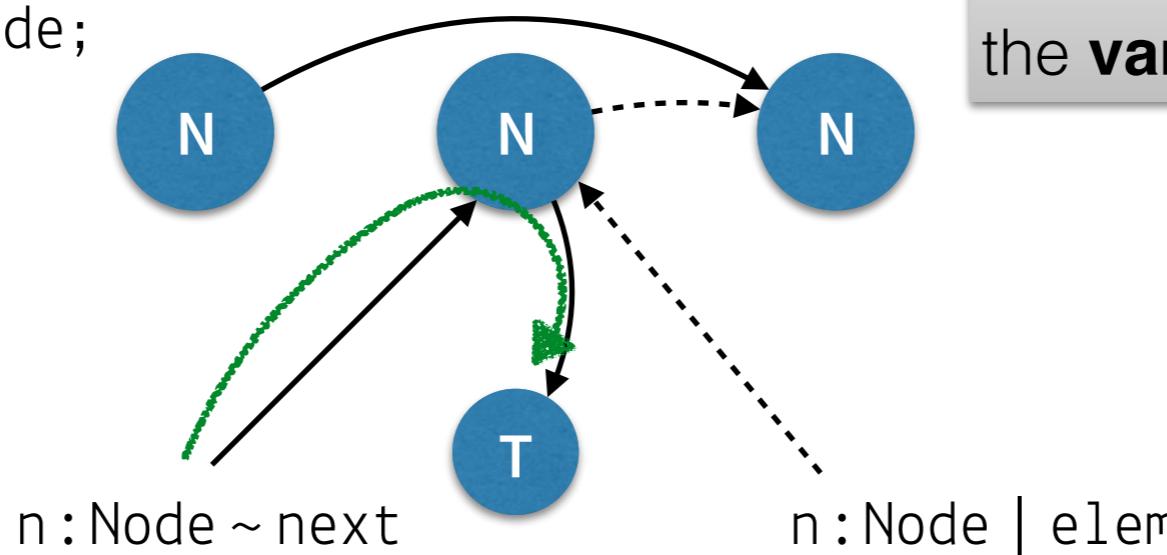
```
struct Node {  
    var elem : T;  
    val next : Node;  
}
```



# Restricted References

- A reference of type Node can access the full node
- A reference of type Node | elem cannot access the var field elem
- A reference of type Node ~ next can access the var field elem, but has a val field next without ownership

```
struct Node {  
    var elem : T;  
    val next : Node;  
}
```

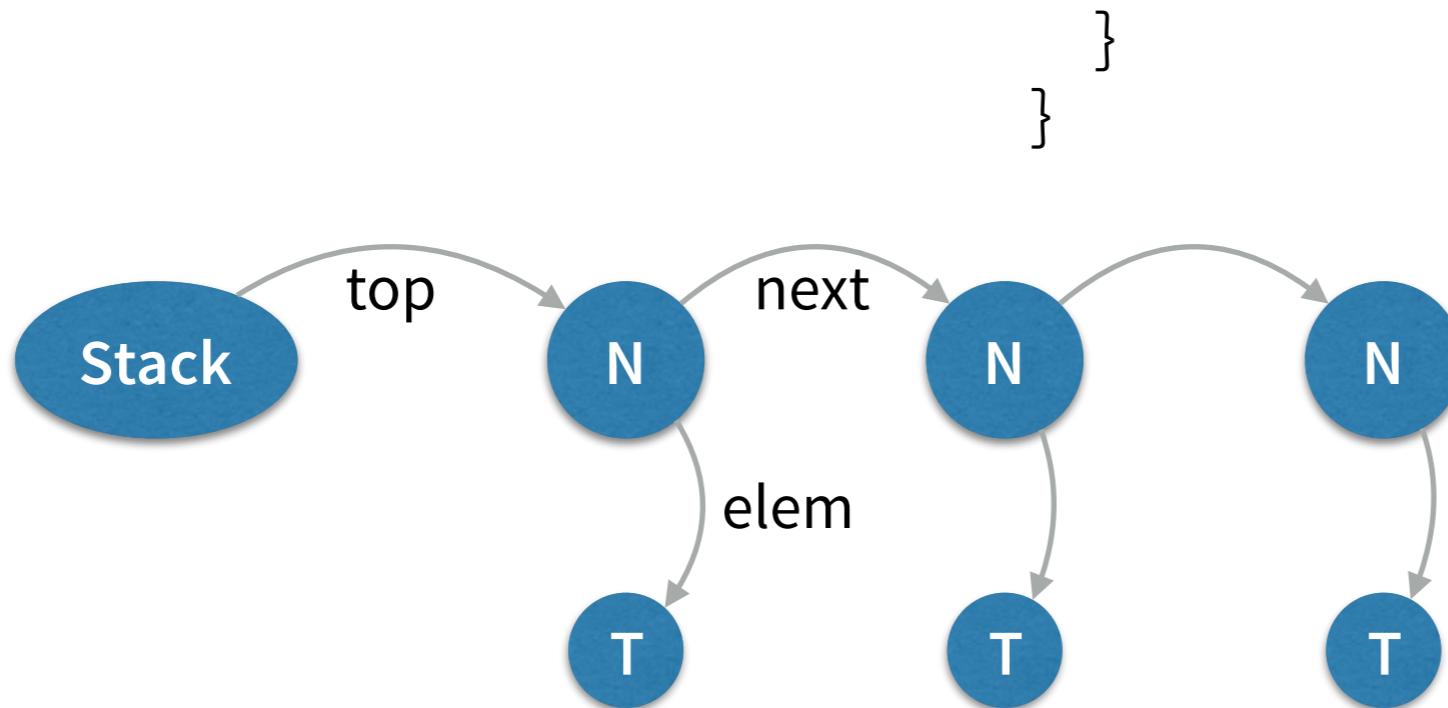


## Main invariant:

For all aliases of a Node, at most one alias has ownership of the **var** field elem

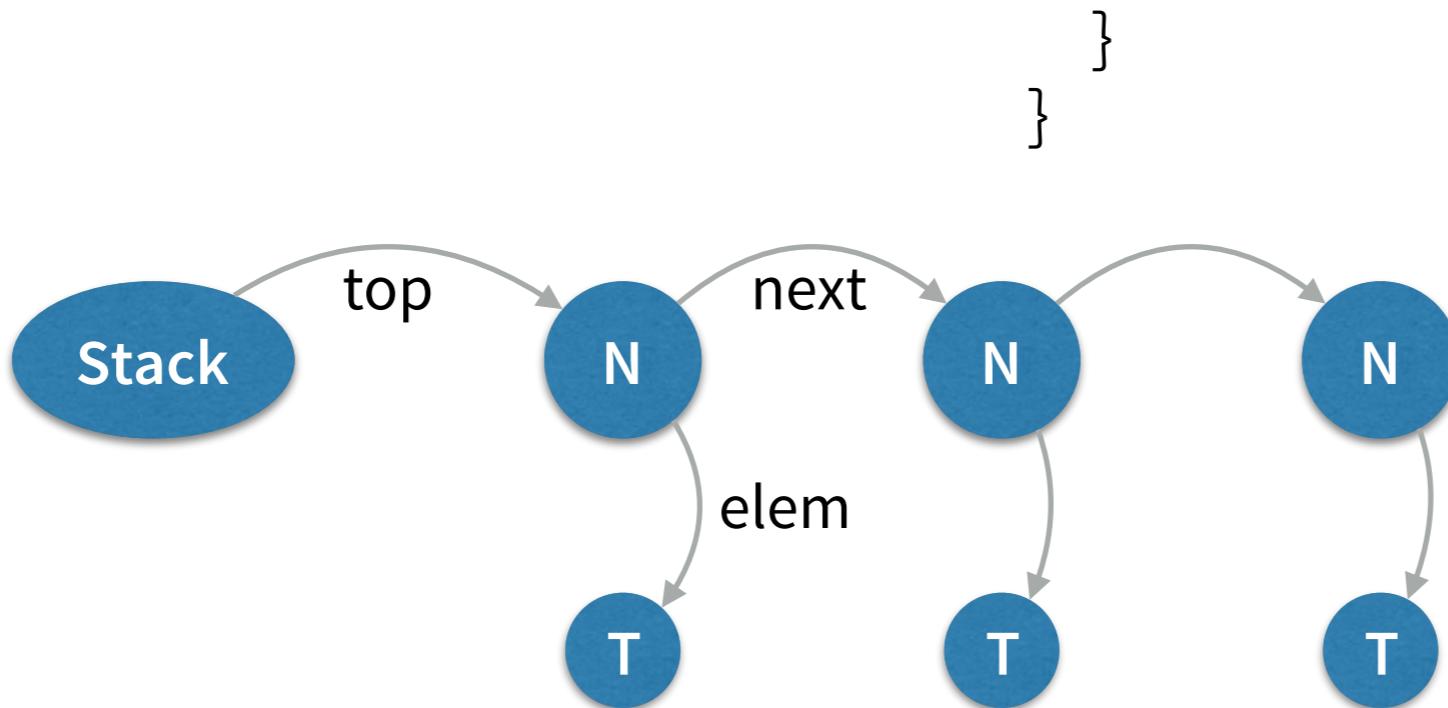
# A Lock-free Stack

```
struct Node {           struct Stack {  
    var elem : T;         var top : Node;  
    val next : Node;     }  
}  
  
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAS(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}  
  
def push(s : Stack, x : T) : void {  
    val newTop = new Node(x);  
    newTop.next = s.top;  
    while(true) {  
        if(CAS(s.top, newTop.next, newTop))  
            return;  
        newTop.next = s.top;  
    }  
}
```



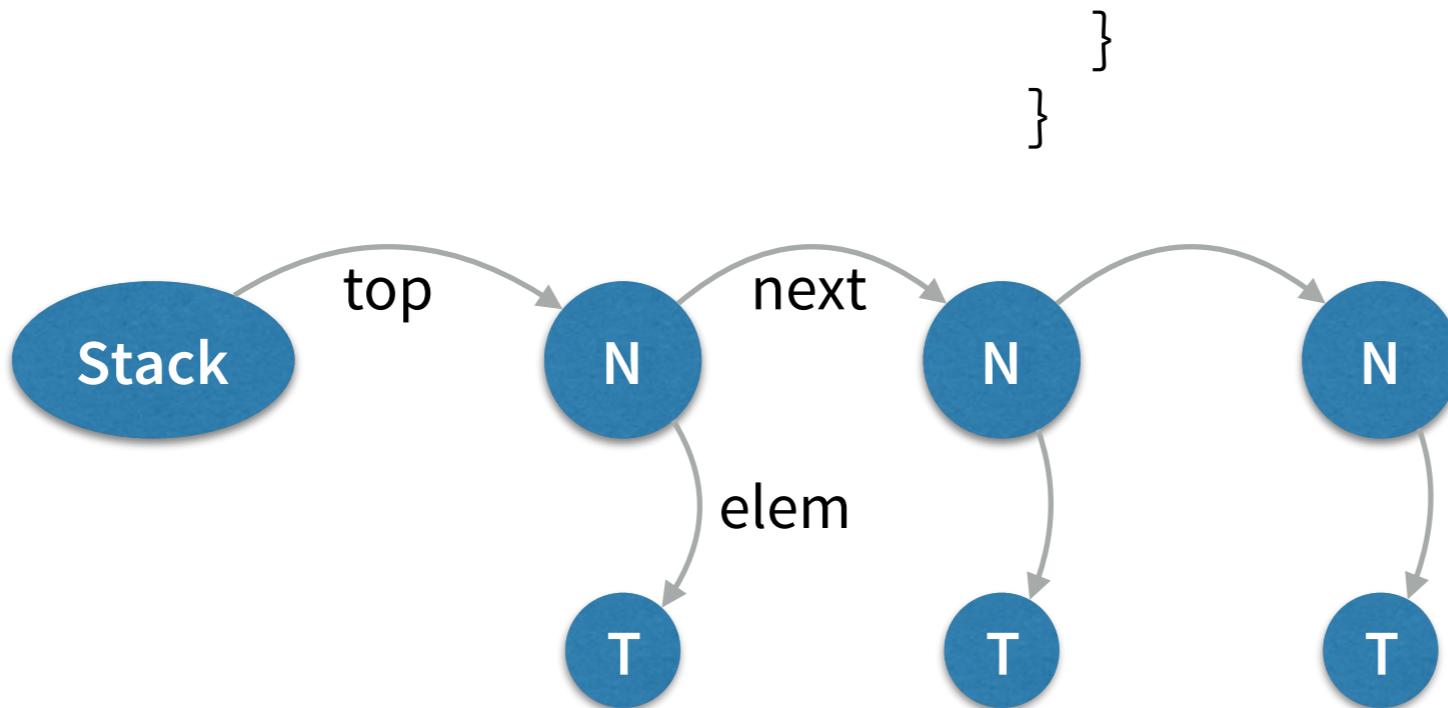
# A Lock-free Stack in LOLCAT

```
struct Node {           struct Stack {  
    var elem : T;         spec top : Node;  
    val next : Node;     }  
}  
  
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return oldTop.elem;  
    }  
}  
  
def push(s : Stack, x : T) : void {  
    val newTop = new Node(x);  
    newTop.next = s.top;  
    while(true) {  
        if(CAT(s.top, newTop.next, newTop))  
            return;  
        newTop.next = s.top;  
    }  
}
```



# A Lock-free Stack in LOLCAT (more explicit)

```
struct Node {           struct Stack {  
    var elem : T;         spec top : Node;  
    val next : Node;     }  
}  
  
def pop(s : Stack) : T {  
    while(true) {  
        val oldTop = speculate s.top;  
        if(CAT(s.top, oldTop, oldTop.next))  
            return consume oldTop.elem;  
    }  
}  
  
def push(s : Stack, x : T) : void {  
    val newTop = new Node(consume x);  
    newTop.next = speculate s.top;  
    while(true) {  
        if(CAT(s.top, newTop.next, newTop))  
            return;  
        newTop.next = speculate s.top;  
    }  
}
```



# Also in the Paper

- LOLCAT formalized and proven sound
  - Linear ownership: at most one alias may access the `var` fields of an object
- Three fundamental lock-free data-structures
  - Treiber Stack
  - Michael—Scott Queue
  - Tim Harris List
- Prototype implementation in Encore
  - OO support



# Not in the Paper

---

- Additional examples (implemented in Encore)
  - Skip list [Fomitchev & Ruppert]
  - BST [Ellen et al.]
  - Lazy list-based set [Heller et al.]
  - Spin-locks
- Type-Assisted Automatic Garbage Collection for Lock-Free Data Structures  
[Yang & Wrigstad, ISMM'17]
  - Pluggable concurrent garbage collector
  - Relies on LOLCAT types to track when objects may be collected



# Summary

---

- Relaxed linearity: unbounded aliasing with linear ownership
  - Access to mutable (**var**) fields is always exclusive
  - Compare-and-swap for atomic ownership transfer
  - Controlled sharing of certain fields (**val** and **spec**)
  - Guarantee absence of (uncontrolled) data-races in lock-free data structures
- LOLCAT provides meaningful types for common patterns in lock-free programming
  - Speculation (Node | elem)
  - Publication (CAT(x.f, y, z))
  - Acquisition (Node ~ next)
  - Types inferred (modulo struct declarations)
- Prototype implementation in Encore
  - Open source: try it yourself!

# Glad Midsommar!



<http://www.kristianstadik.se/www/wordpress/wp-content/uploads/2016/06/midsommar.jpg>



# Summary

---

- Relaxed linearity: unbounded aliasing with linear ownership
  - Access to mutable (**var**) fields is always exclusive
  - Compare-and-swap for atomic ownership transfer
  - Controlled sharing of certain fields (**val** and **spec**)
  - Guarantee absence of (uncontrolled) data-races in lock-free data structures
- LOLCAT provides meaningful types for common patterns in lock-free programming
  - Speculation (Node | elem)
  - Publication (CAT(x.f, y, z))
  - Acquisition (Node ~ next)
  - Types inferred (modulo struct declarations)
- Prototype implementation in Encore
  - Open source: try it yourself!

